SOFTWARE UNIVERSITY

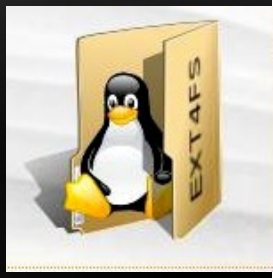

4 – Linux File Systems

Marian Marinov
CEO of 1H Ltd.
mm@1h.com

Stoyan Stoyanov
System Administrator
sto [at] softuni.bg

SOFTWARE UNIVERSITY

Agenda

- File System Architecture
- Virtual File System Layer
- Directory Structure
- Mount operations

SOFTWARE UNIVERSITY

Agenda

- File system types
  - Local
  - Log-structured - NAND
  - Pseudo
  - Network
  - Cluster
  - Distributed

File System Architecture

There is NO silver bullet!

SOFTWARE UNIVERSITY

File System Architecture

There is NO DIRECTORY!

SOFTWARE UNIVERSITY

File System Architecture

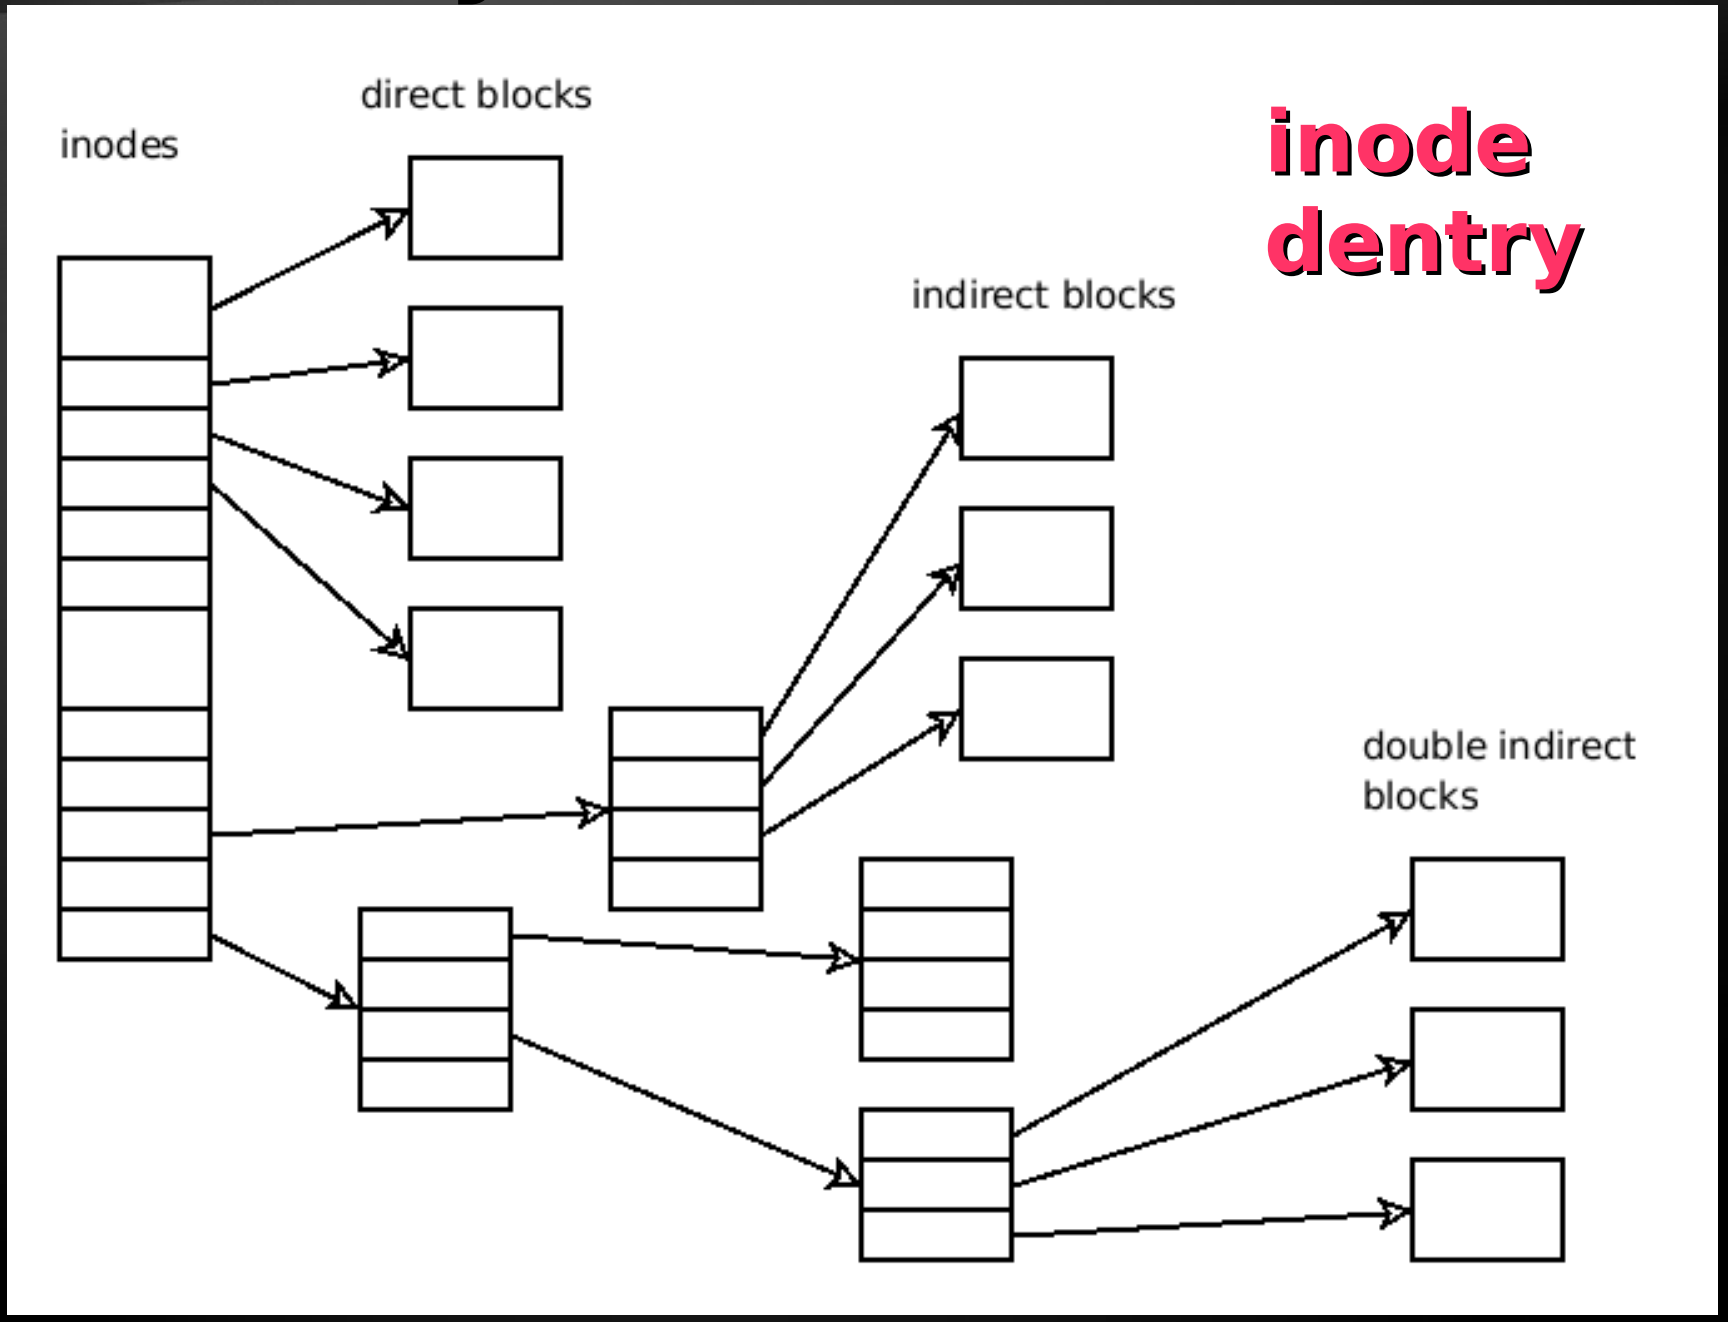

SOFTWARE UNIVERSITY

File System Architecture

Super Blocks

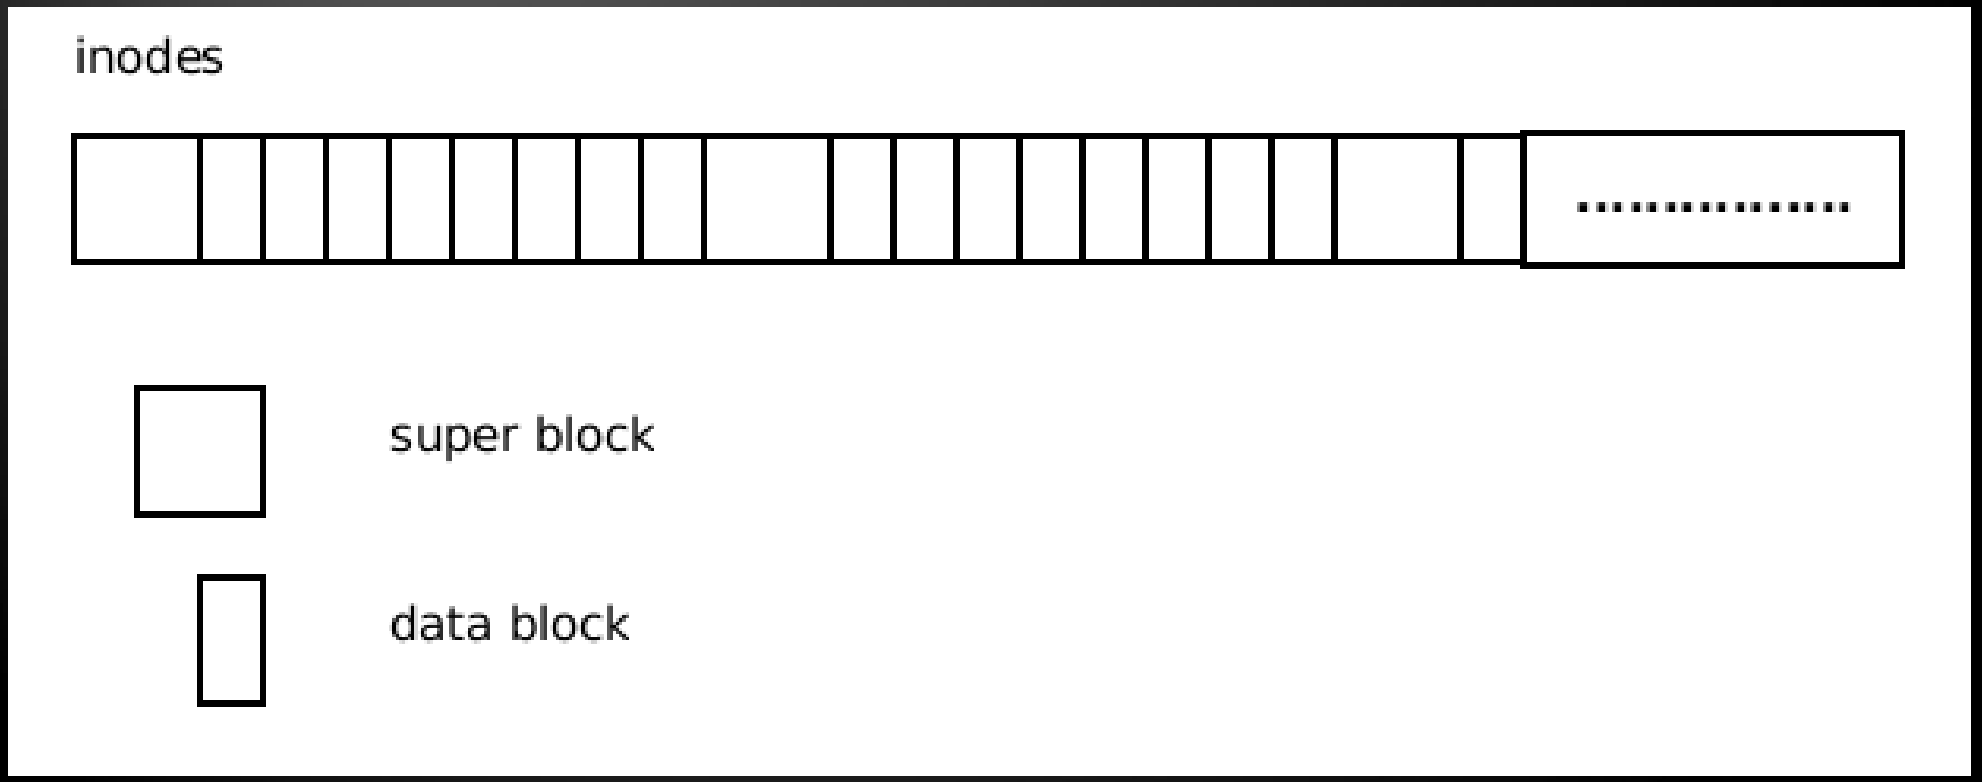

SOFTWARE UNIVERSITY

File System Architecture

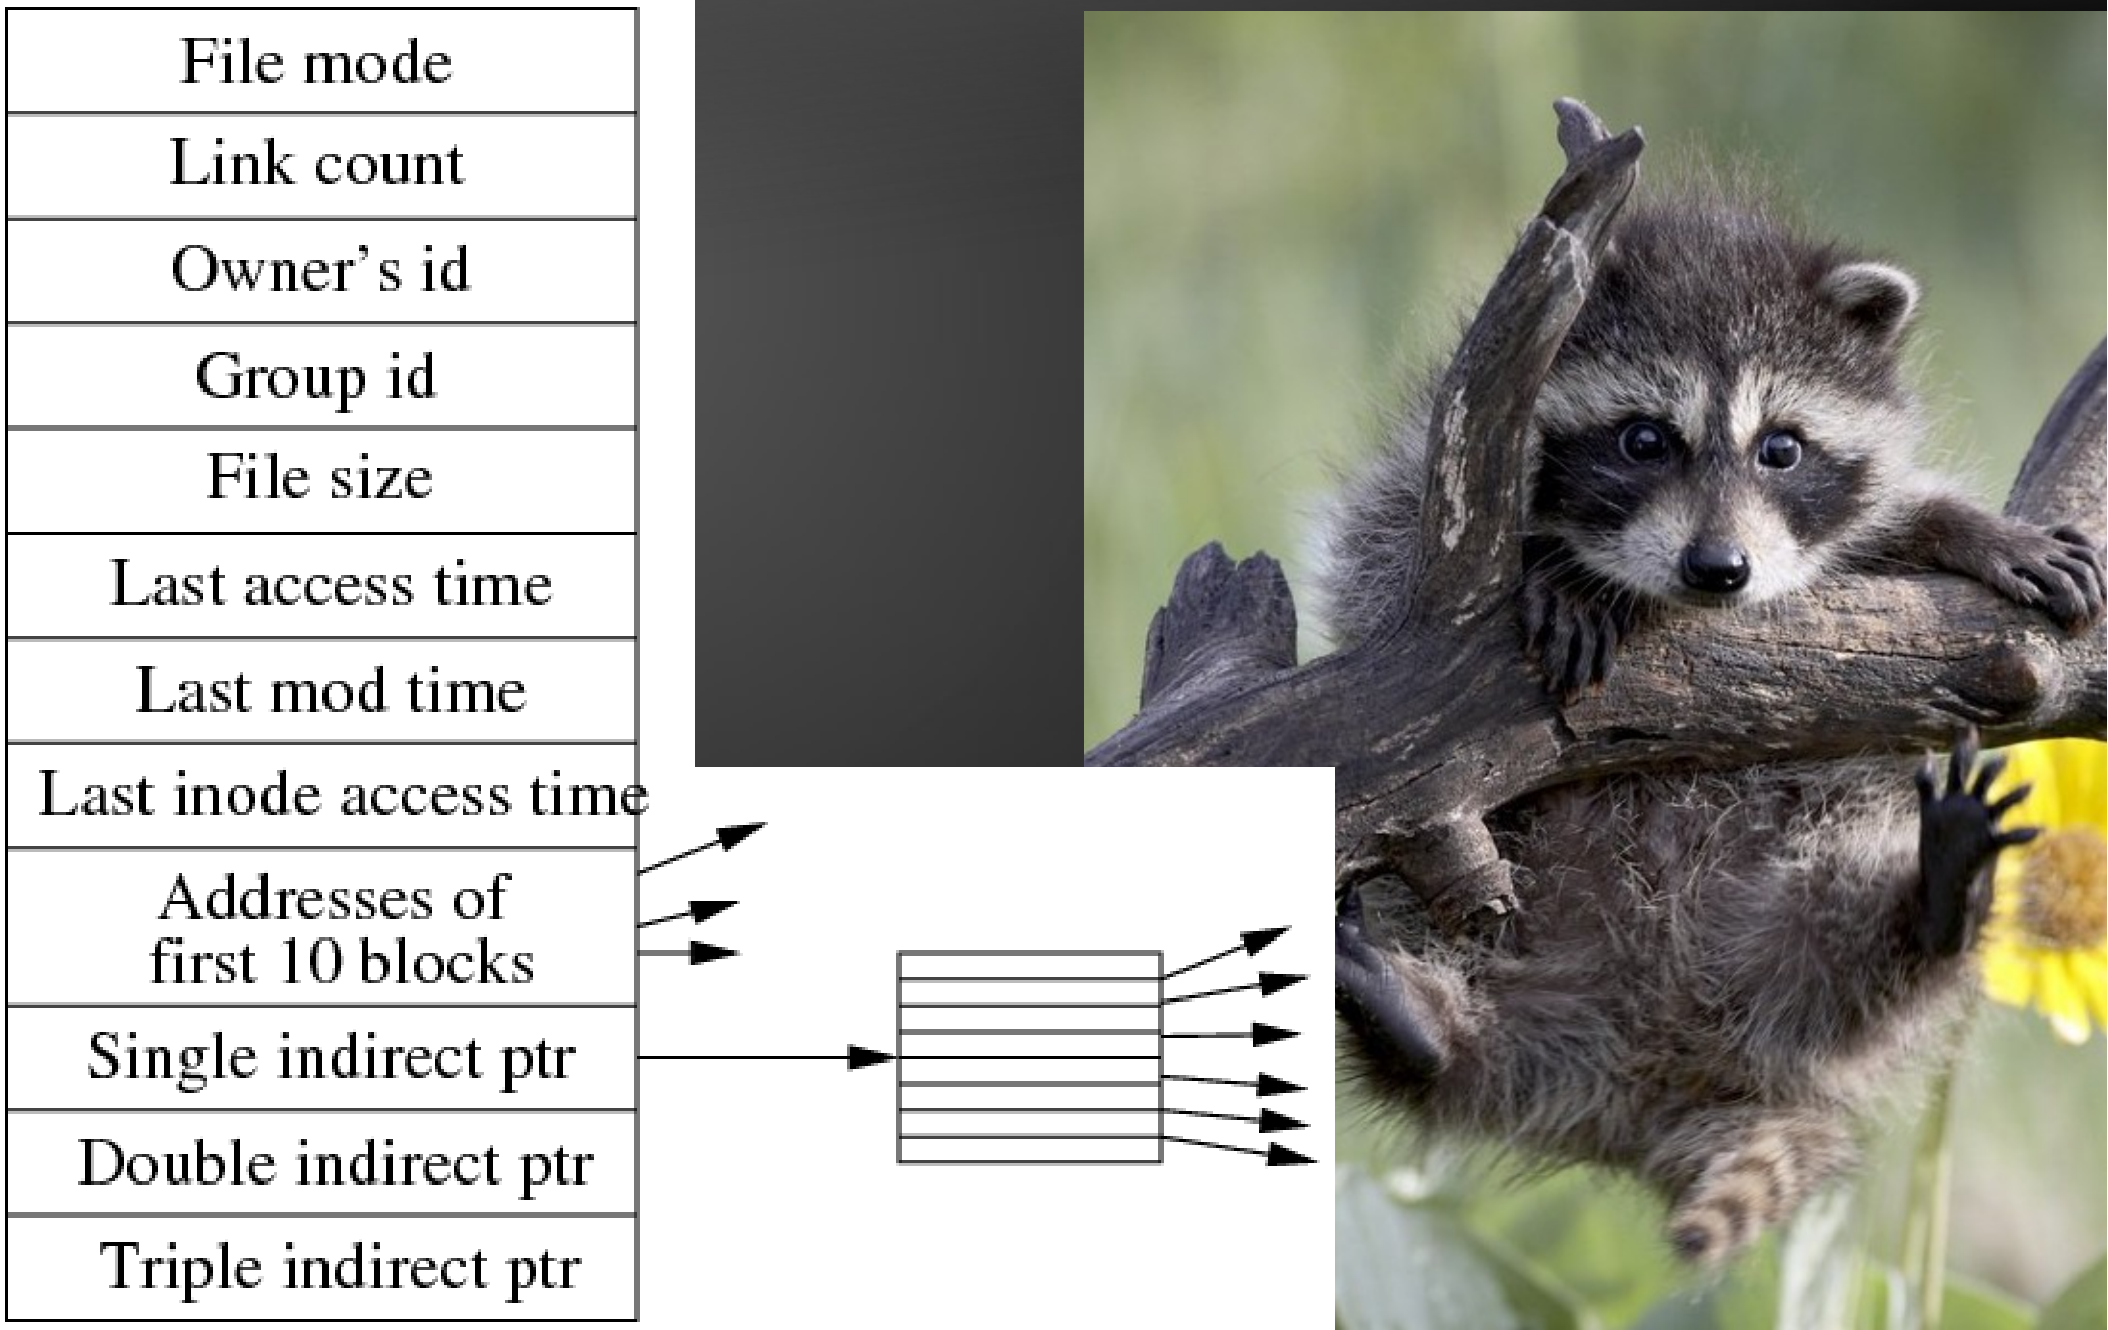

SOFTWARE UNIVERSITY

File System Architecture

Soft/symlinks
Hardlinks

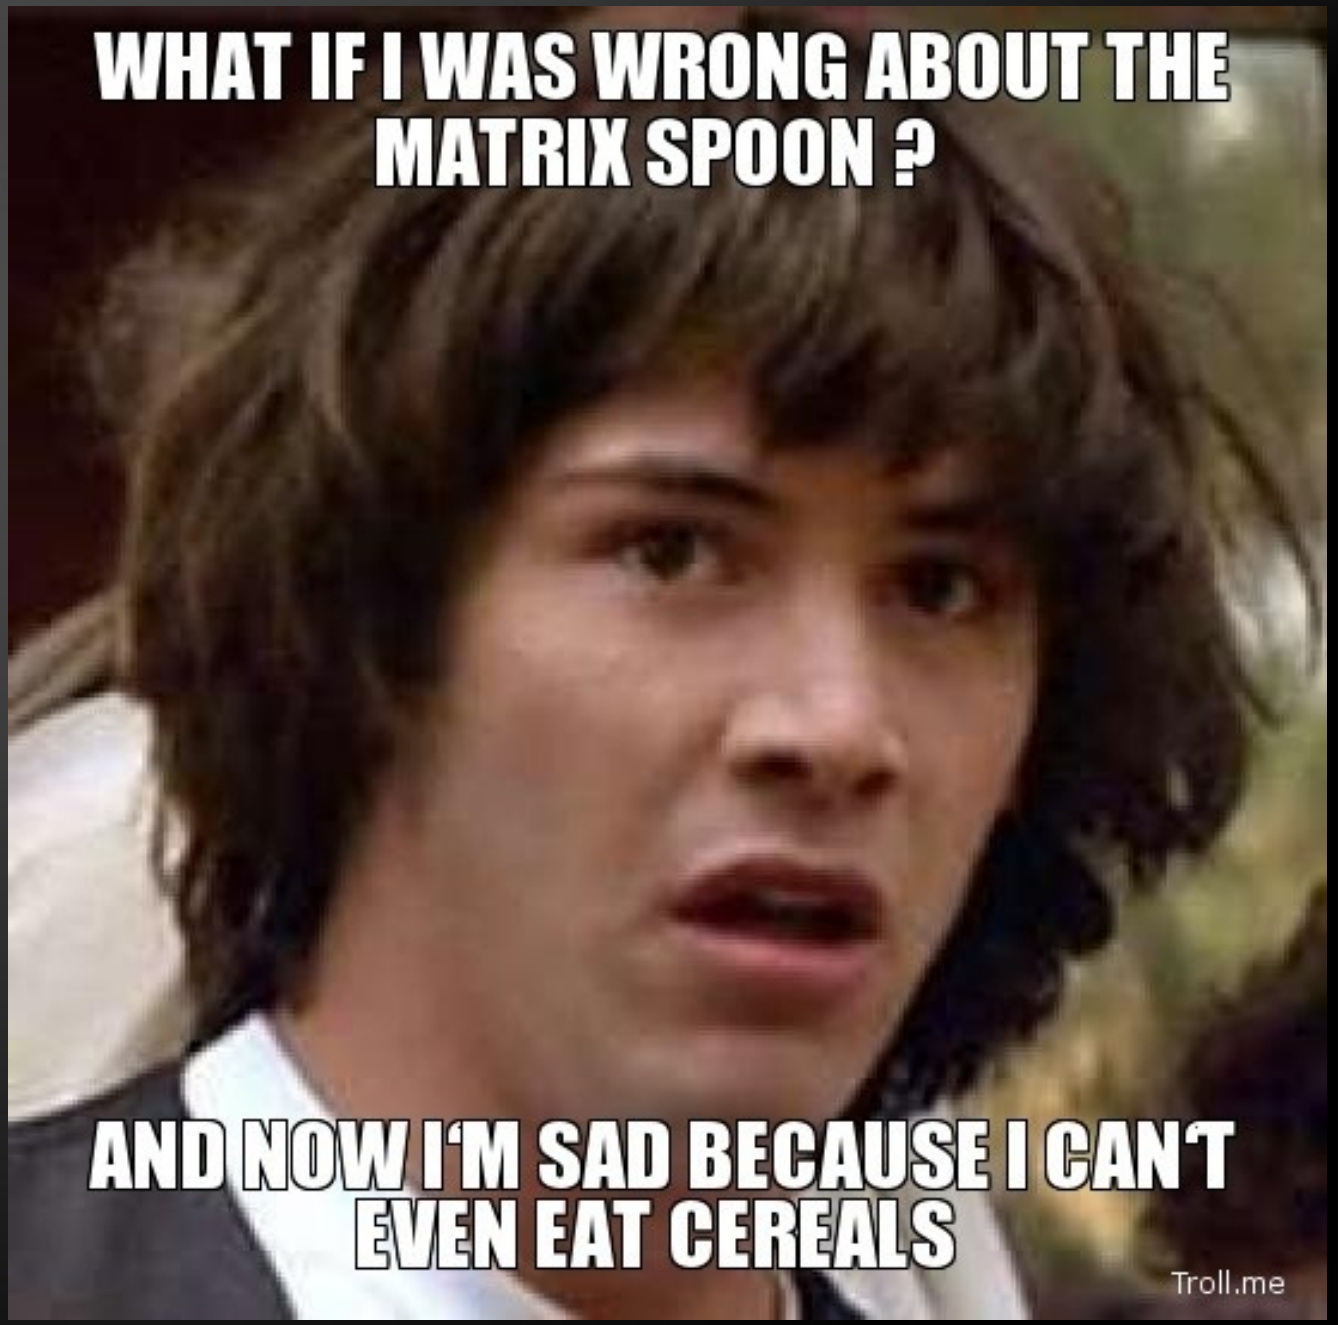

SOFTWARE UNIVERSITY

File System Architecture

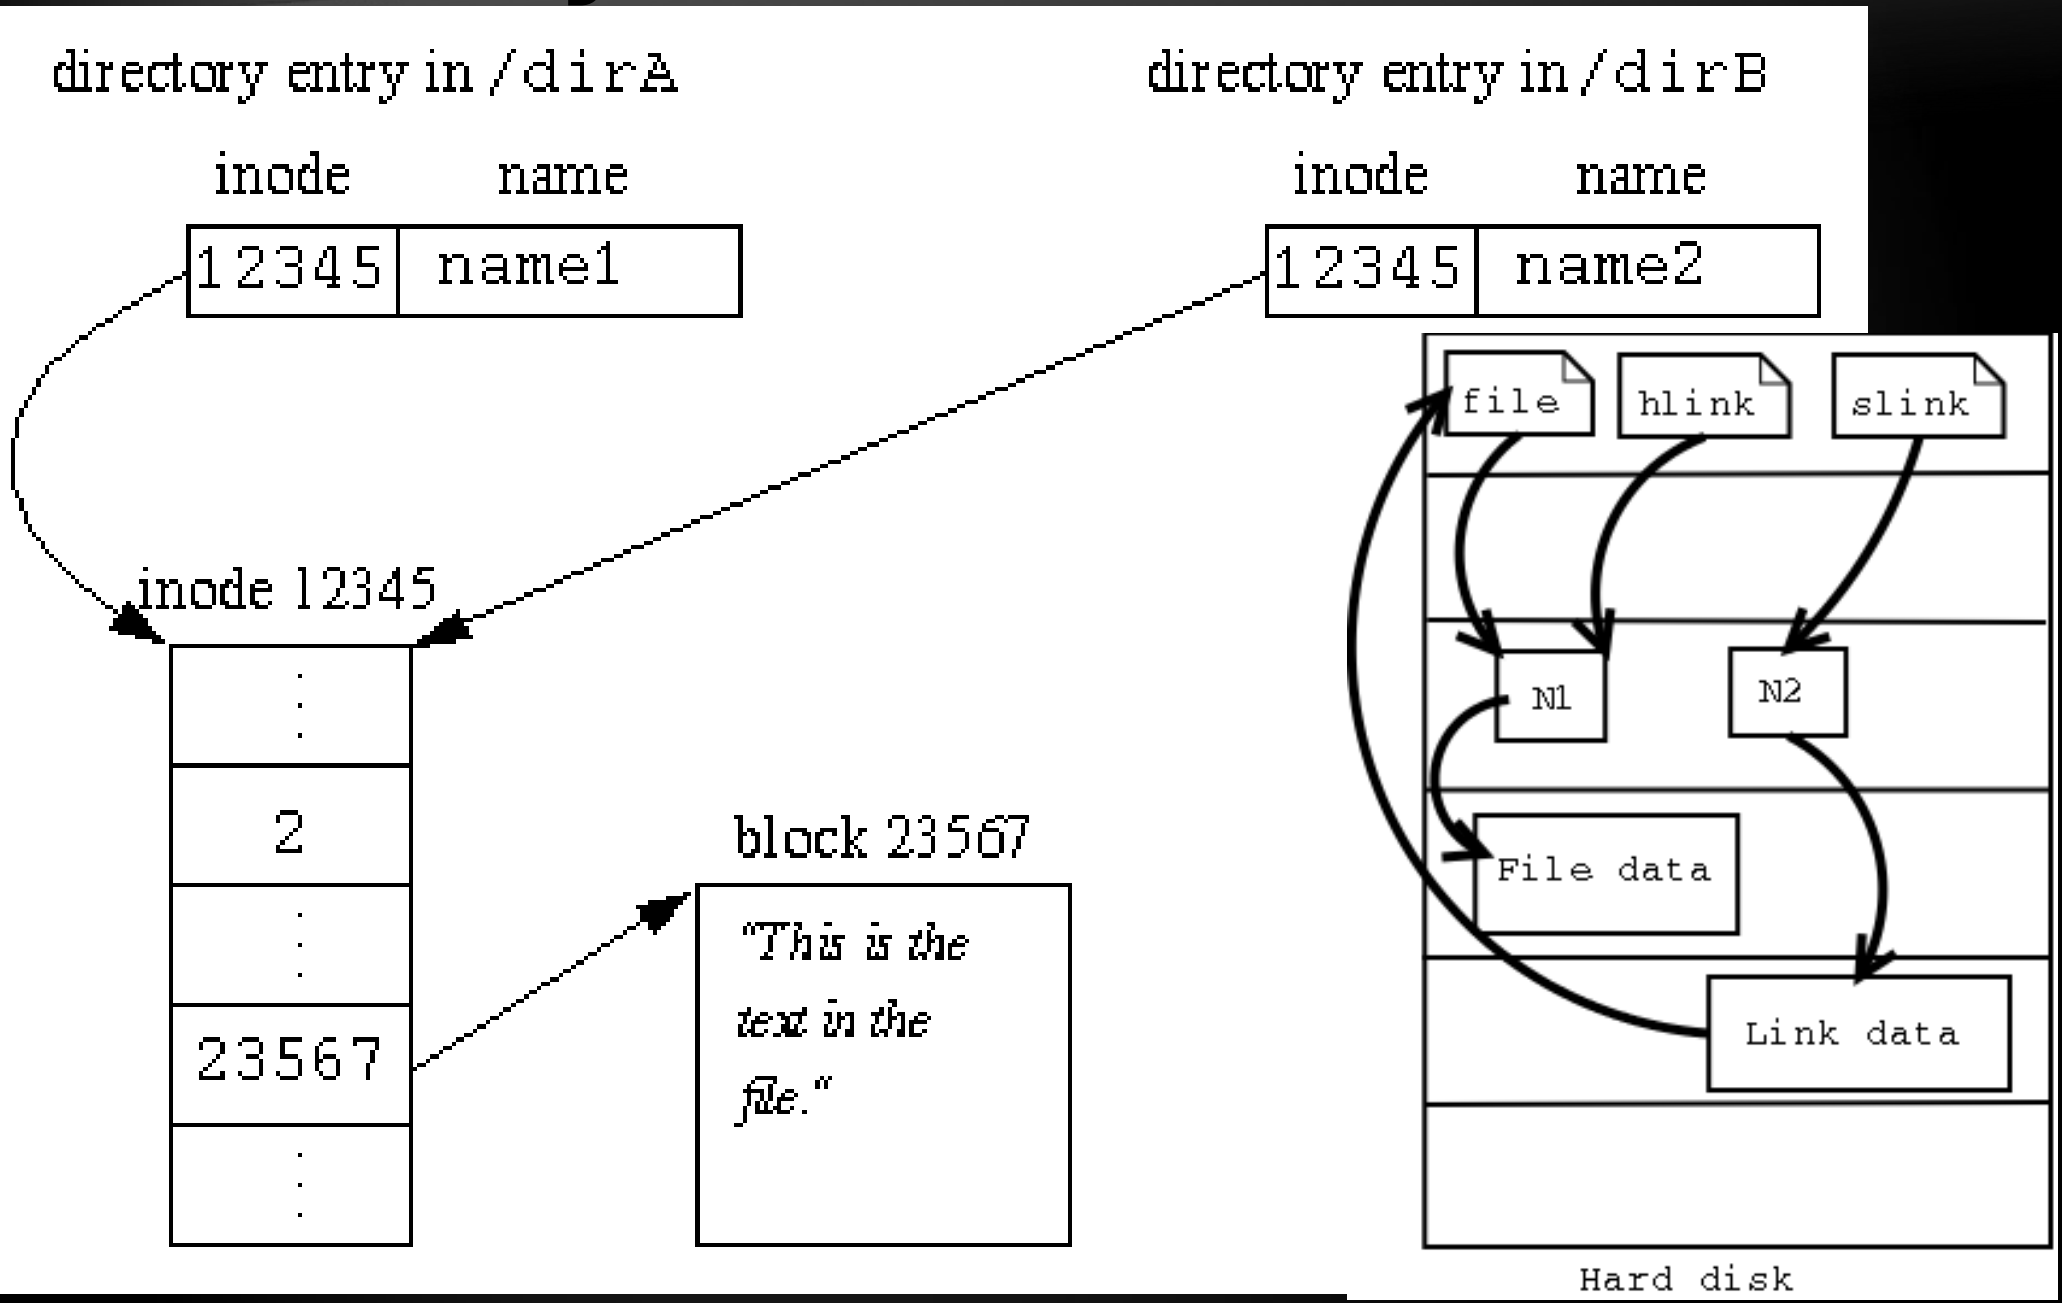

SOFTWARE UNIVERSITY

File System Architecture

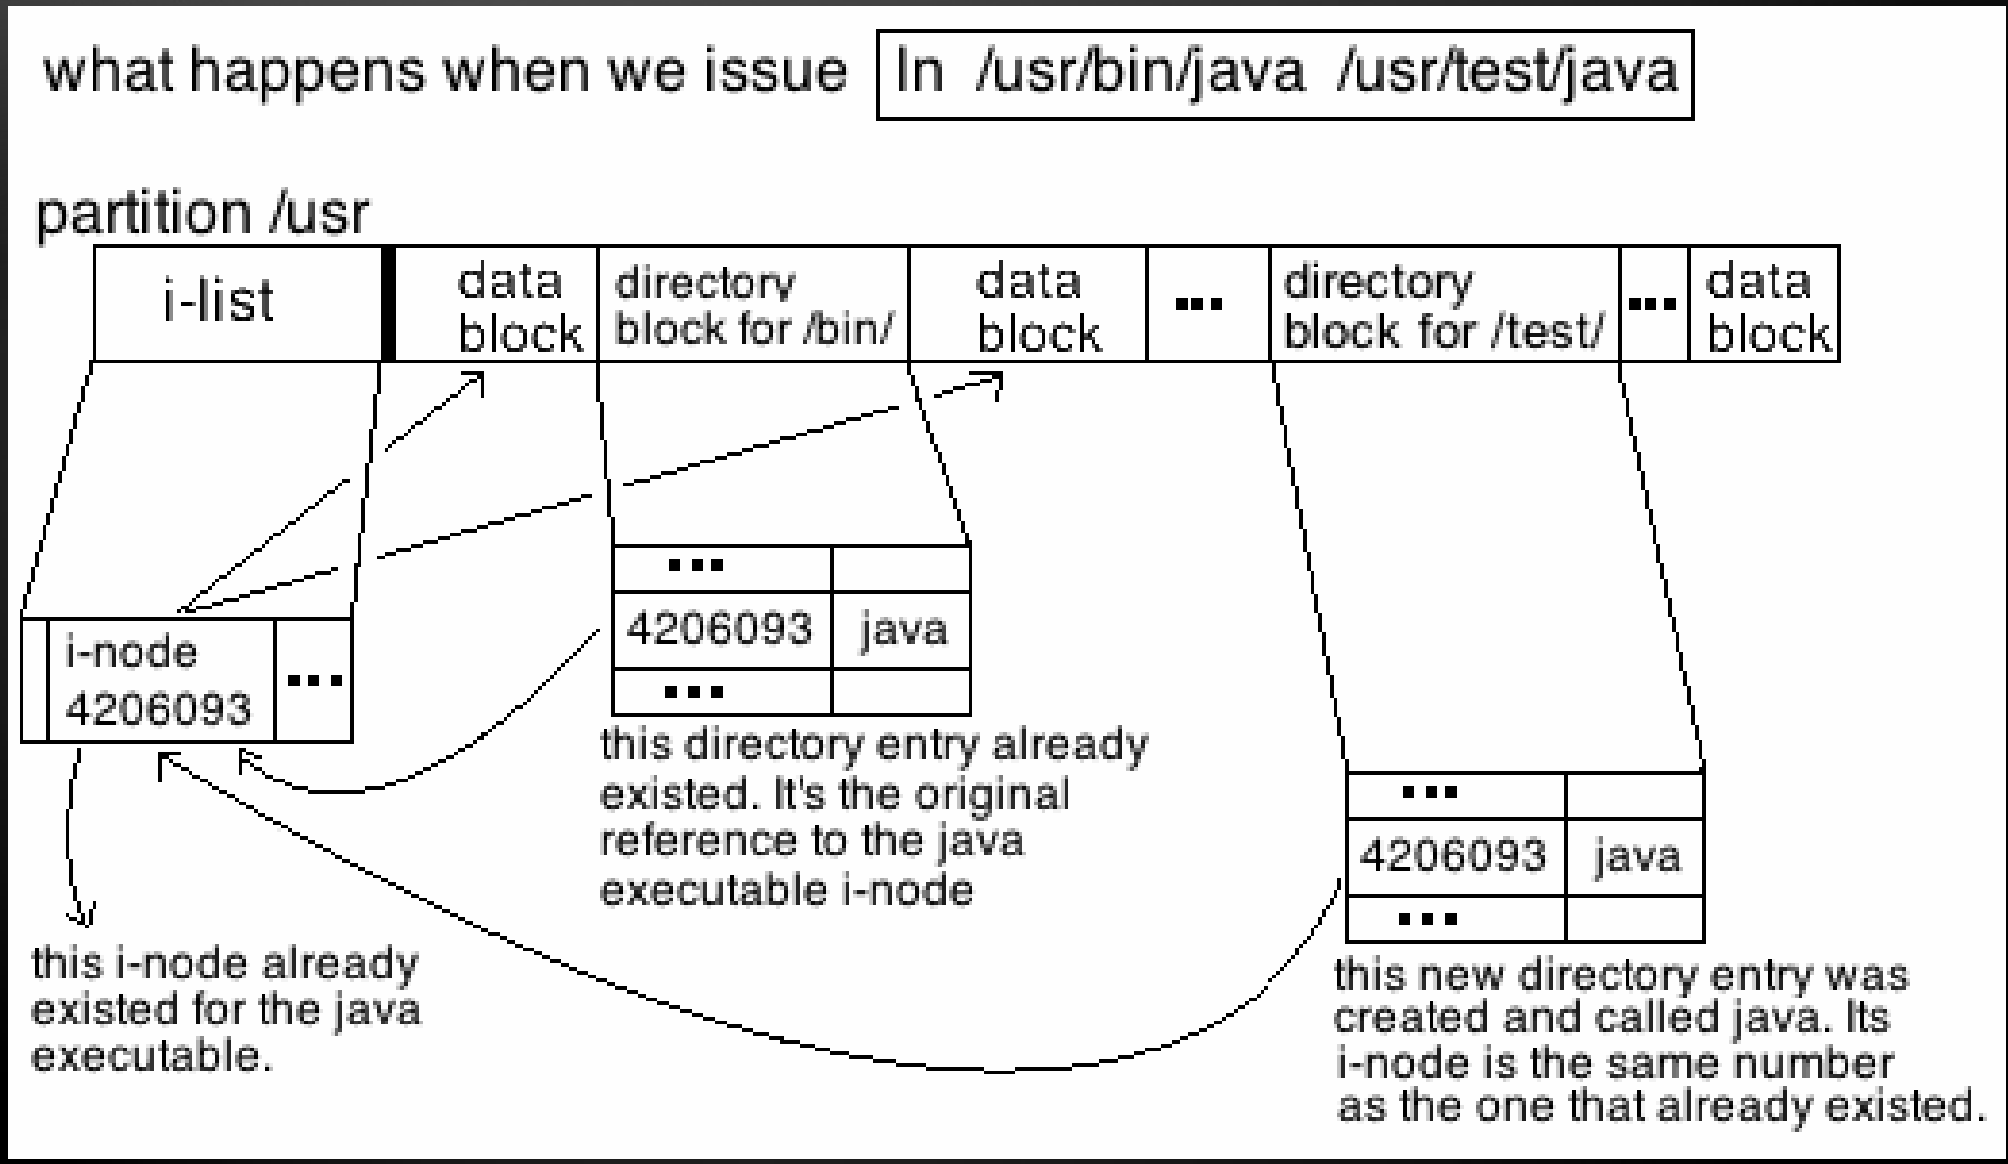

ln source destination **ln -s source destination**

SOFTWARE UNIVERSITY

File System Architecture

DATA is never erased... it gets OVERWRITTEN!

SOFTWARE UNIVERSITY

Local File System

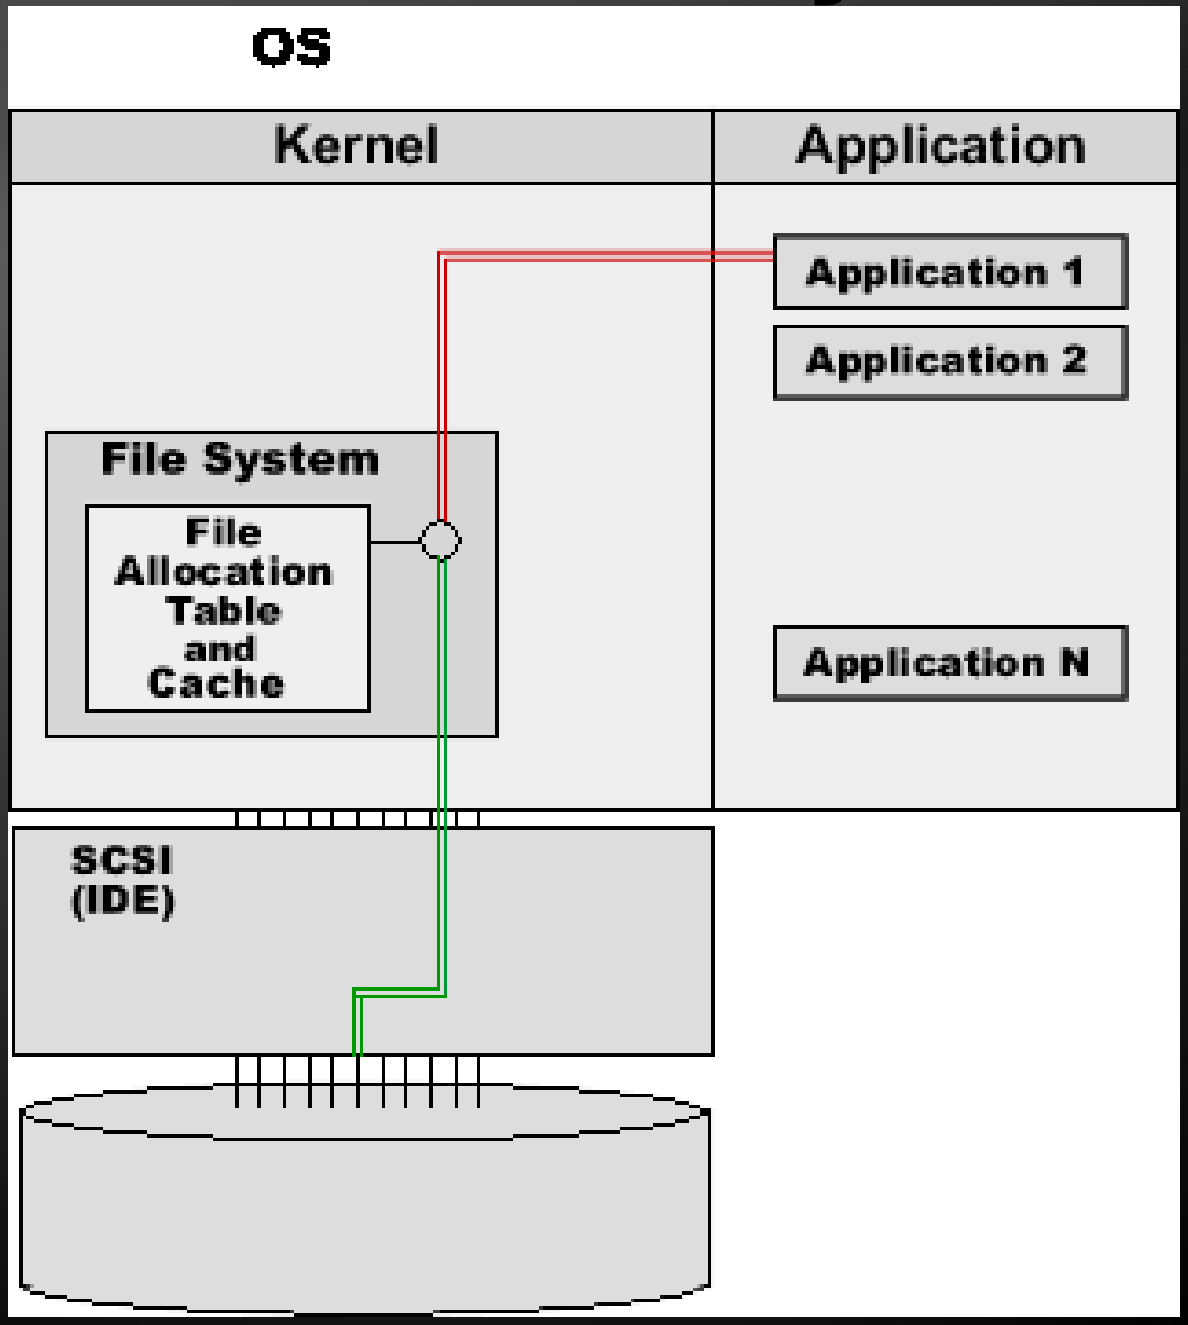

SOFTWARE UNIVERSITY

Network File System

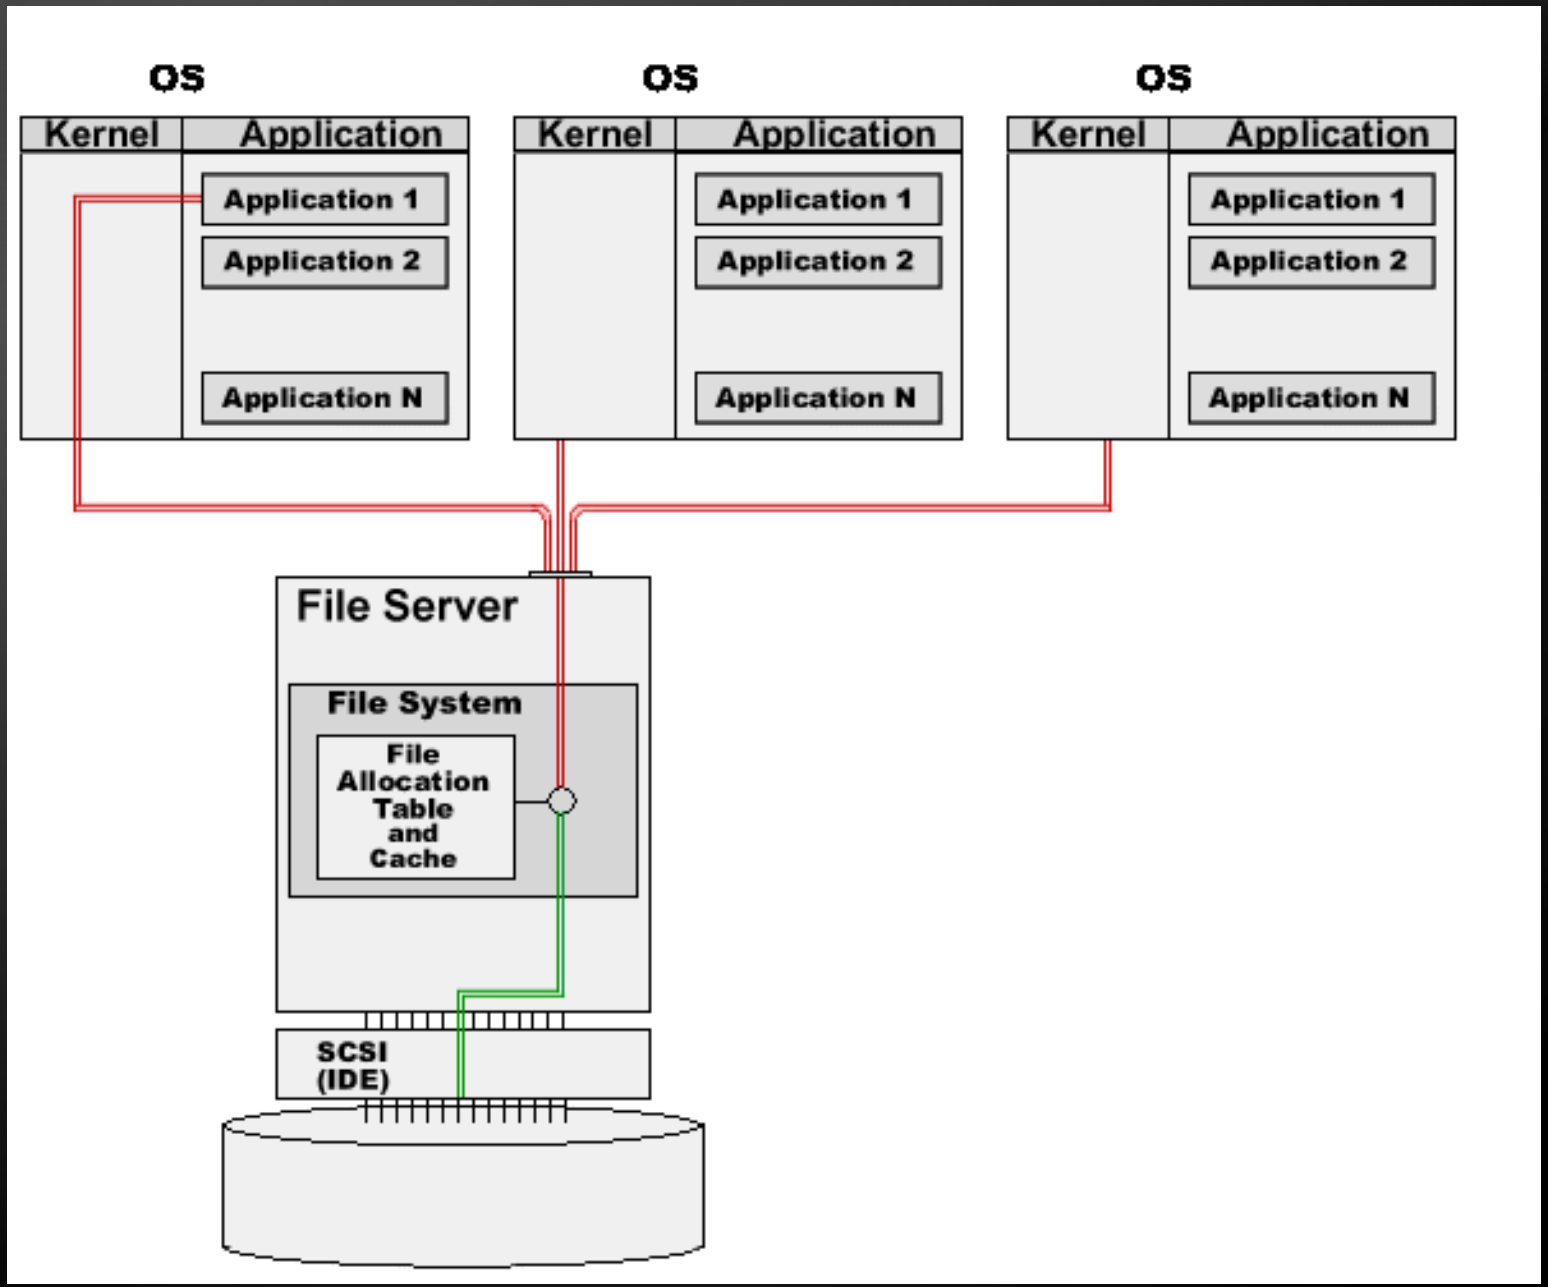

SOFTWARE UNIVERSITY

Cluster File System

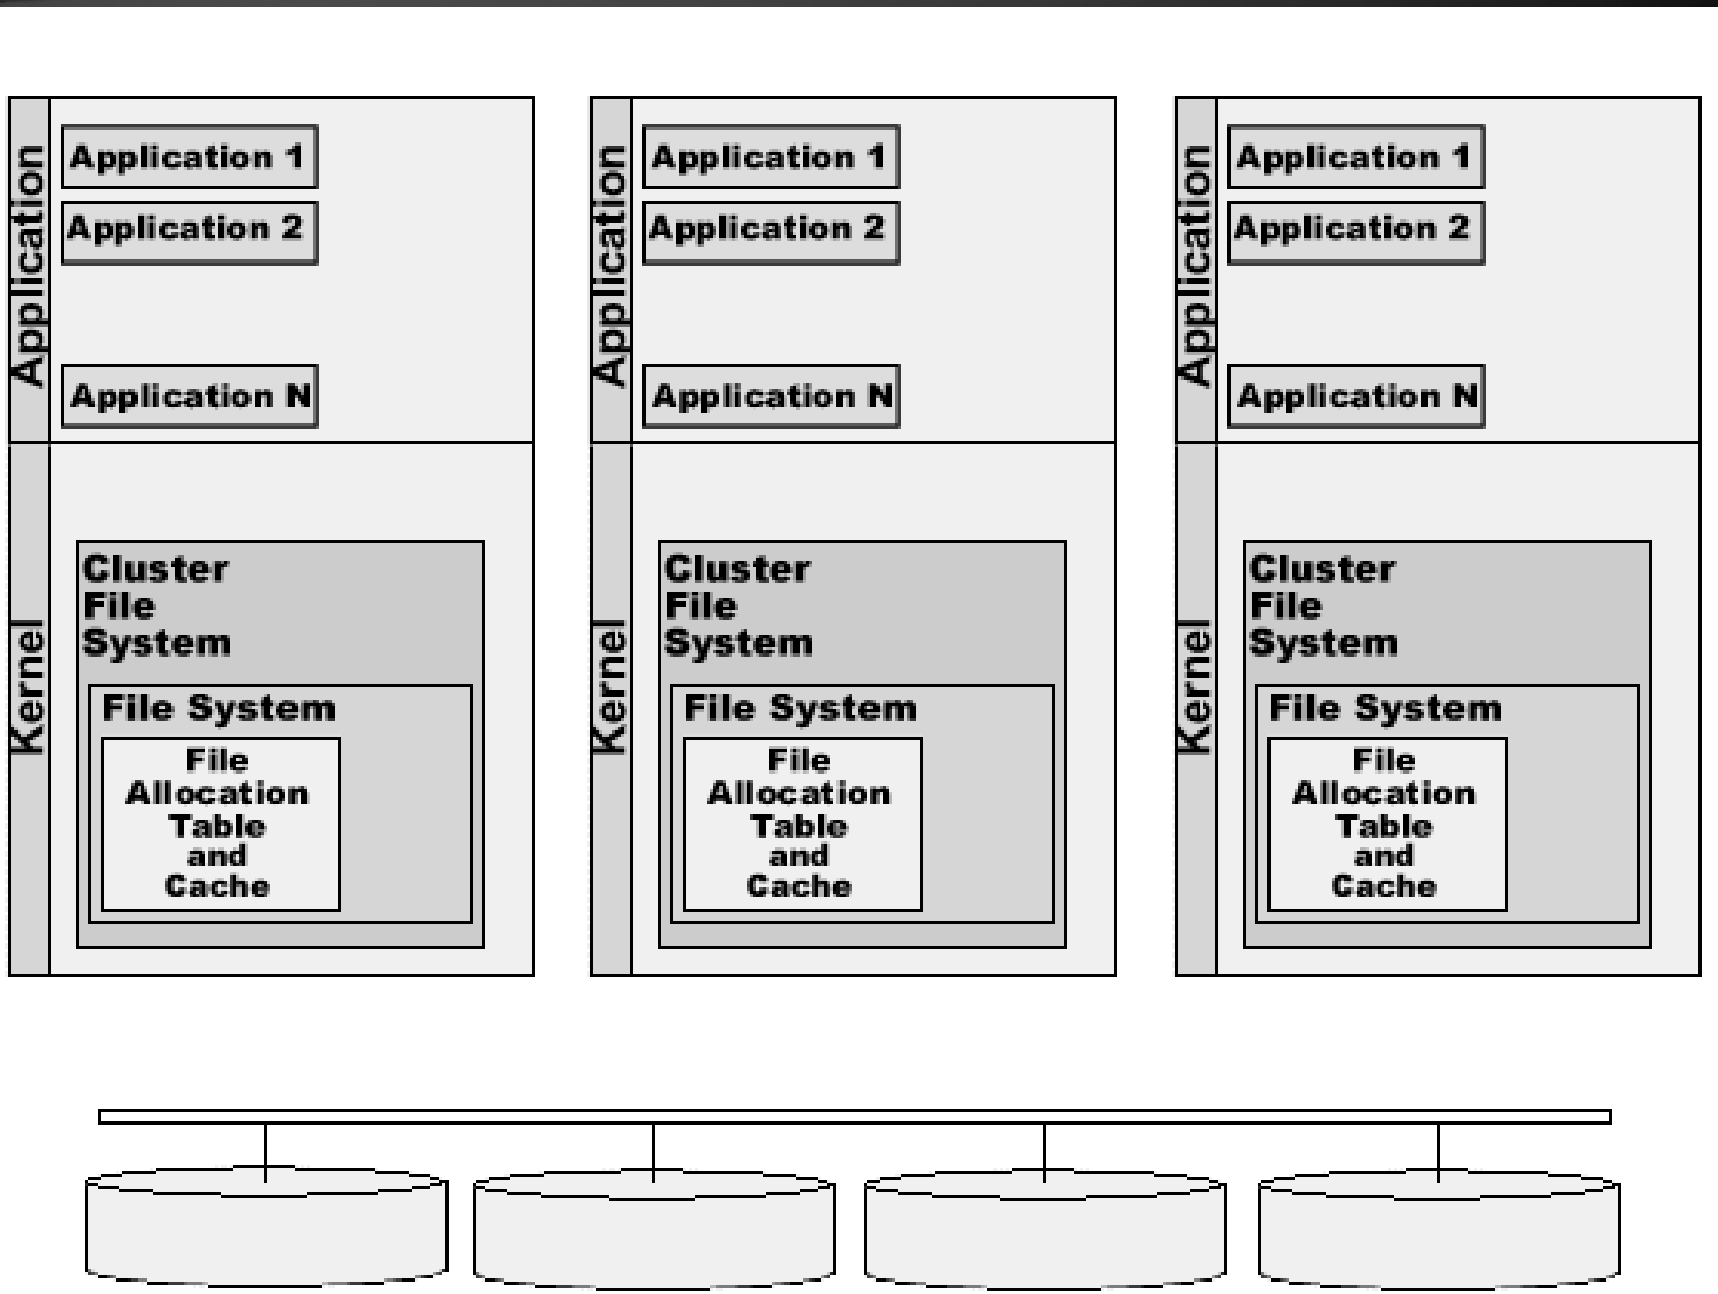

SOFTWARE UNIVERSITY

Cluster File System

“One DLM to rule them all, one DLM to mind them, one DLM to sync them all, and in the cluster, bind them”

SOFTWARE UNIVERSITY

Distributed File System

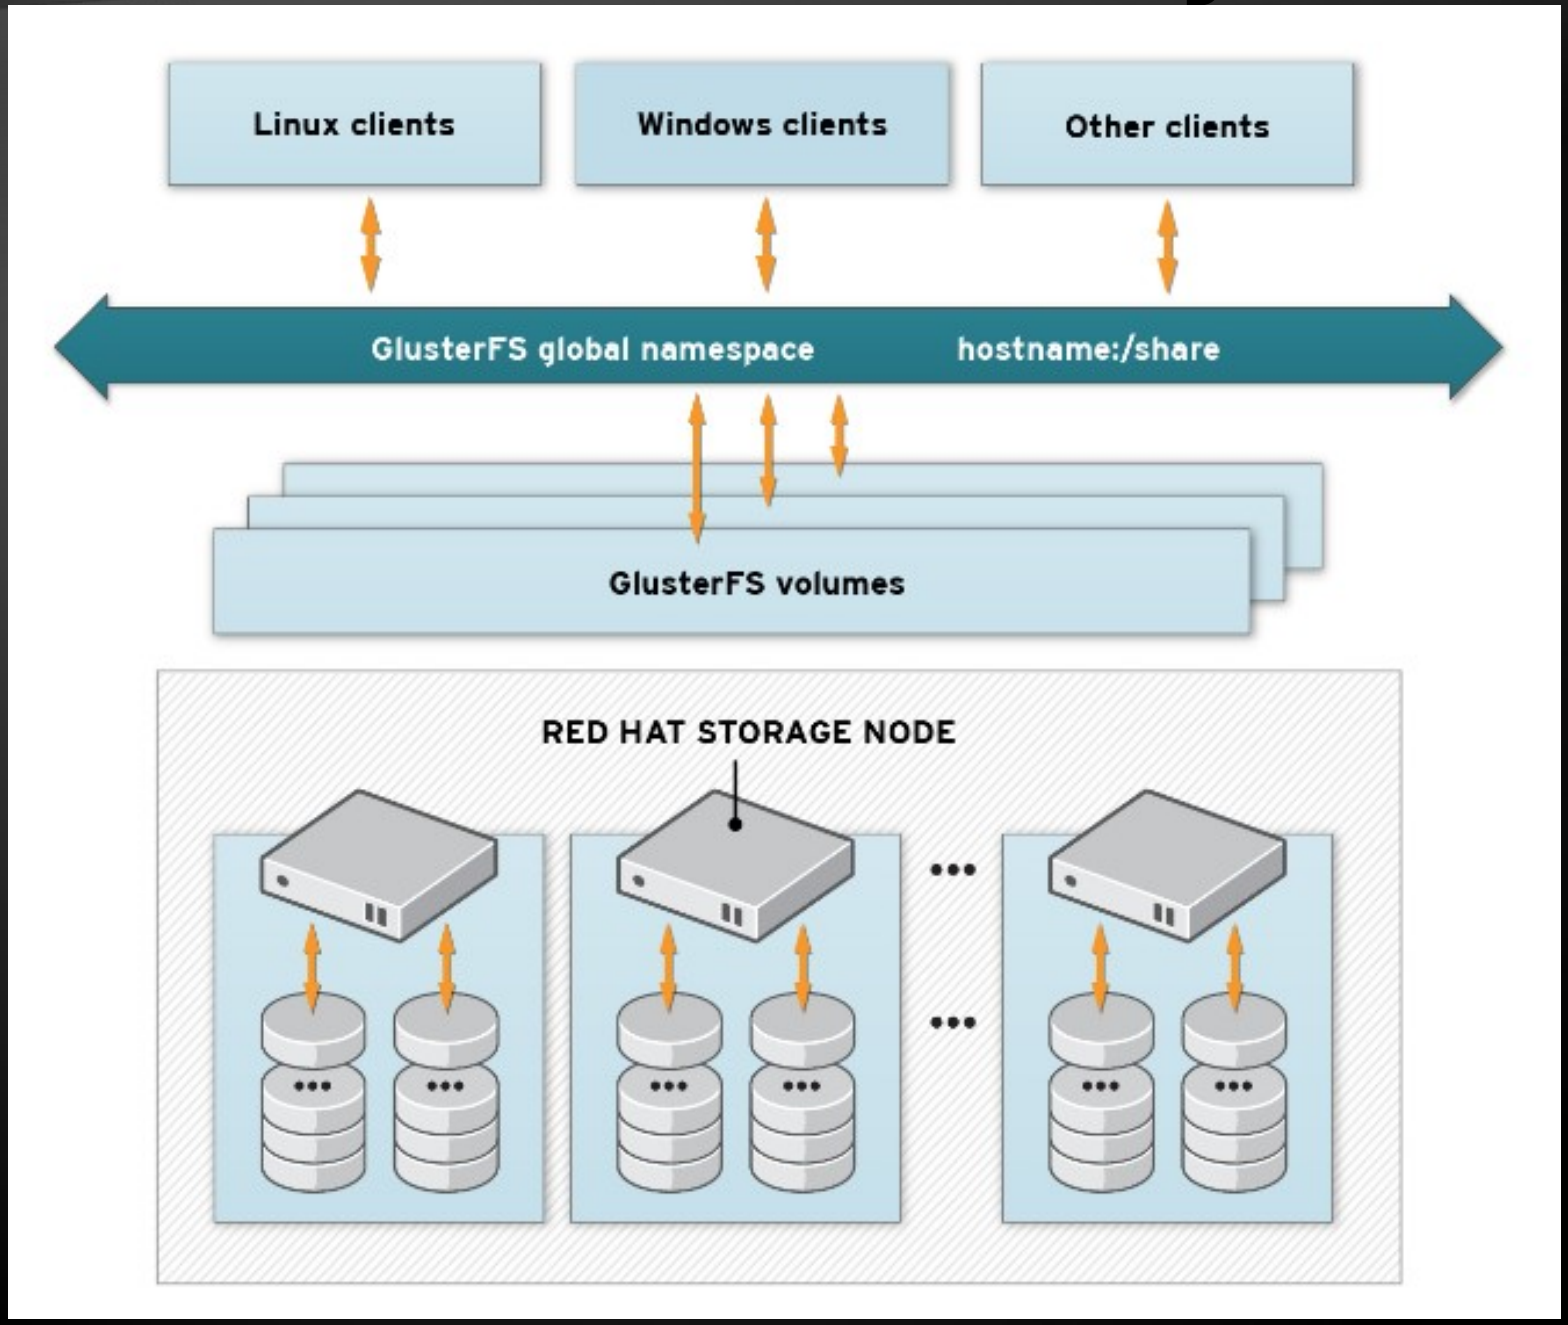

SOFTWARE UNIVERSITY

Virtual File System Layer

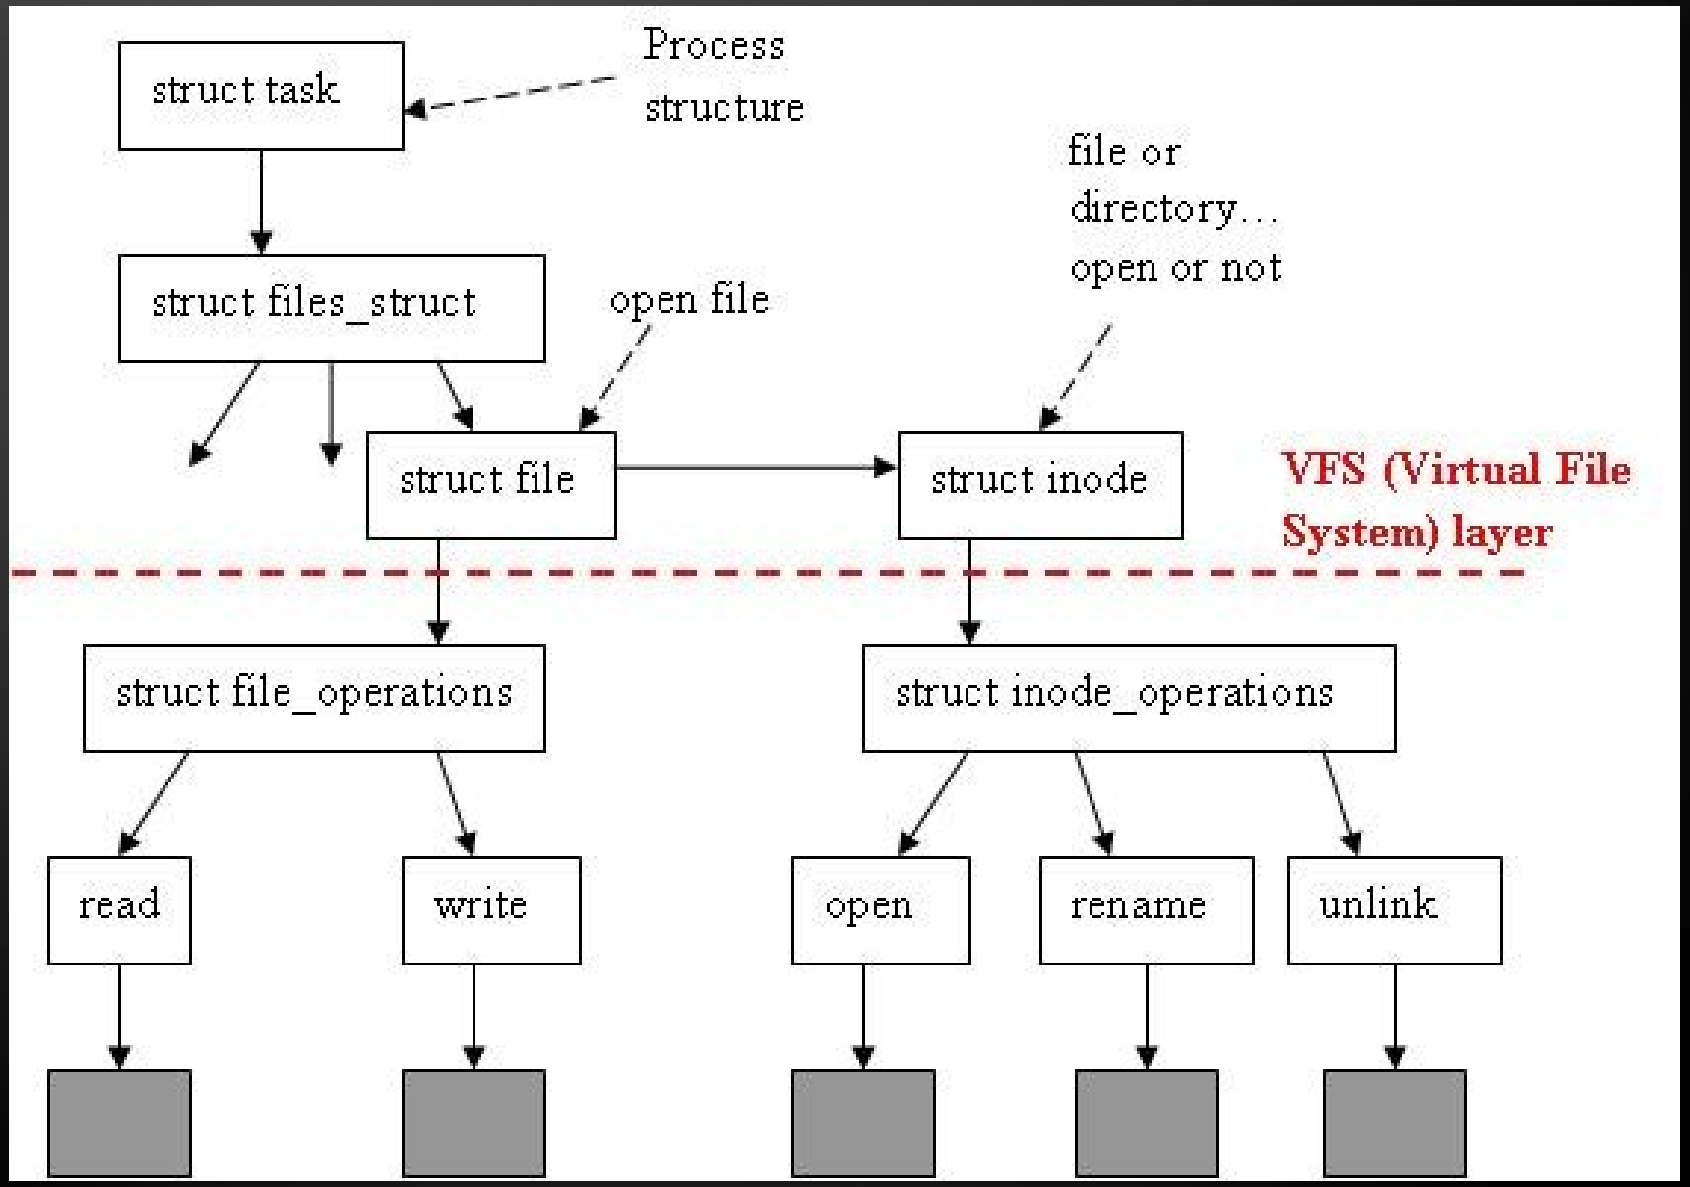

SOFTWARE UNIVERSITY

Virtual File System Layer

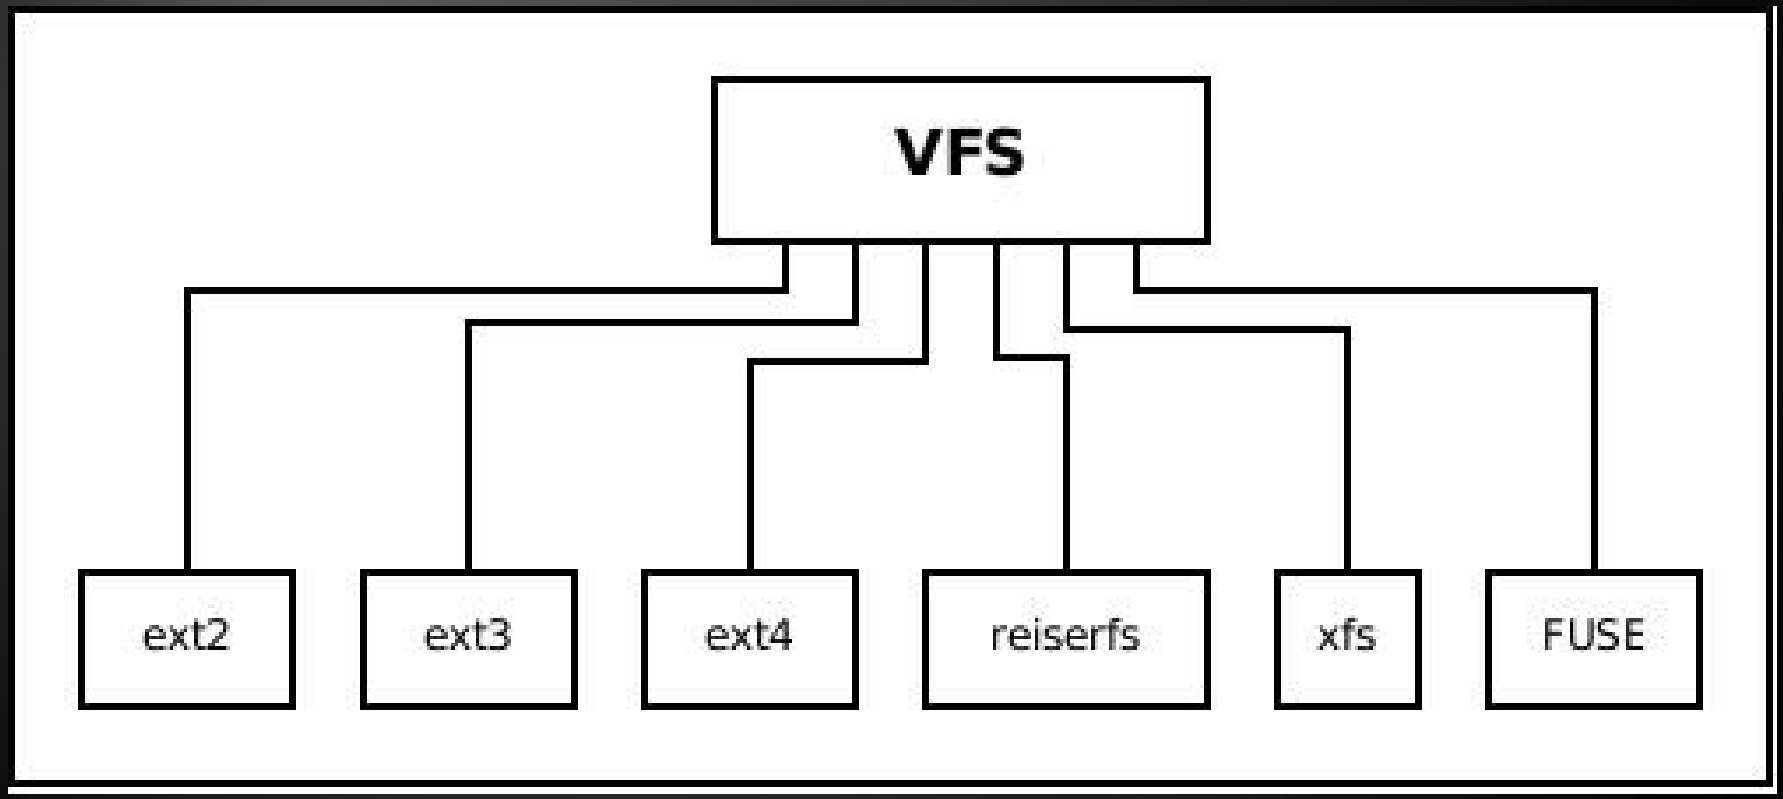

Introduced April 1992

SOFTWARE UNIVERSITY

VFS Caches

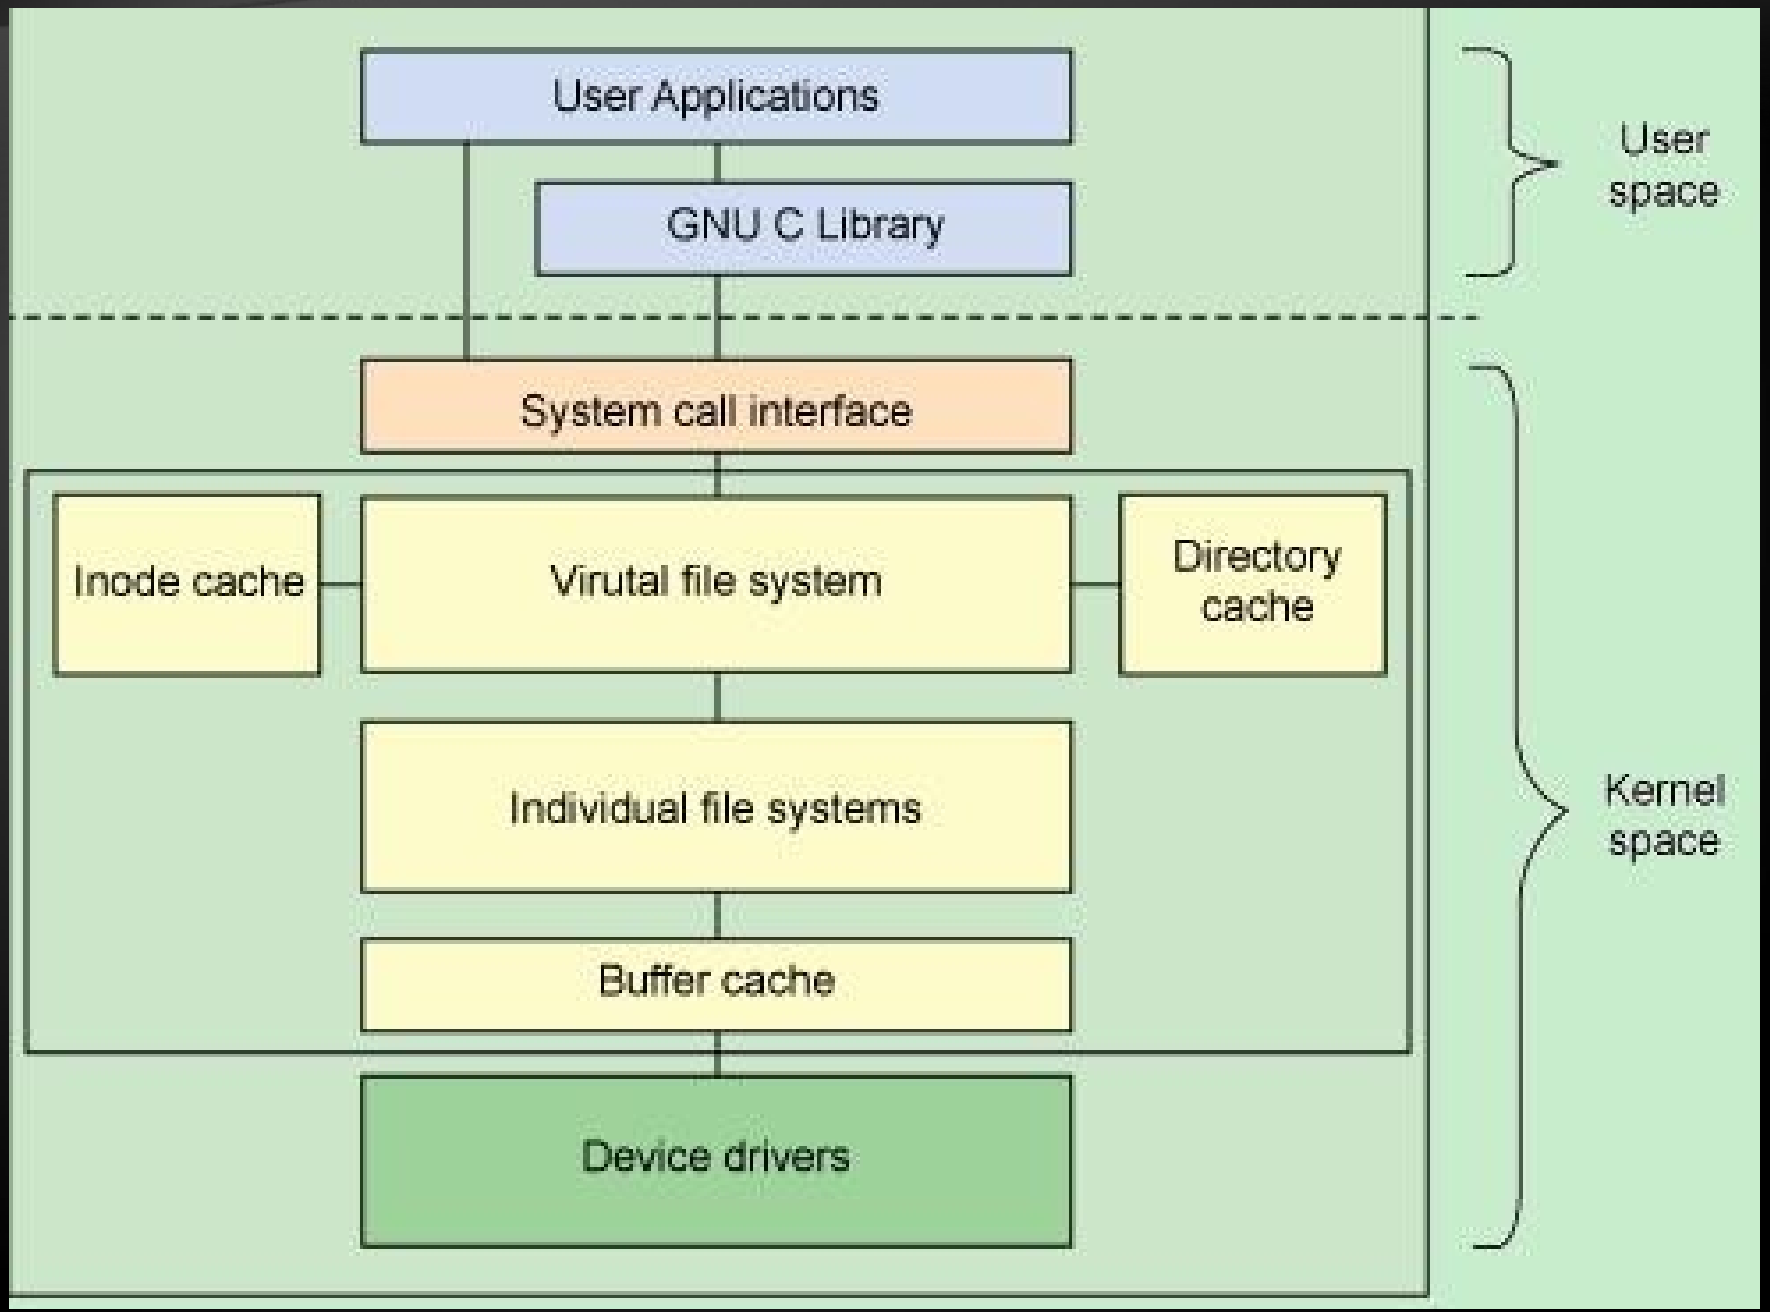

SOFTWARE UNIVERSITY

FUSE

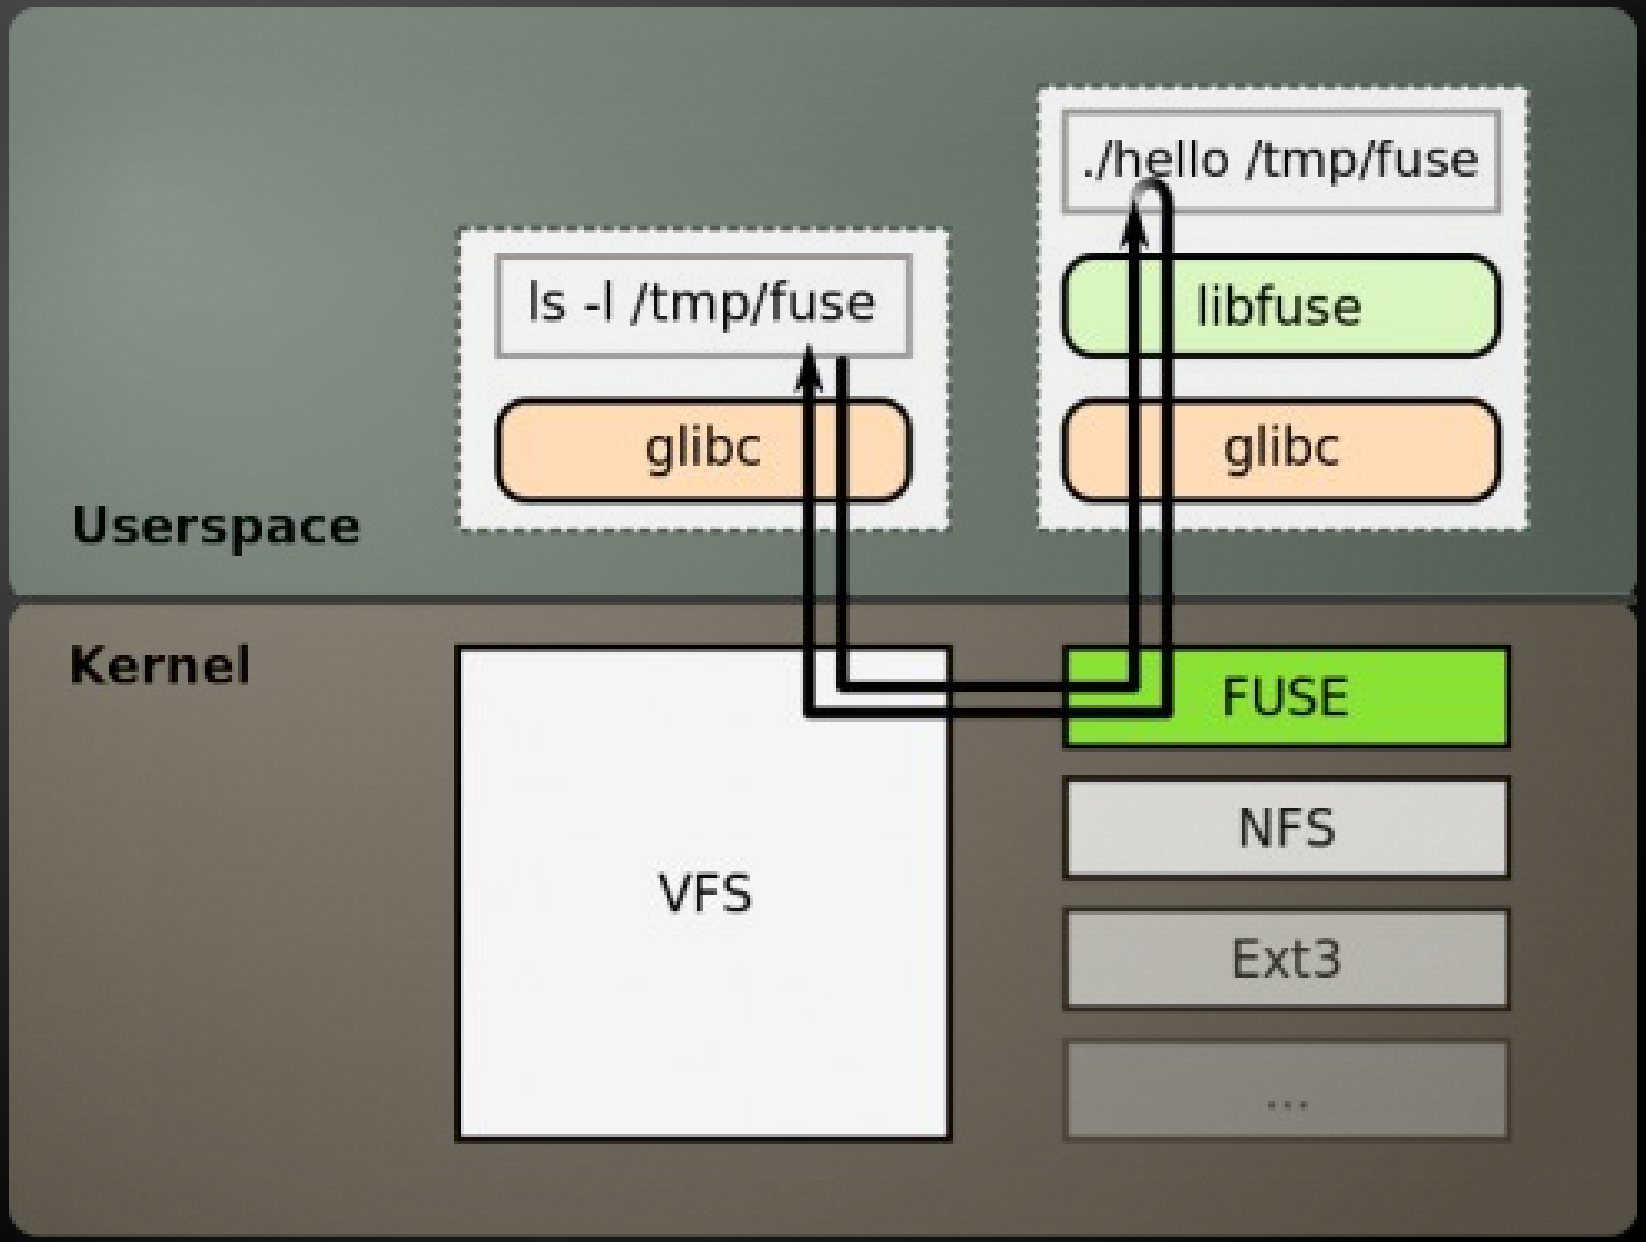

SOFTWARE UNIVERSITY

Mounting

- Attaching a device into the directory tree
- Mount point – a destination directory where a device is mounted
- Creates an entry in the kernel for each mounted device/dir
- /proc/mounts
  /etc/fstab, /etc/mtab

SOFTWARE UNIVERSITY

Mounting - CMD

- cat /proc/partitions
- cat /proc/mounts
- mount
- umount
- /etc/fstab

```
/dev/sdb2       /          ext4    defaults,noatime,nodiratime     0 0
/dev/sdb1       /boot      ext2    defaults,noatime,nodiratime     0 0
proc            /proc      proc    defaults                        0 0
tmpfs           /dev/shm   tmpfs   defaults                        0 0
/home/hackman   /fedora/home/hackman   none   rw,bind,auto   0 0
//10.2.0.11/share      /storage/beast        cifs
user=hackman,password=p1r@tk3,uid=1000,gid=1000,noauto 0 0
```

SOFTWARE UNIVERSITY

Ext/2/3/4

- First Linux FS – MinixFS
- And there we go….
  - Ext - April 1992, Linux 0.96c
  - Ext2 – January 1993
  - Ext3 – November 2001
  - Ext4 – October 2006

SOFTWARE UNIVERSITY

MinixFS

- Max. partition size – 64MB
- Max. file name size – 14 chars
- Ownership – uid, gid
- Permissions – user, group, others

SOFTWARE UNIVERSITY

Ext

- Max. partition size – 2GB
- Max. file name size – 255 chars
- No support for time stamps
  - Access
  - Inode modification
  - Data modification

SOFTWARE UNIVERSITY

Ext2

- Max. partition size – 32TB
- Max. file name size – 255 chars
- Max. file size – 2TB
- Max. Number of files - 10^{18}
- FS Perms
- Time stamps

SOFTWARE UNIVERSITY

Ext3

- Max. partition size – 32TB
- Max. file name size – 255 chars
- Max. file size – 2TB
- Max. Number of files - 10^{18}
- Sub directory limit - 32,000
- Time stamps
- Has Jurnal

SOFTWARE UNIVERSITY

Ext4

- Max. partition size – 1EB
- Max. file name size – 255 chars
- Max. file size – 2TB
- Max. Number of files - 10^{18}
- Sub directory limit - 64,000
- File space pre-allocation
- File space delayed allocation

Ext/2/3

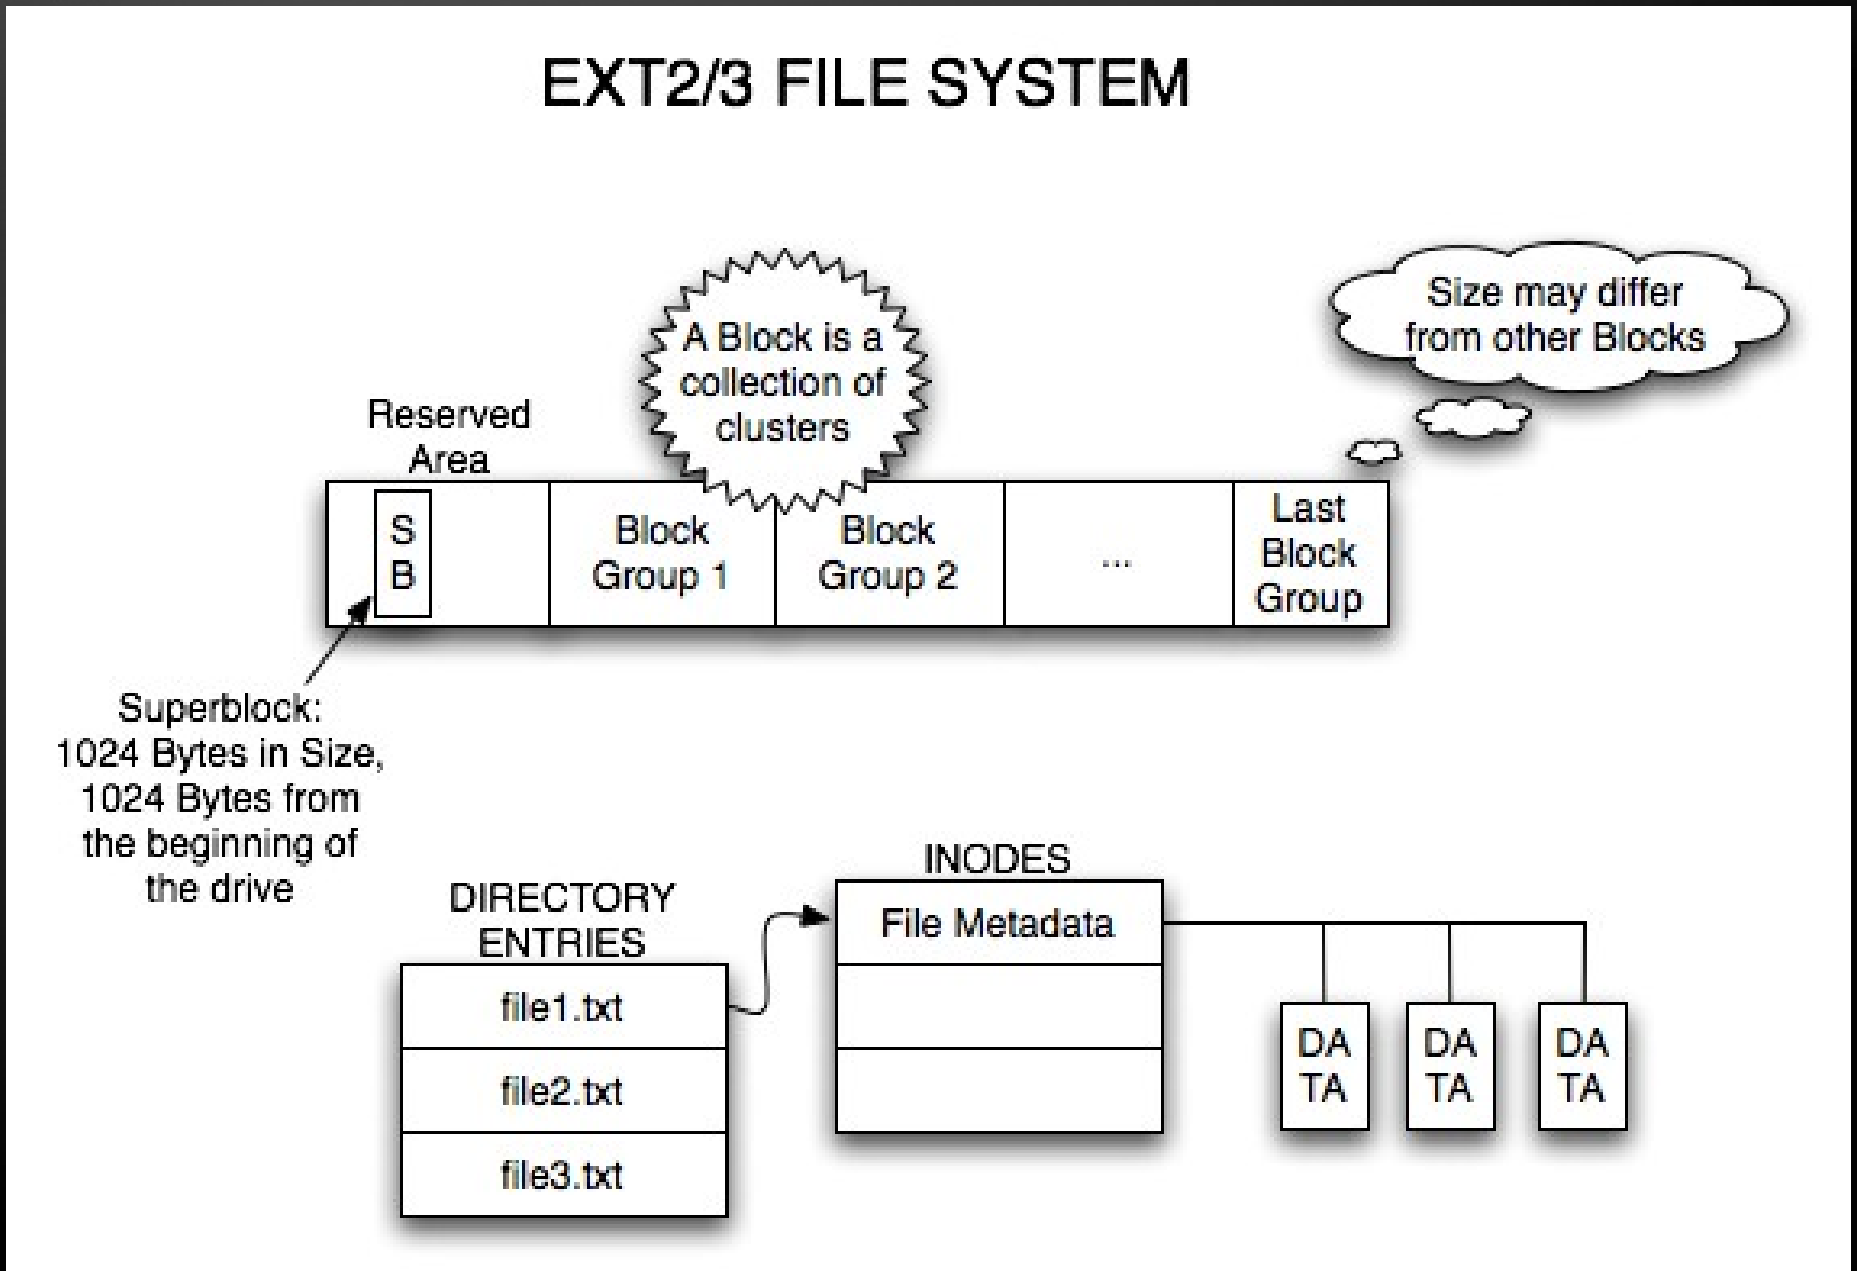
EXT2/3 FILE SYSTEM
A Block is a collection of clusters
Size may differ from other Blocks
Reserved Area
S B
Block Group 1
Block Group 2
...
Last Block Group
Superblock: 1024 Bytes in Size, 1024 Bytes from the beginning of the drive
INODES
DIRECTORY ENTRIES
file1.txt
file2.txt
file3.txt
File Metadata
DA TA
DA TA
DA TA

SOFTWARE UNIVERSITY

Ext3

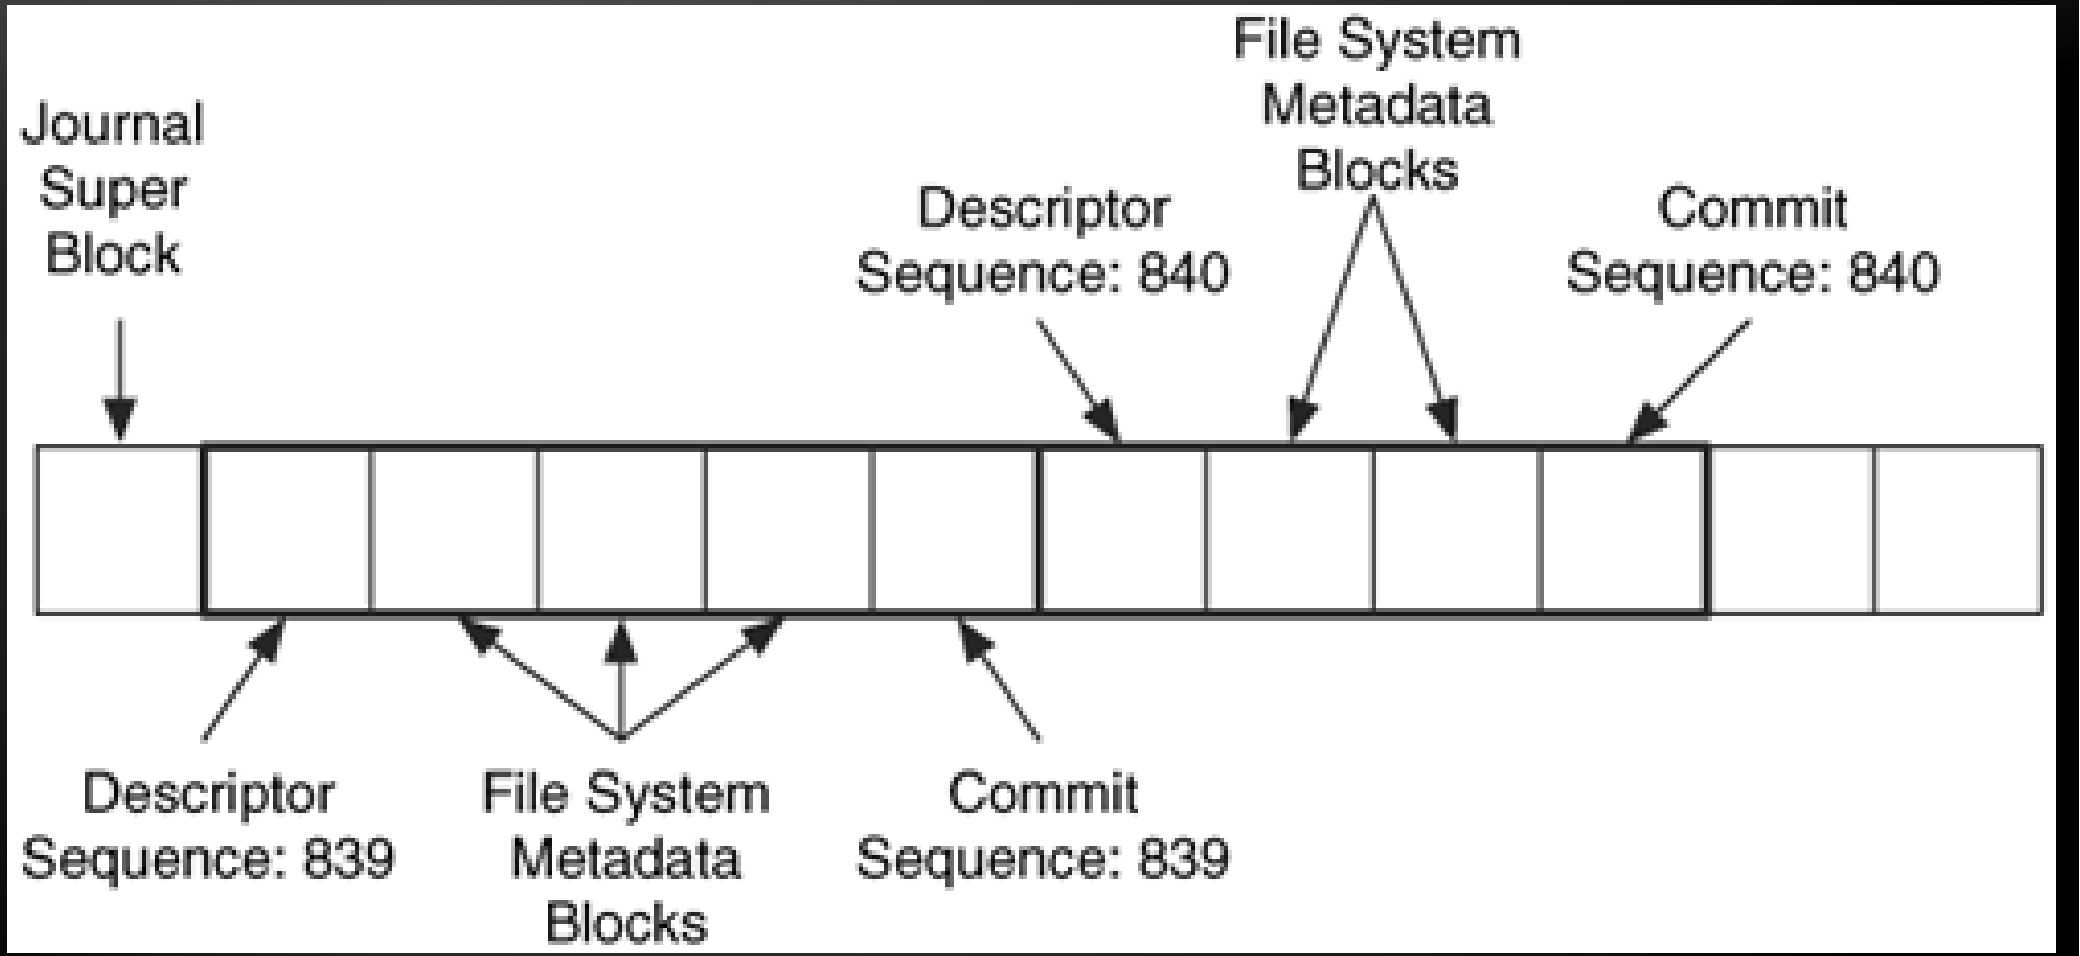

SOFTWARE UNIVERSITY

Ext4

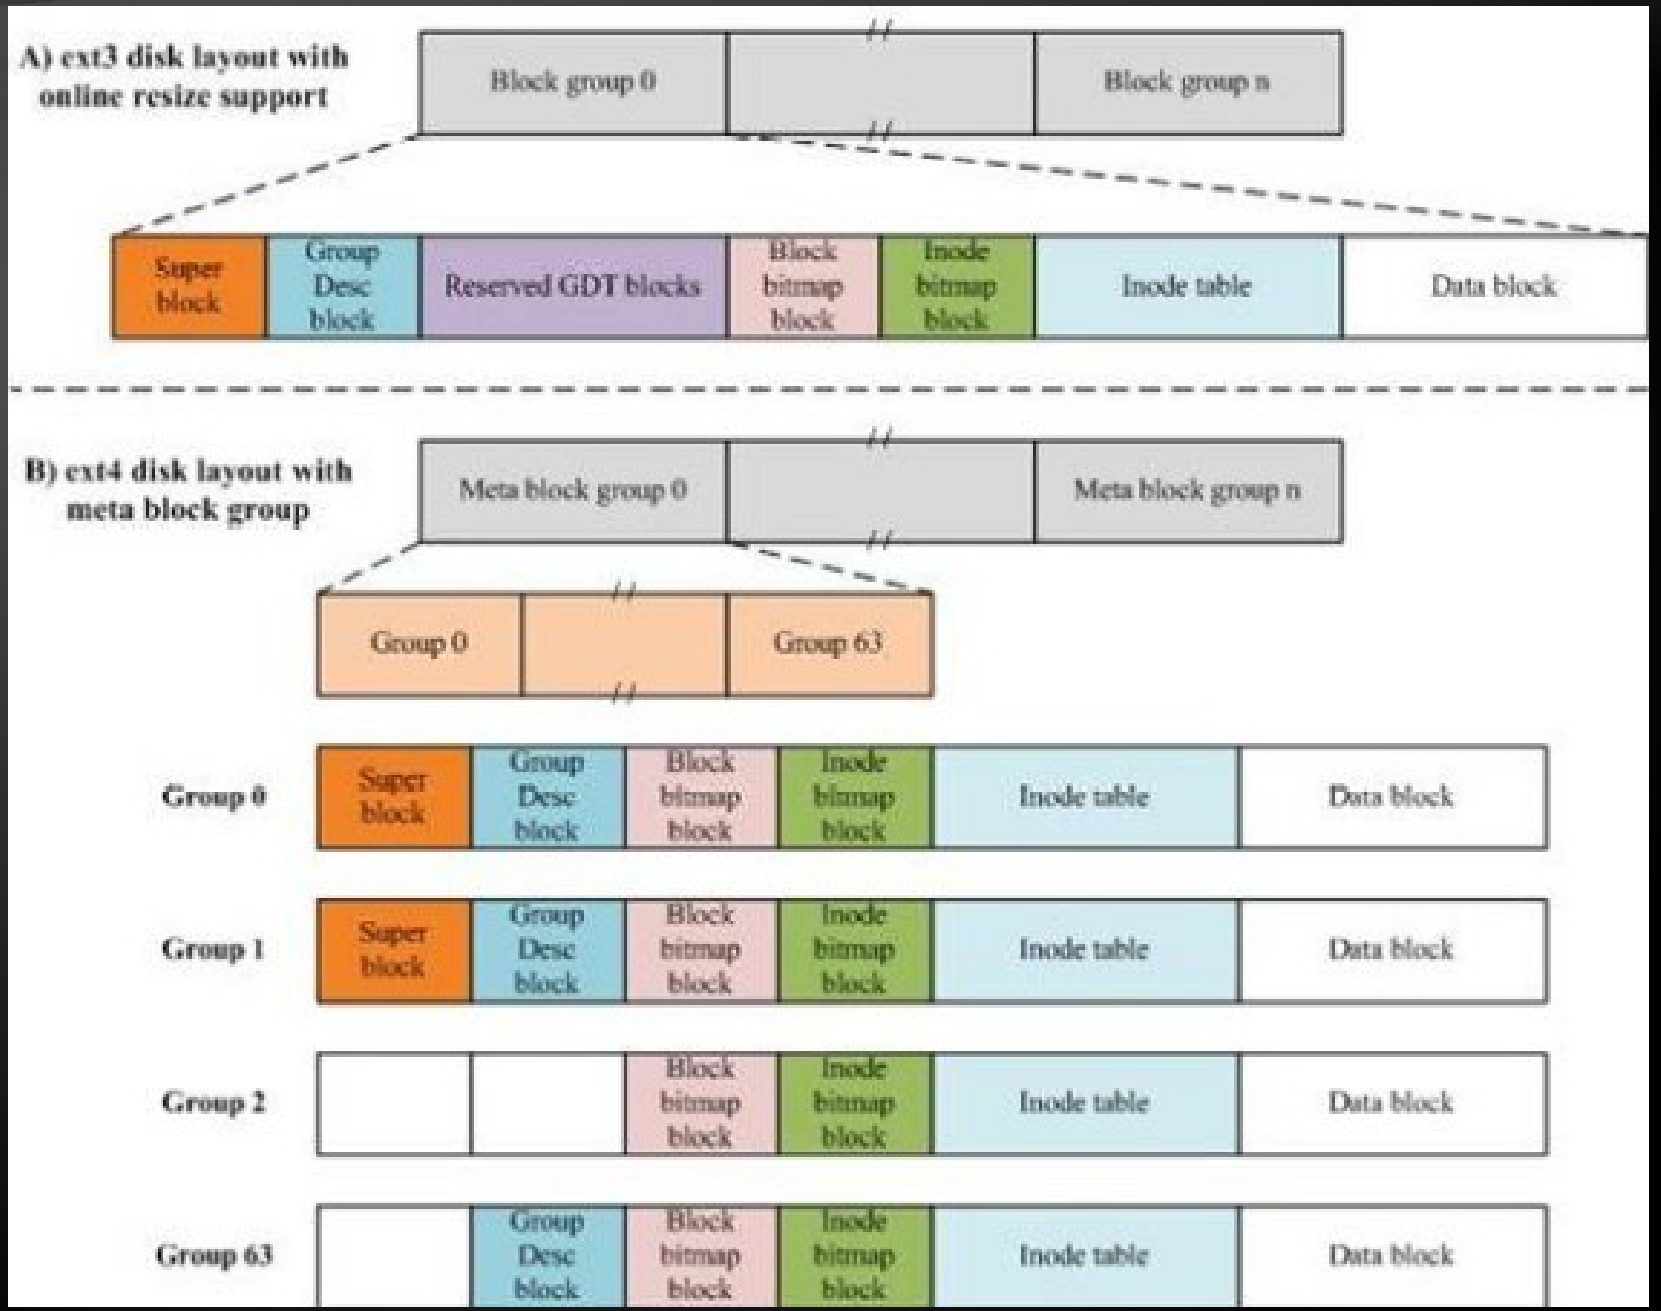

SOFTWARE UNIVERSITY

ReiserFS

- Introduced 2001
- Metadata-only journaling
- Online resizing (growth only)
- Tail packing, a scheme to reduce internal fragmentation.
- Max. file size - 1EB
- Max. number of files - 2^{32}

SOFTWARE UNIVERSITY

ReiserFS

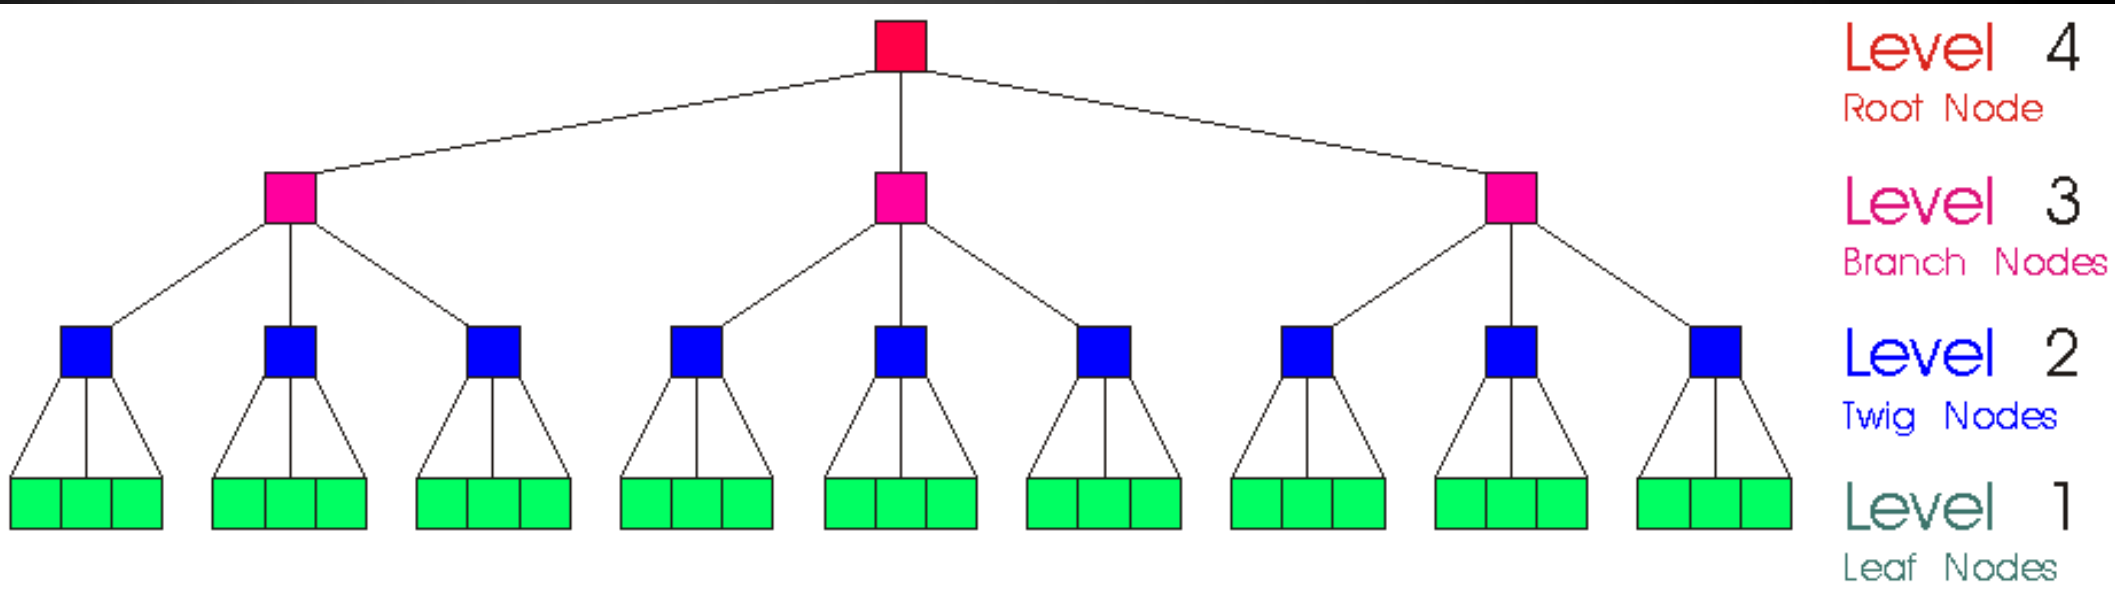

SOFTWARE UNIVERSITY

ReiserFS

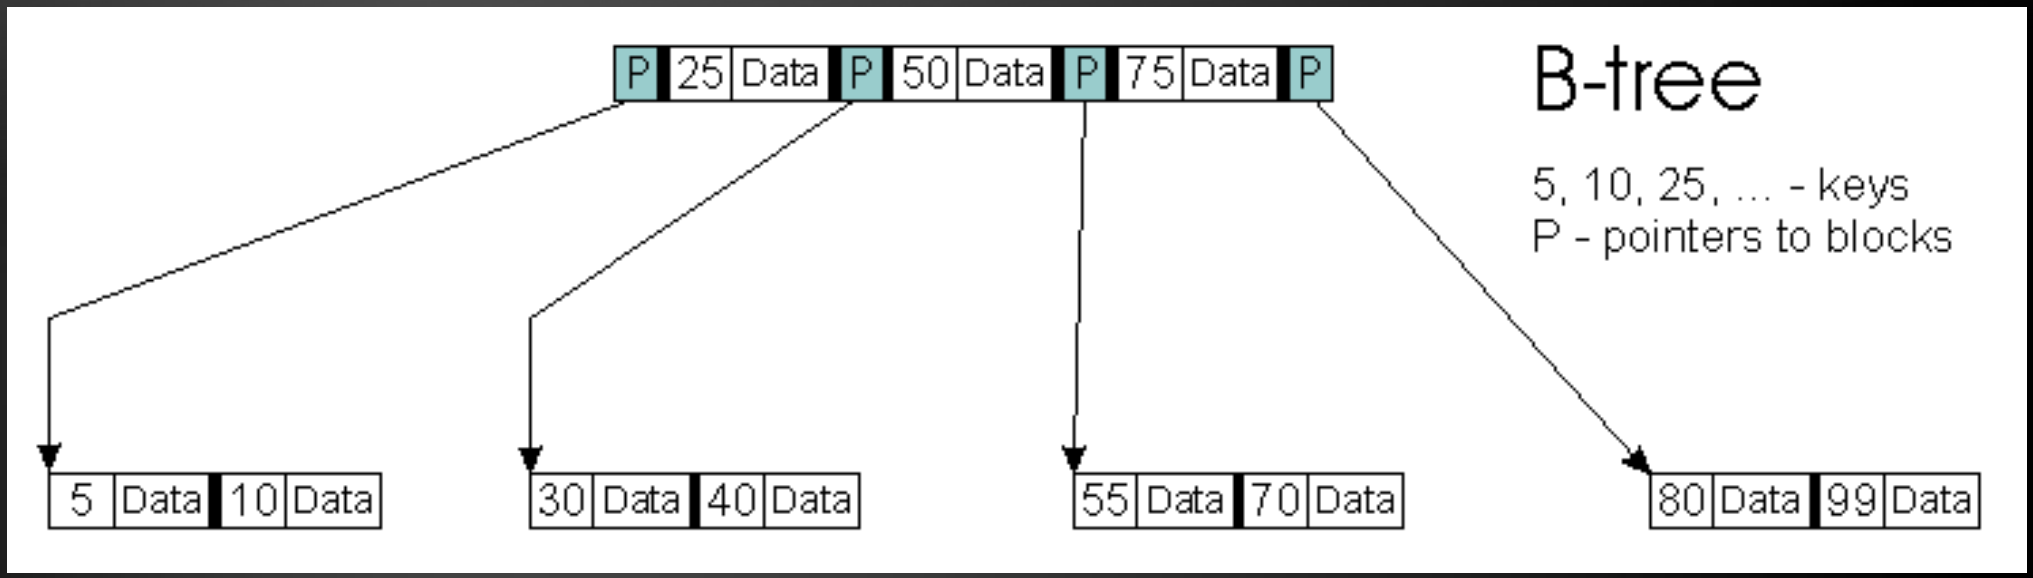

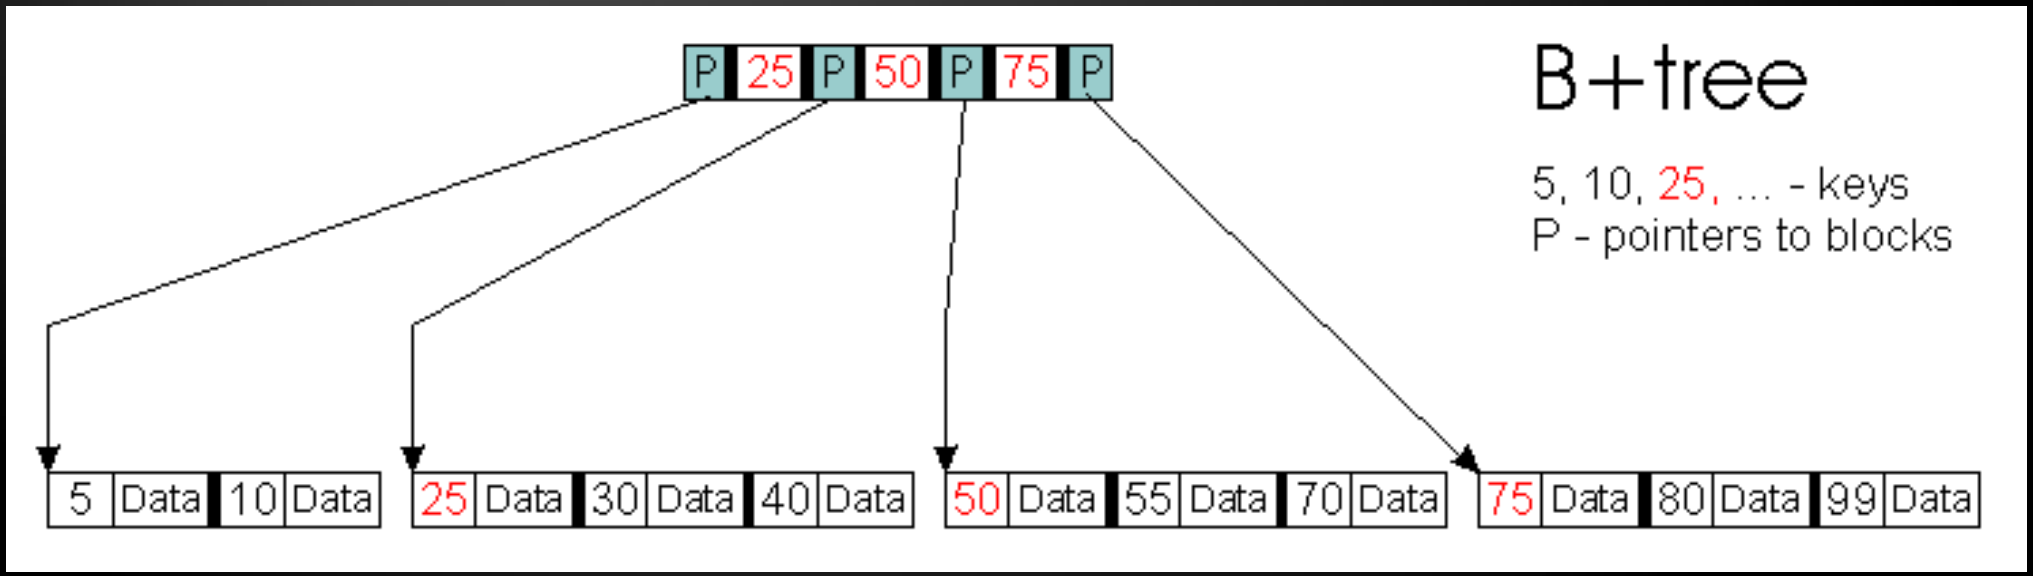

SOFTWARE UNIVERSITY

XFS

- Introduced 2001
- Max file size 8 EB
- Max volume size 16 EB
- Online resize(growth only)
- Online defragmentation

SOFTWARE UNIVERSITY

XFS

Equally sized chunks

Allocation groups - AG

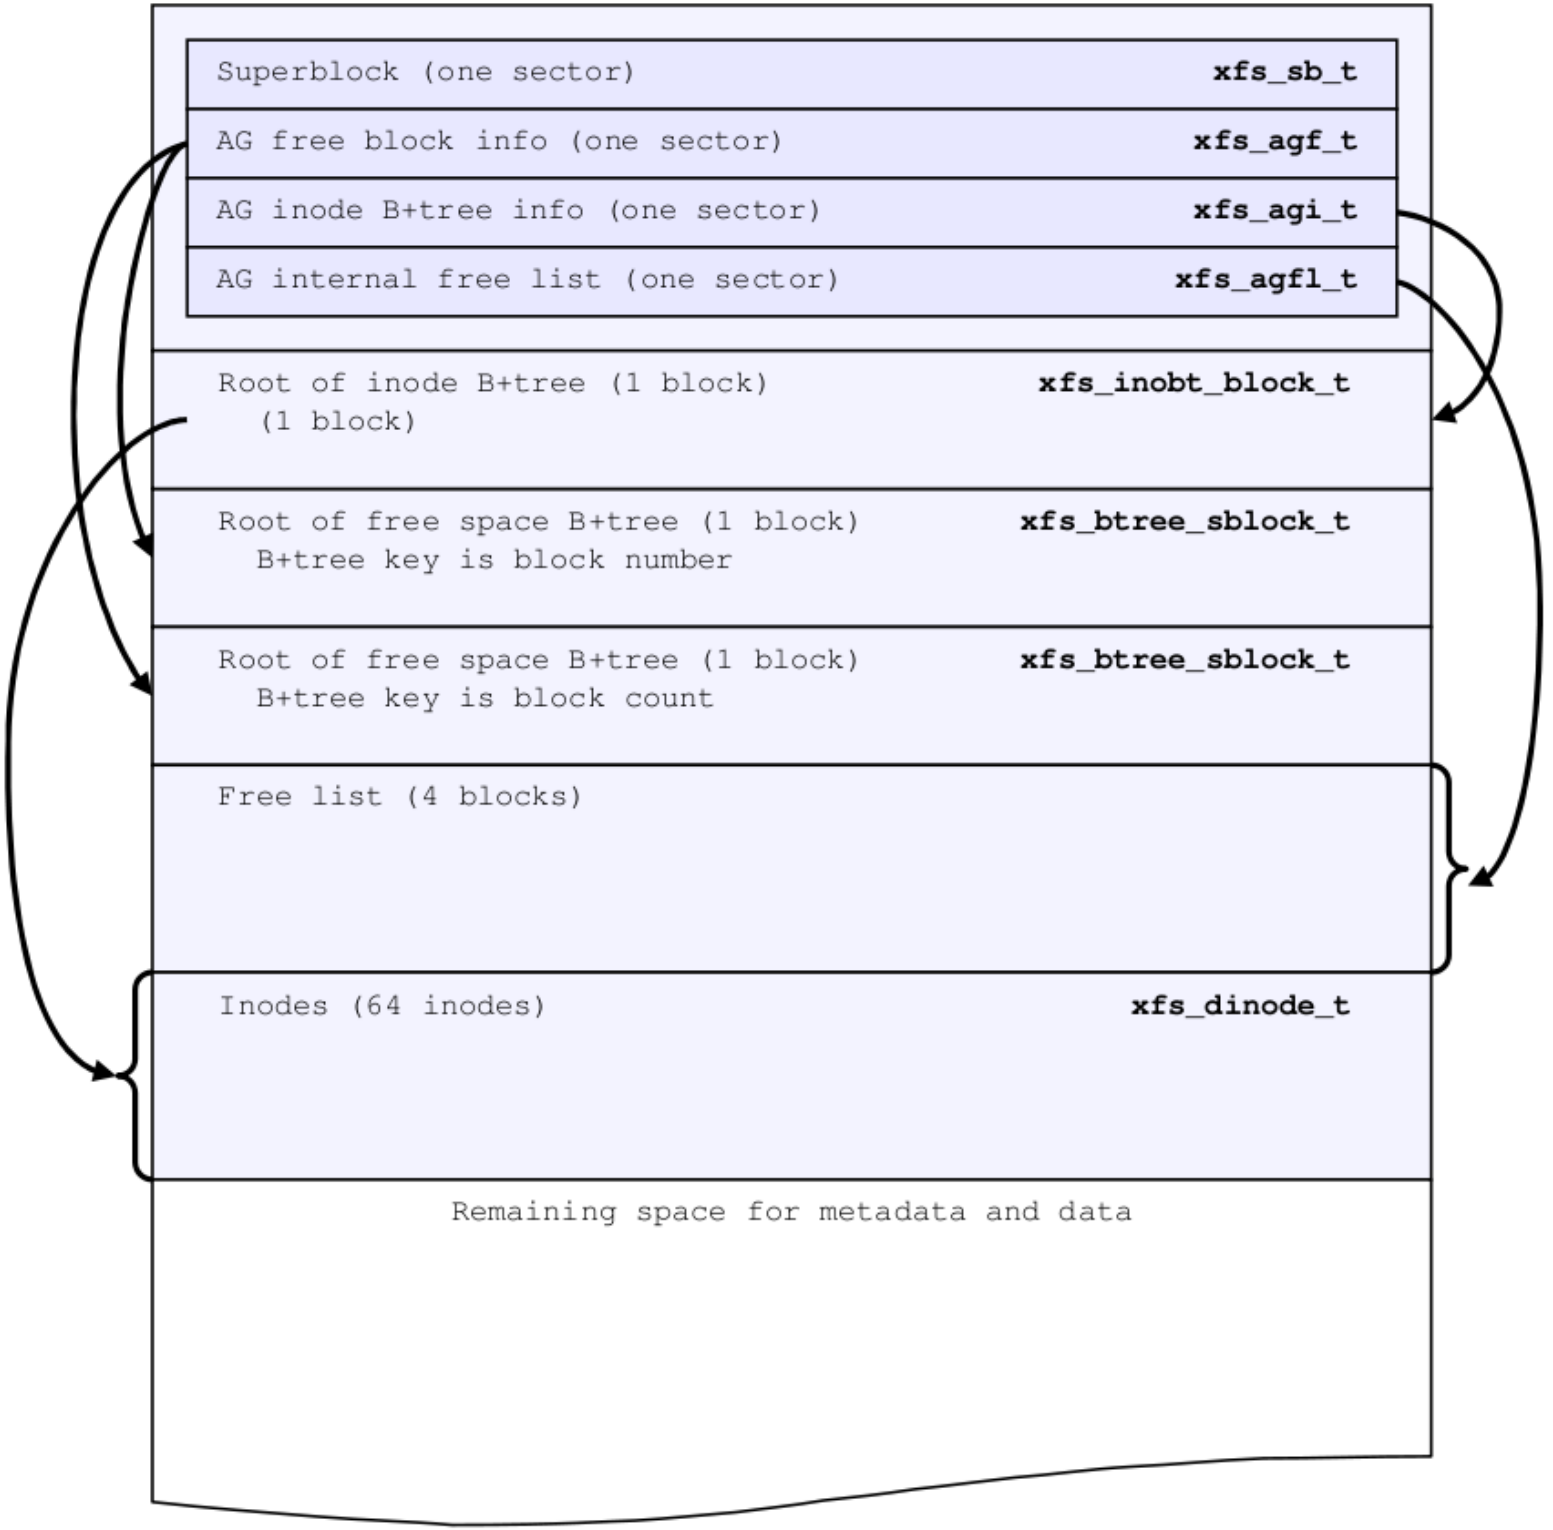

SOFTWARE UNIVERSITY

Log-structured File Systems Architect-ure

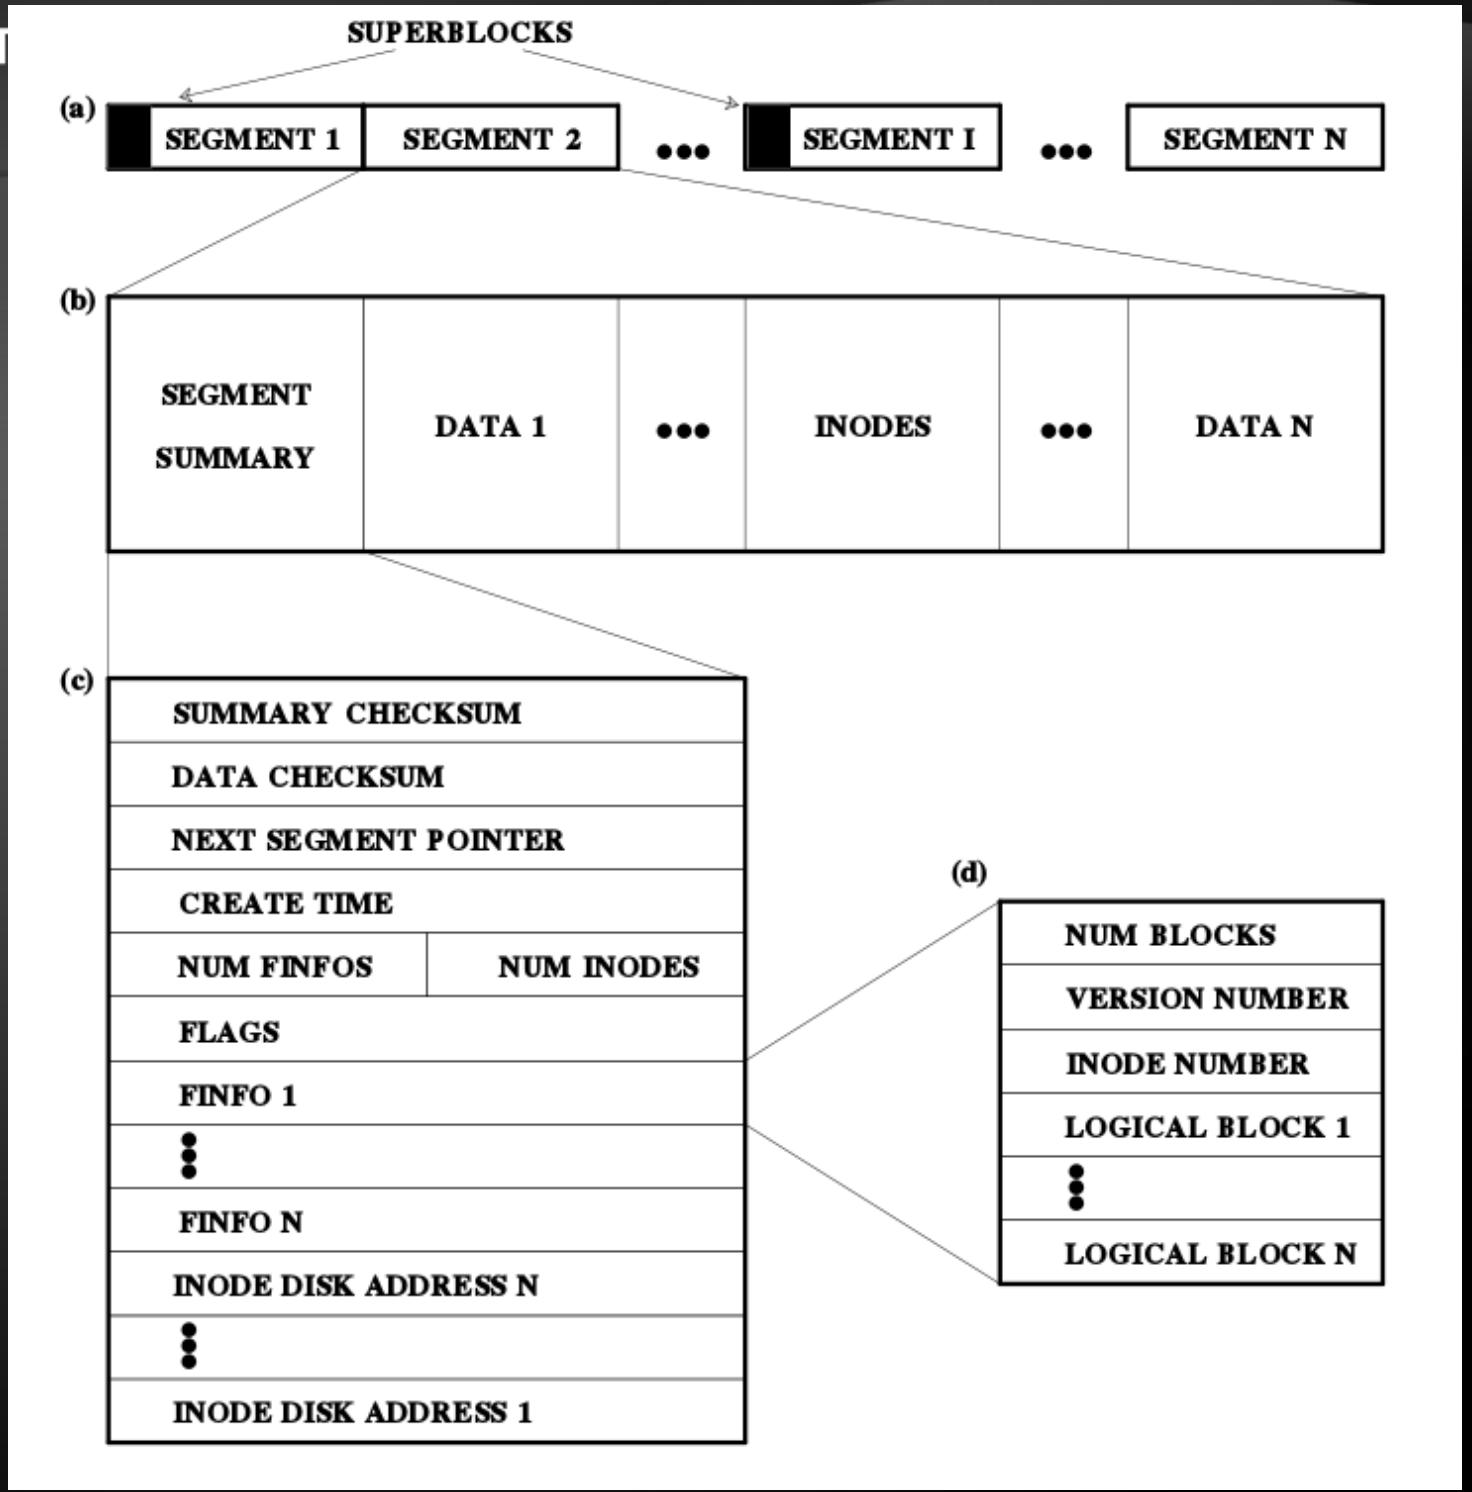

SOFTWARE UNIVERSITY

Log-structured File Systems Architecture

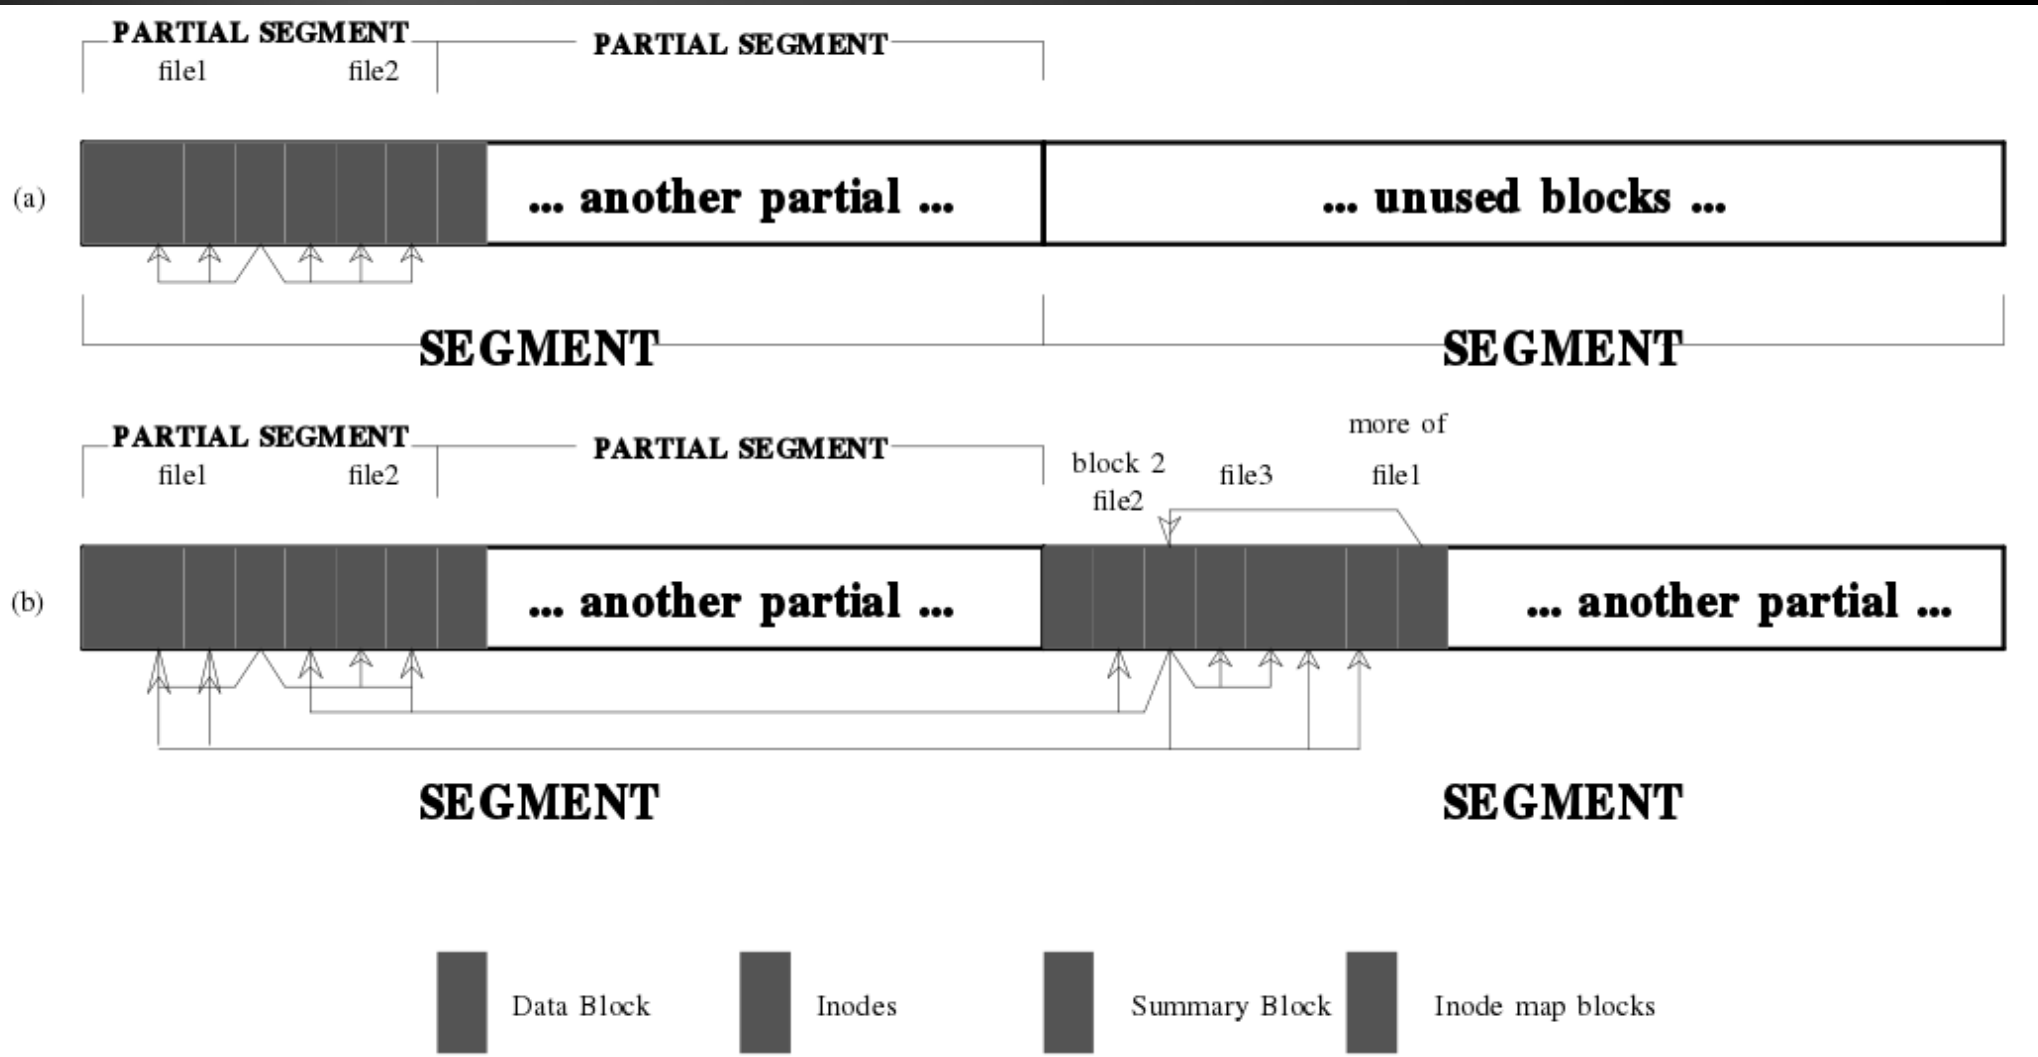

SOFTWARE UNIVERSITY

Log-structured NAND File Systems

	System requirement	JFFS2	YAFFS2	LogFS	UBIFS
1	Boot time	Poor	Good	Excellent	Good
2	I/O performance	Good	Good	Fair	Excellent
3	Resource usage	Fair	Excellent	Good	Fair
4	NAND device life expectancy	Good	Fair	N/A	Excellent
5	Tolerance for unexpected power-off	Good	Good	Poor	Good
6	Integrated in mainline	Yes	No	Yes	Yes

NILFS2
F2FS

SOFTWARE UNIVERSITY

Pseudo File Systems

- procfs
- sysfs
- debugfs
- configfs
- tmpfs
- others

SOFTWARE UNIVERSITY

Pseudo File Systems

- debugfs is designed to provide Kernel Devs with simple way to push data into User space
- configfs is for creating, managing and destroying kernel objects from user-space
- sysfs is for viewing and manipulating objects from user-space which are created and destroyed by kernel space

SOFTWARE UNIVERSITY

Pseudo File Systems

- procfs is the first FS to provide easy access to kernel-space from user-space
- tmpfs is a very fast in-memory file system

SOFTWARE UNIVERSITY

Network File Systems

- Network File System – NFS v3/v4
- Common Internet File System - CIFS

SOFTWARE UNIVERSITY

Network File Systems

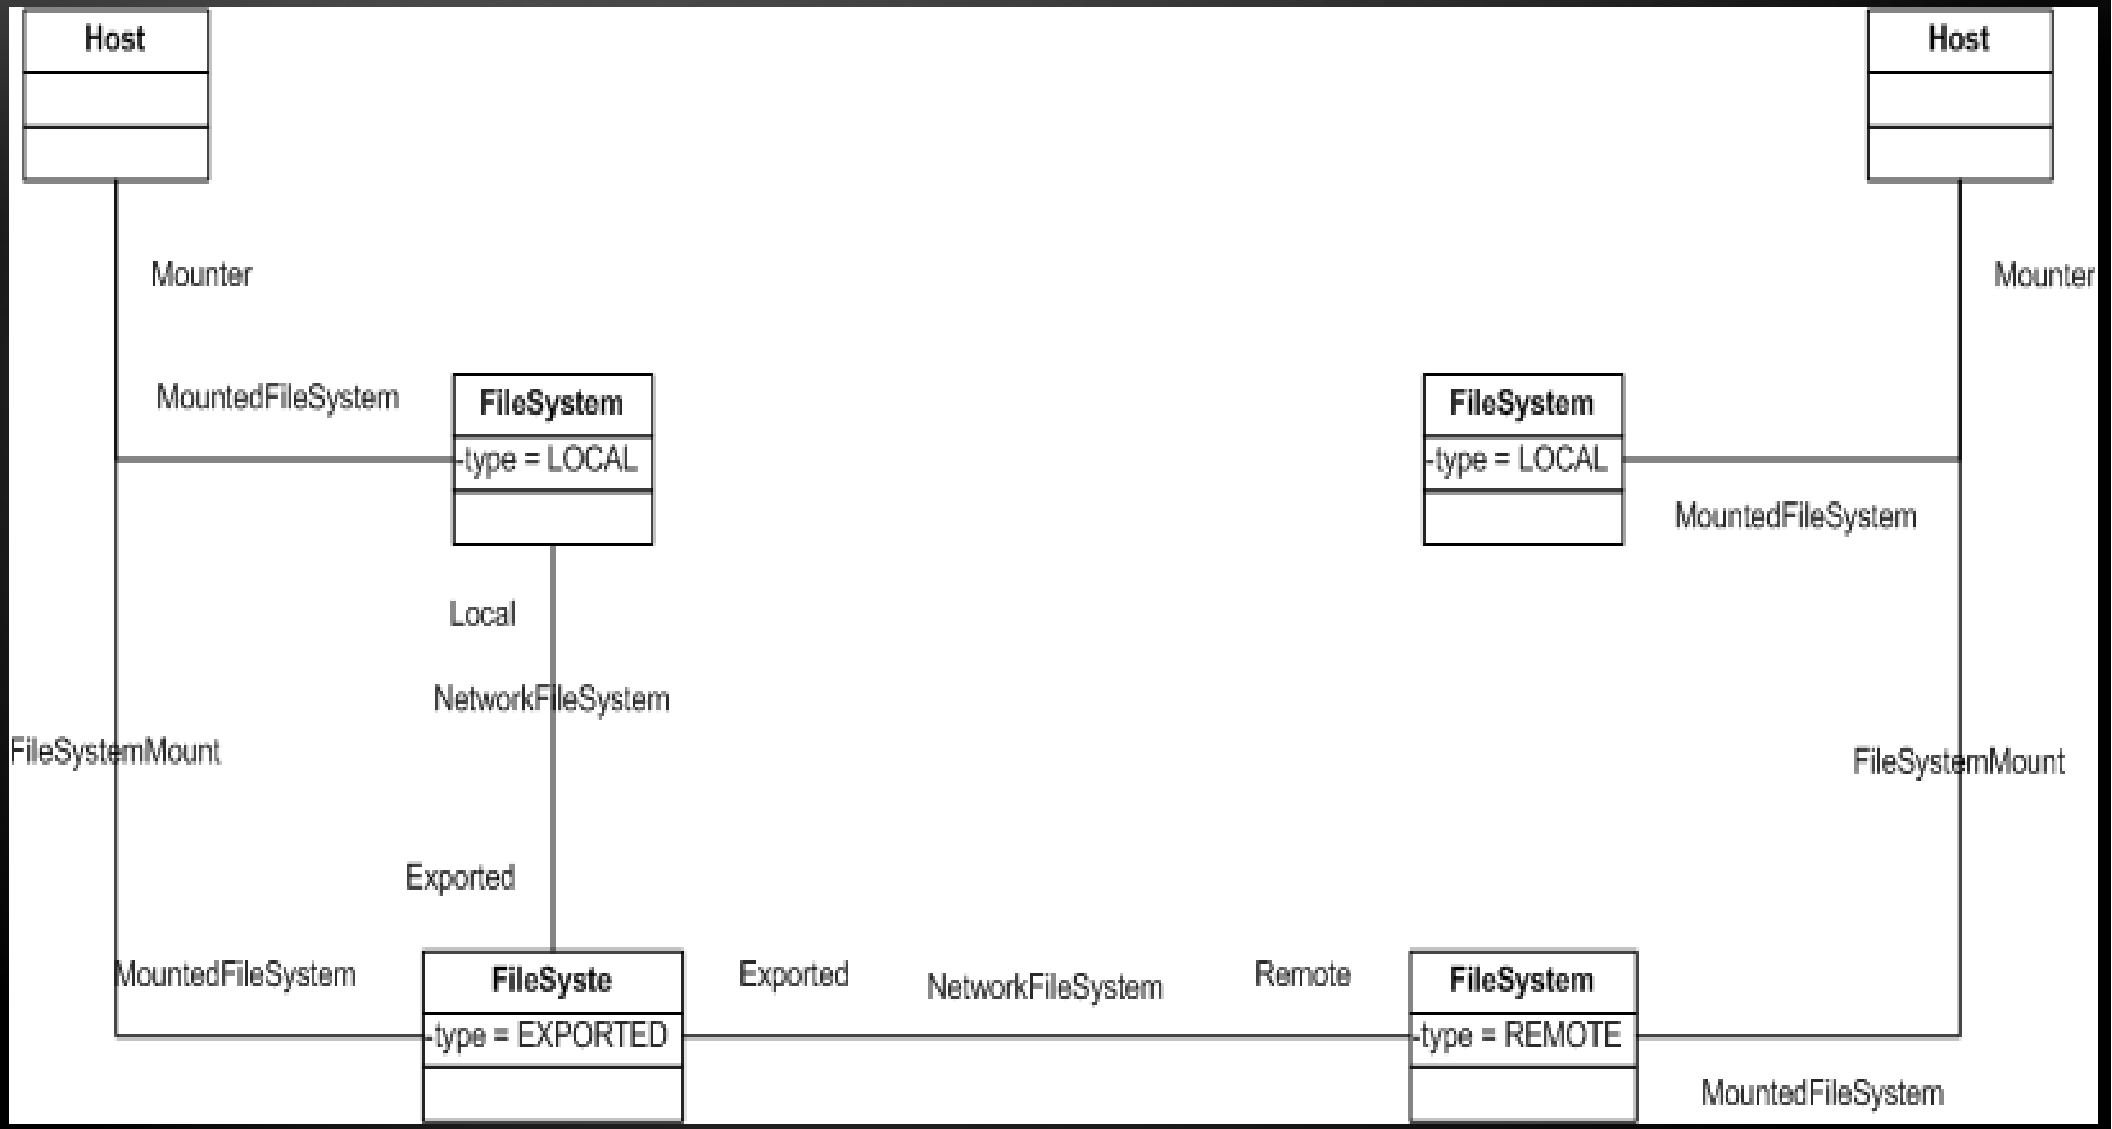

SOFTWARE UNIVERSITY

Network File Systems

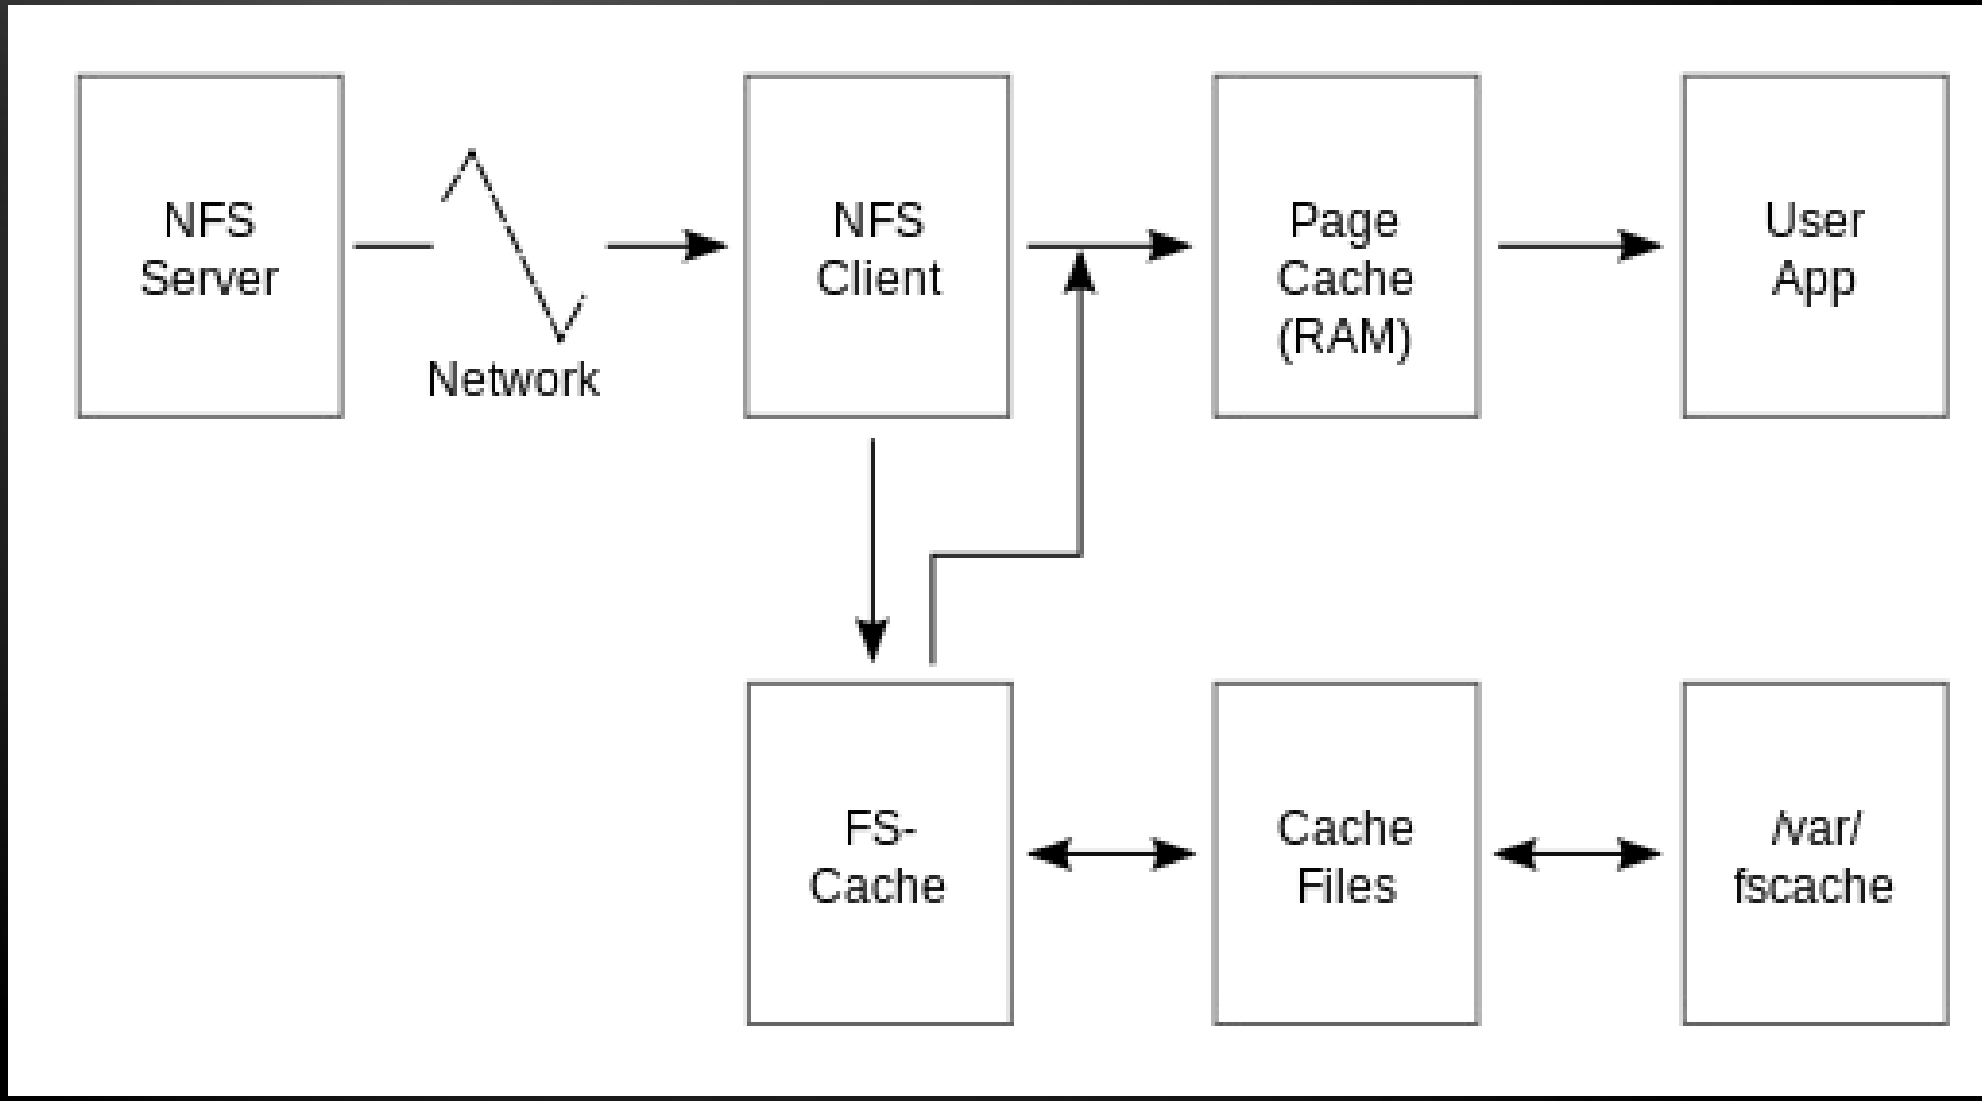

SOFTWARE UNIVERSITY

Network File System

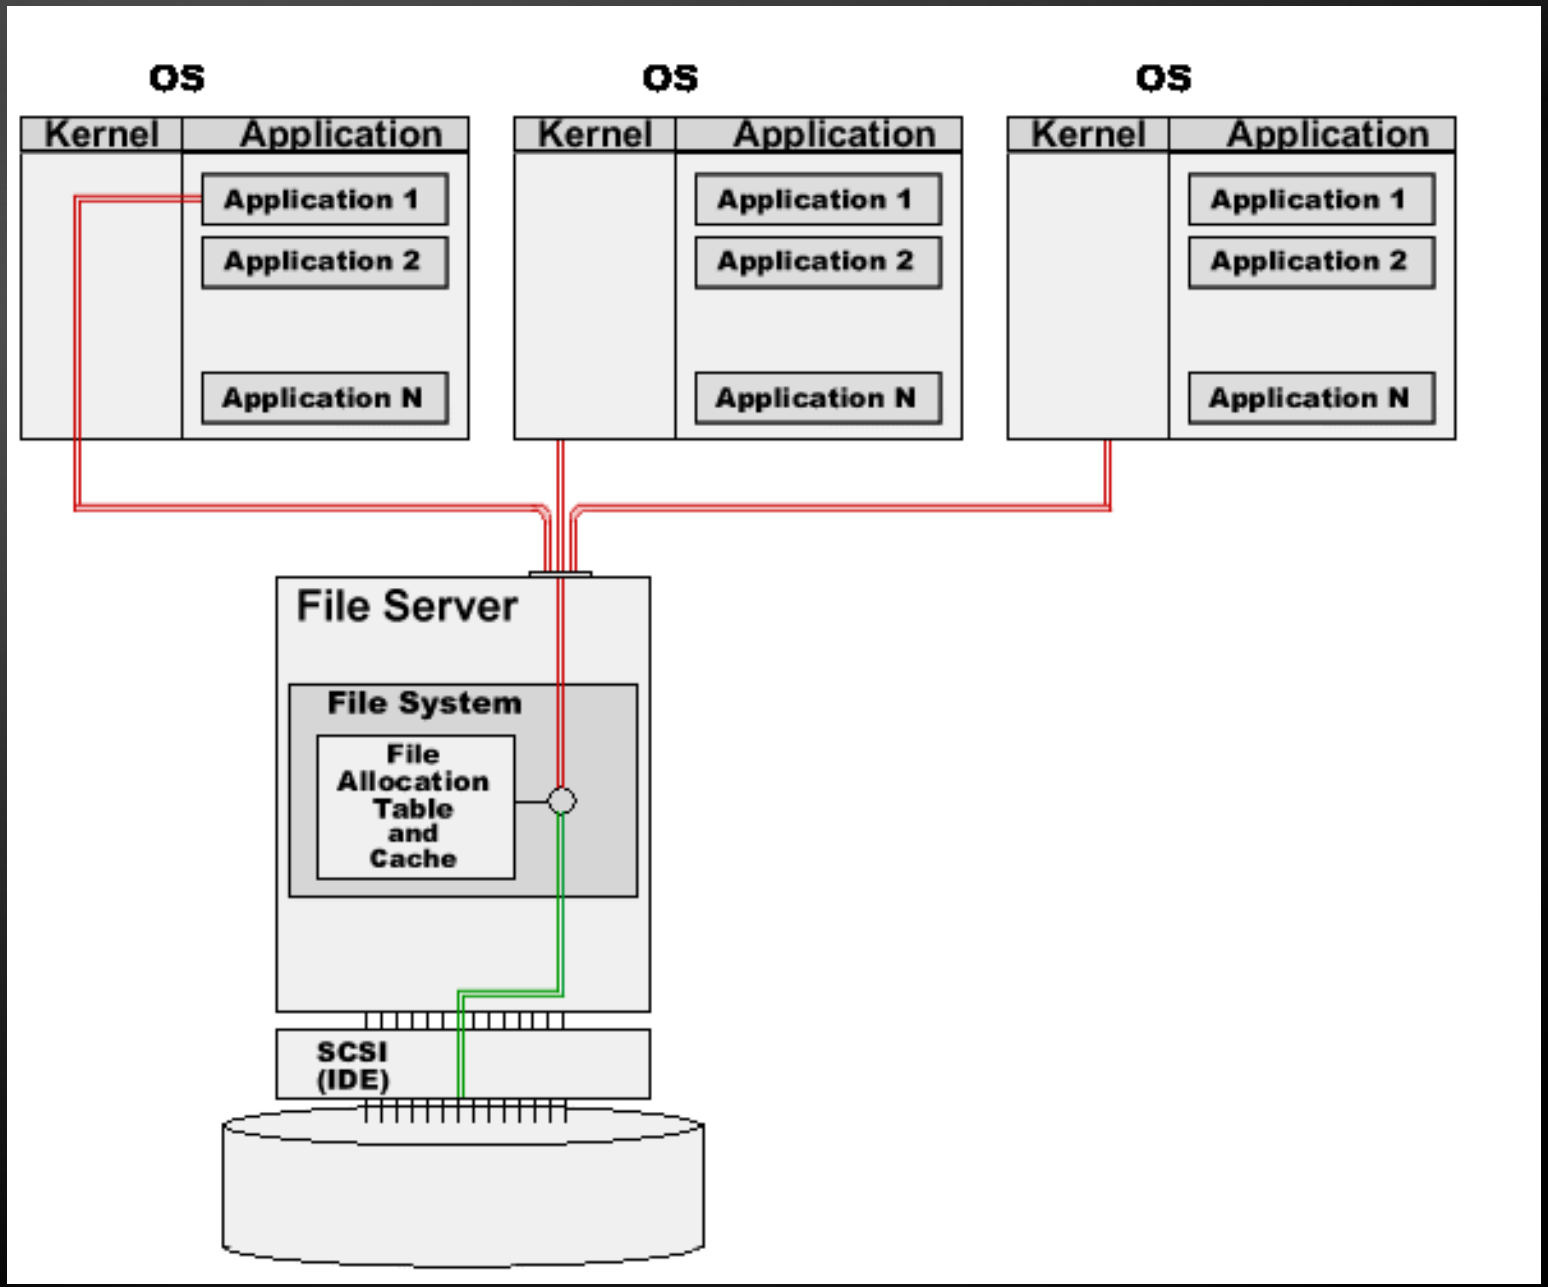

SOFTWARE UNIVERSITY

Cluster File System

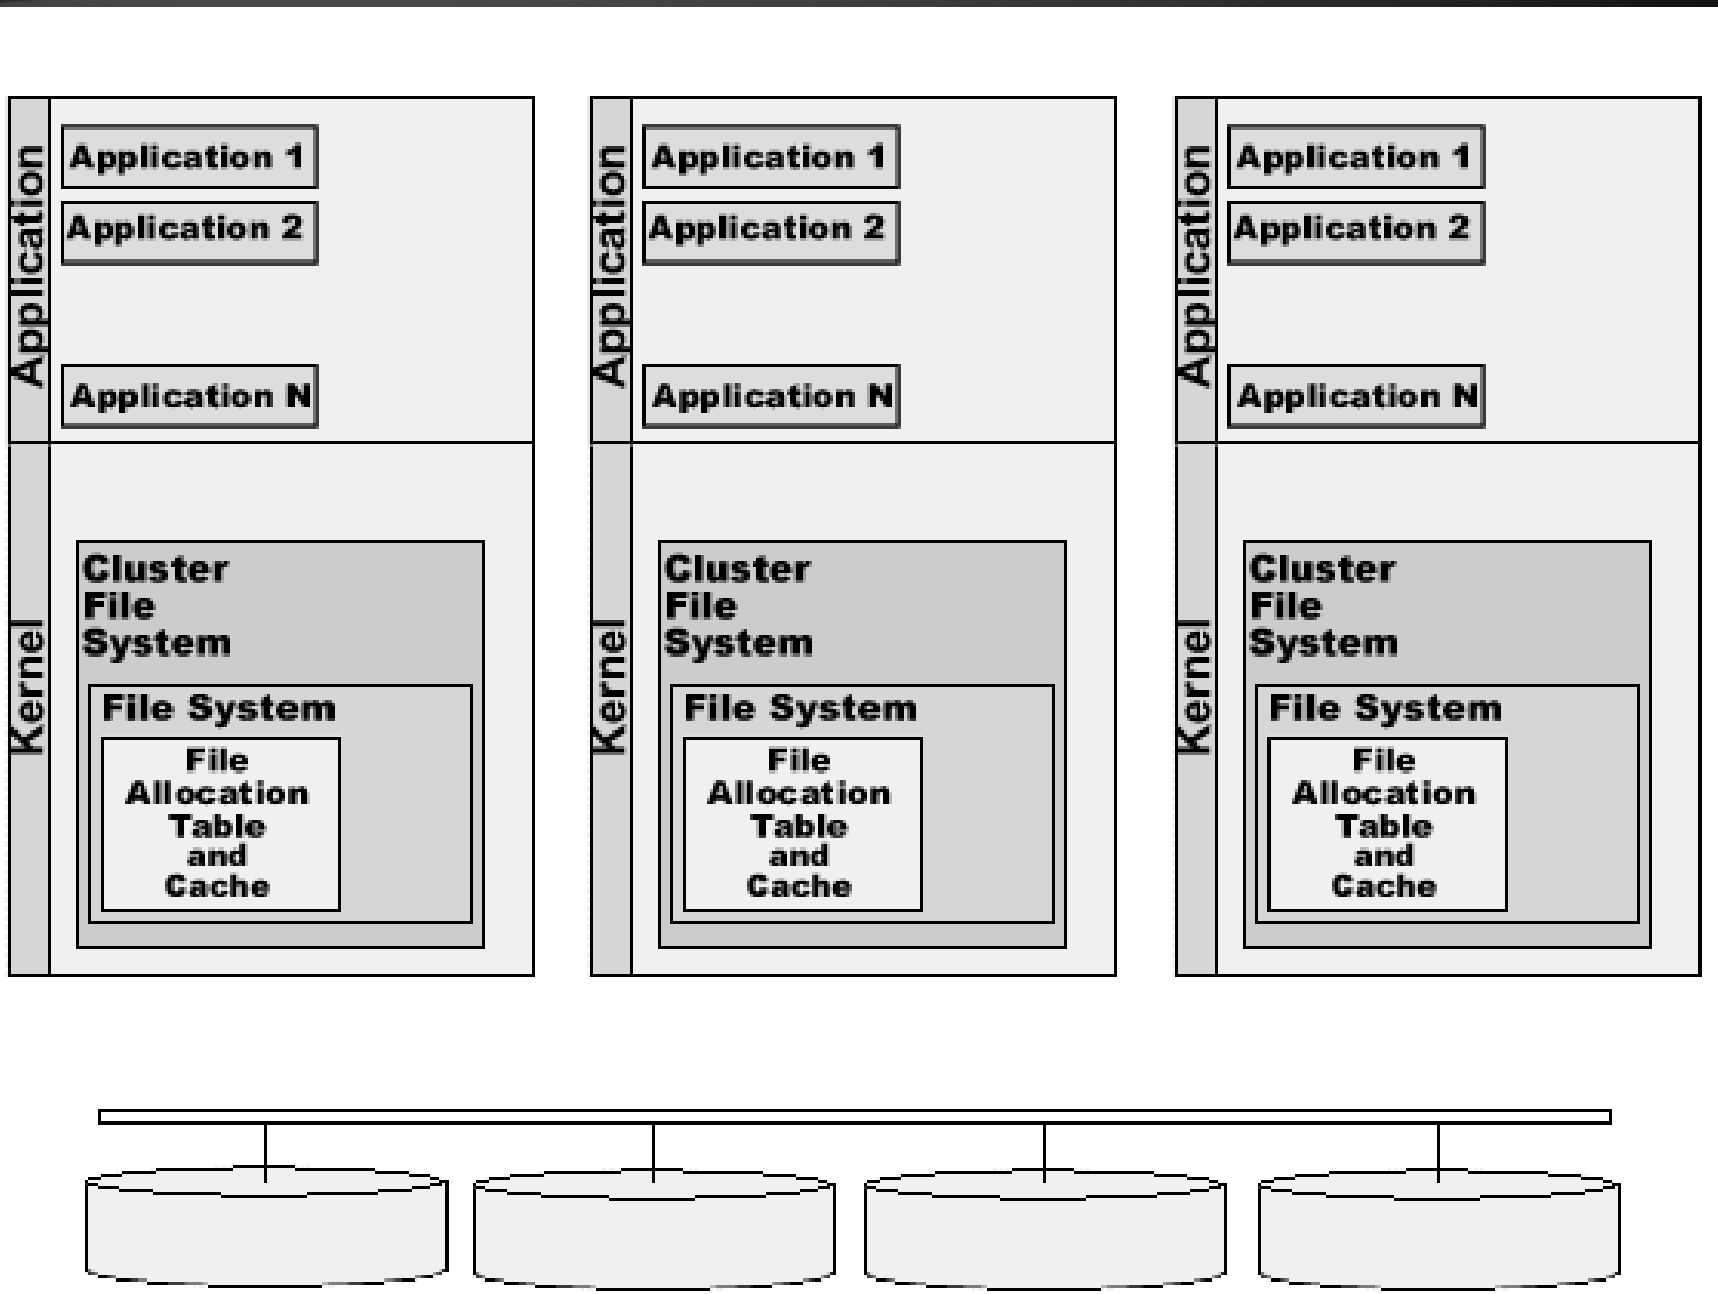

SOFTWARE UNIVERSITY

Cluster File Systems

- GFS, GFS2
- OCFS2

SOFTWARE UNIVERSITY

Cluster File Systems - GFS2

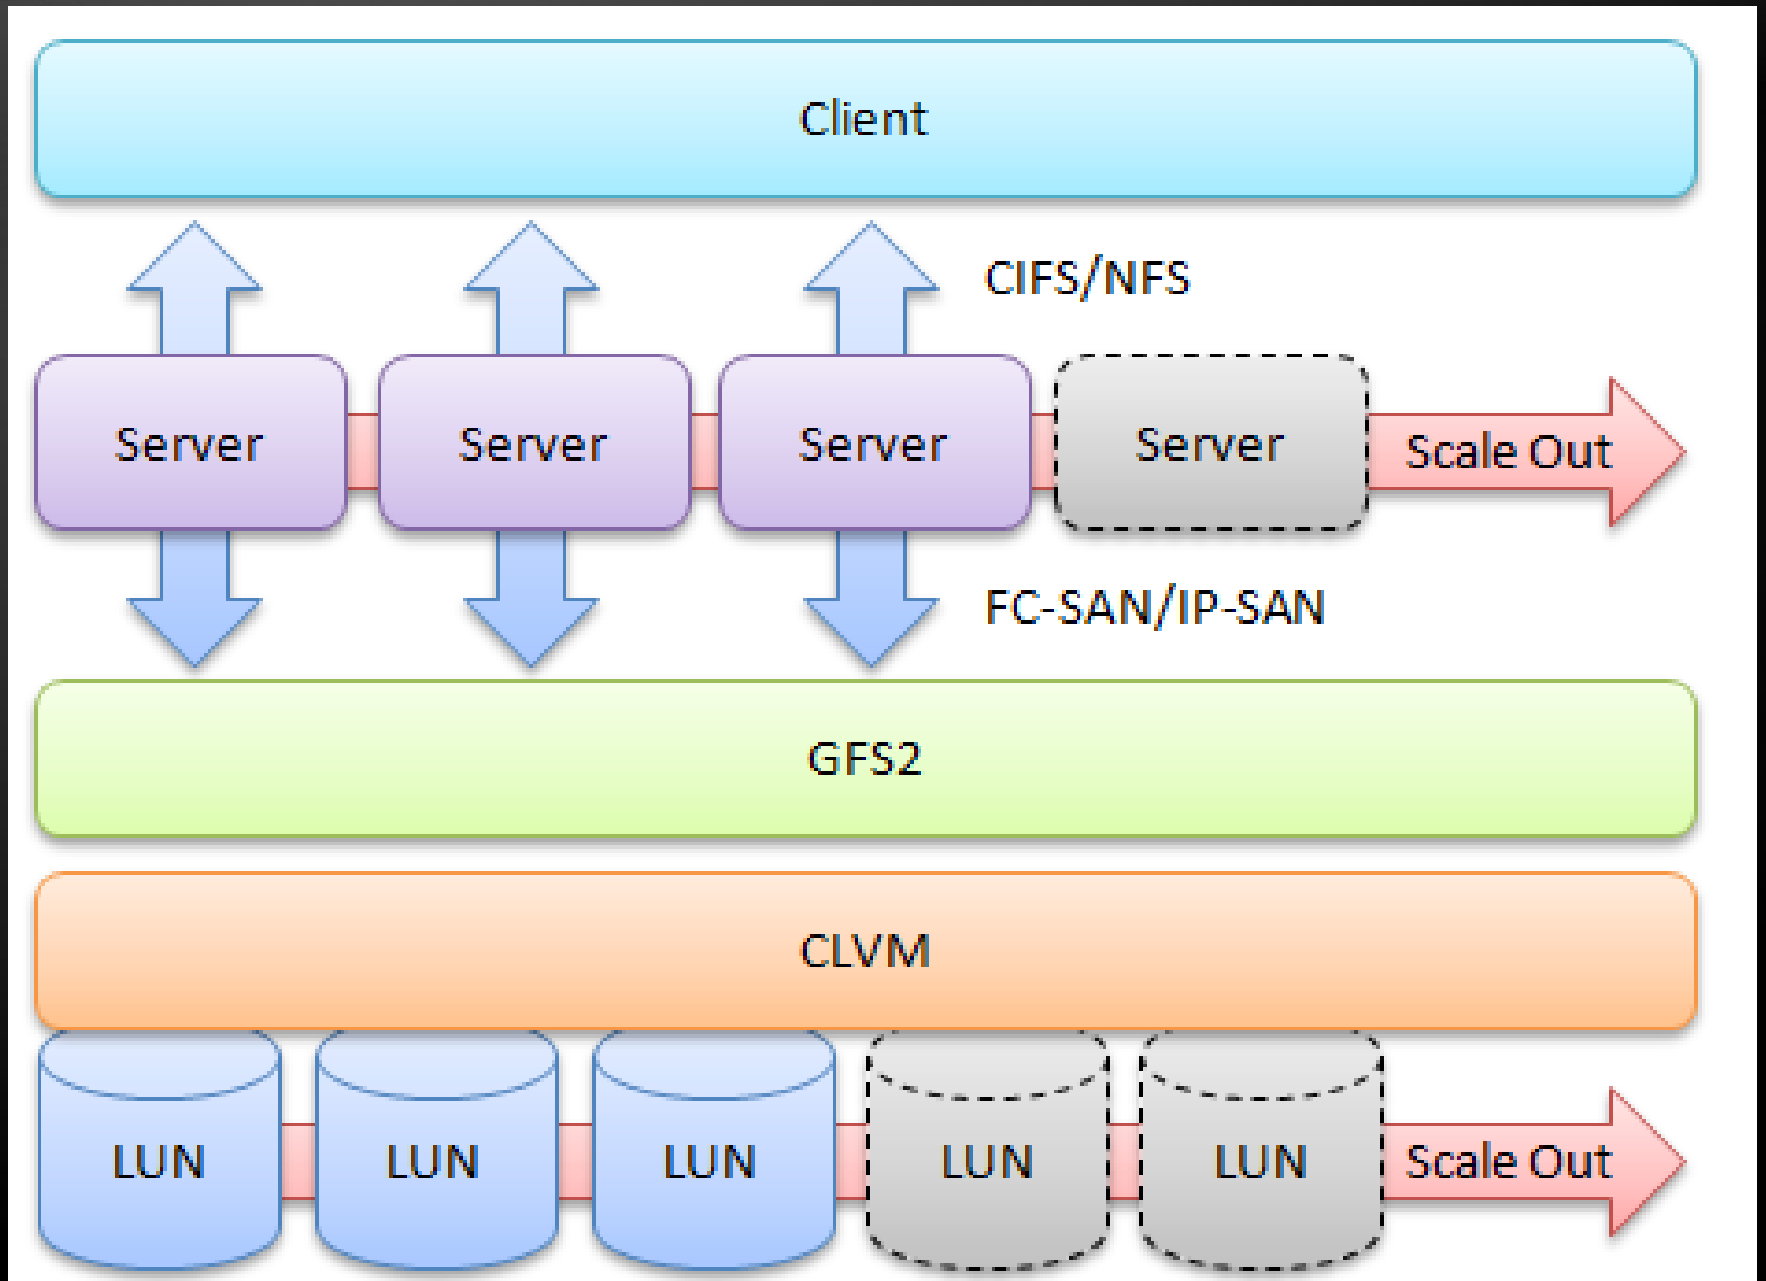

SOFTWARE UNIVERSITY

Cluster File Systems - GFS2

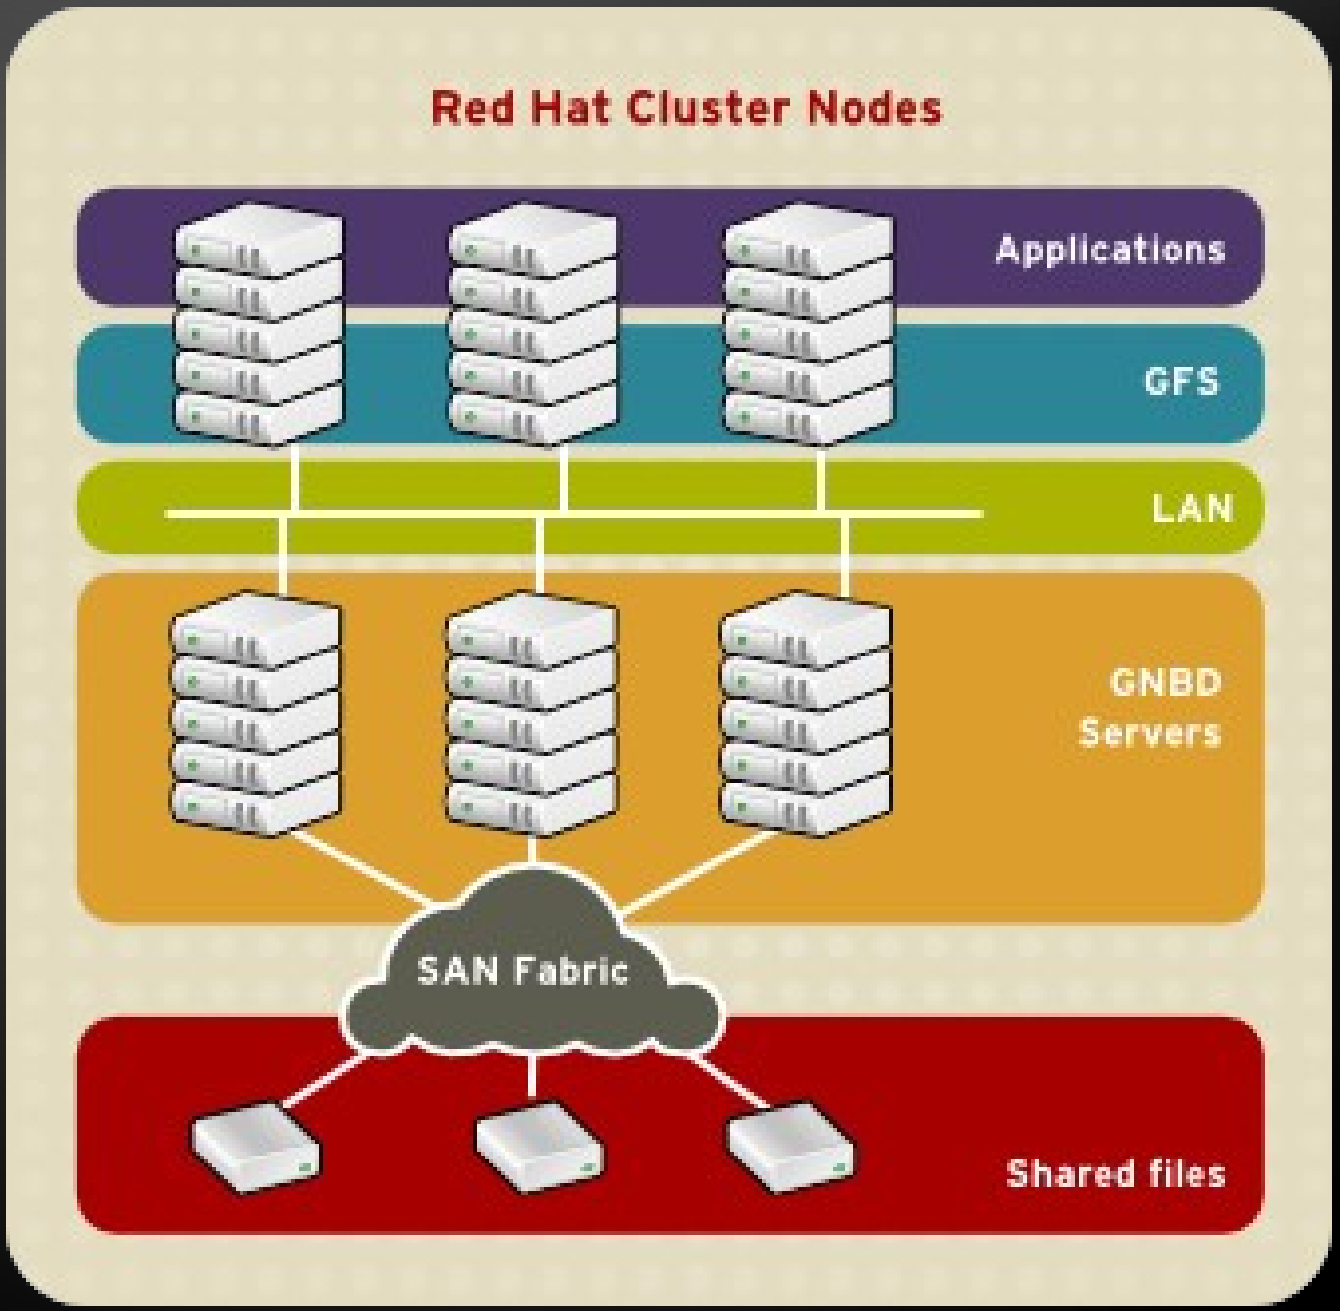

SOFTWARE UNIVERSITY

Cluster File Systems - OCFS2

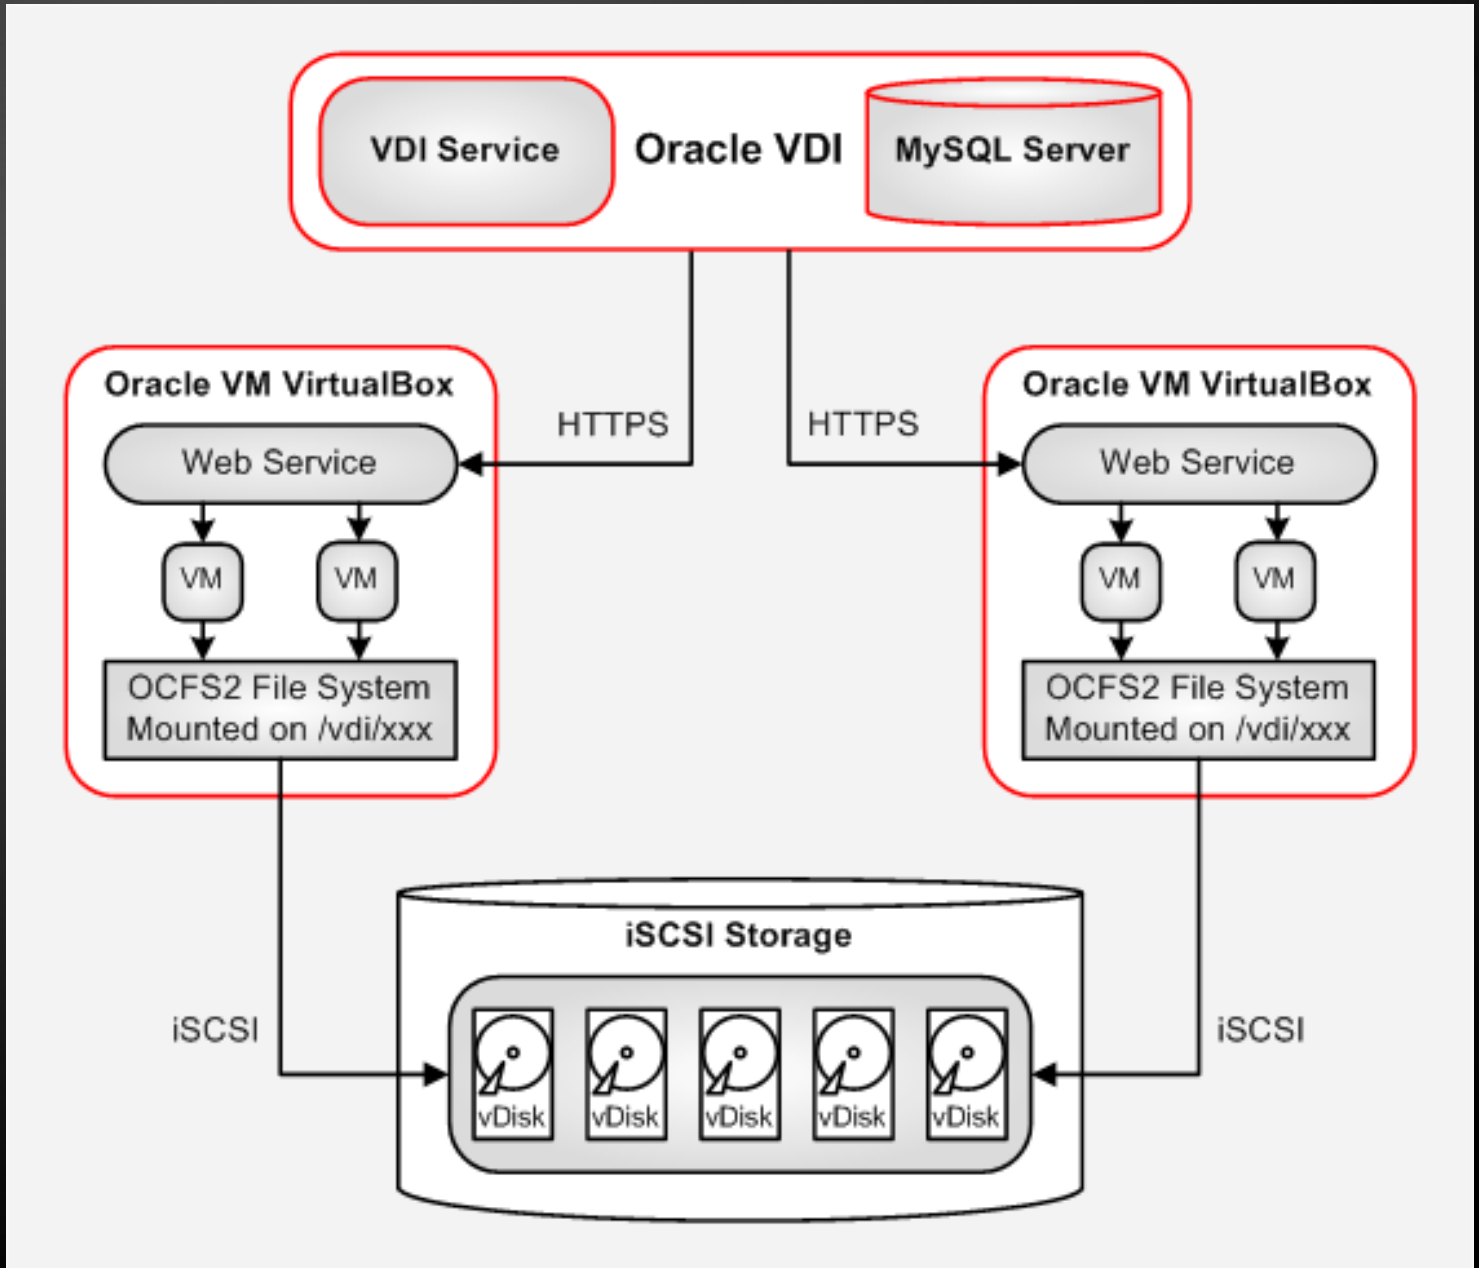

SOFTWARE UNIVERSITY

Cluster File Systems - OCFS2

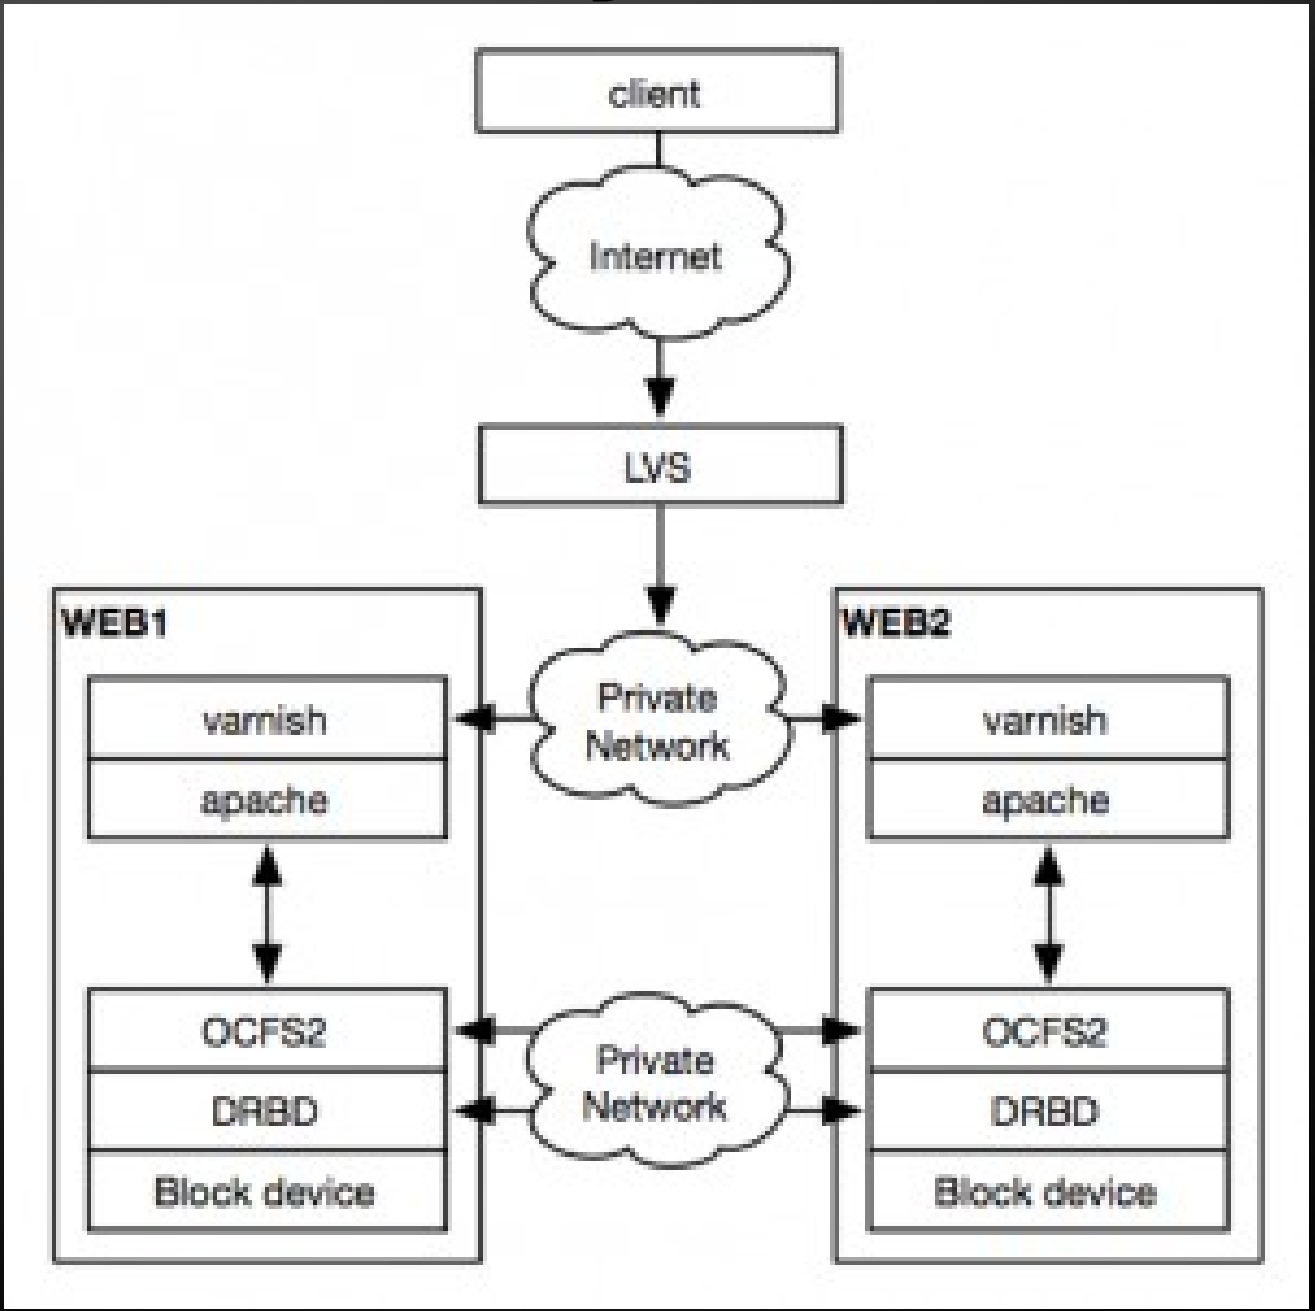

SOFTWARE UNIVERSITY

Distributed File Systems

- Hadoop
- Lustre
- GlusterFS
- GFarm
- FhgFS
- PohmelFS
- Ceph
- PVFS2

SOFTWARE UNIVERSITY

Hadoop

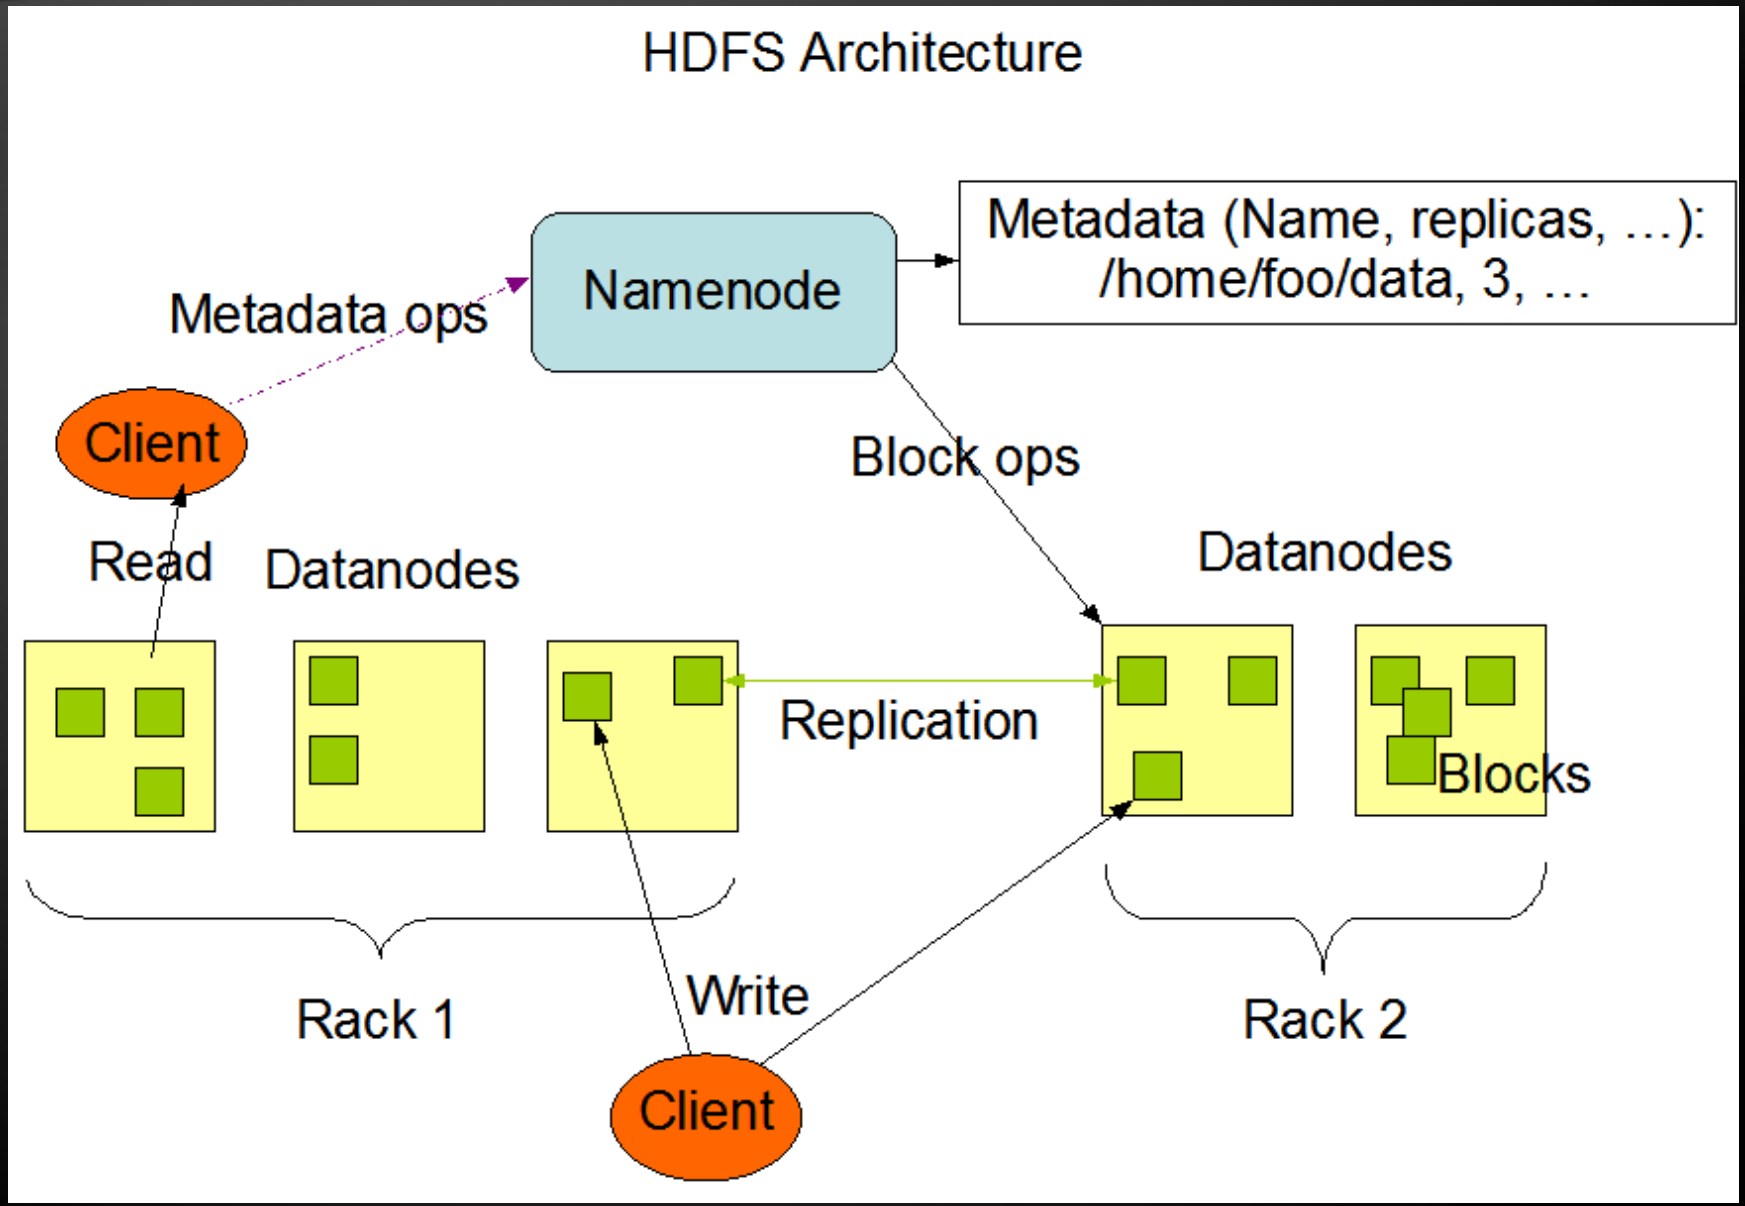

SOFTWARE UNIVERSITY

Hadoop

- Large block FS – 64MB
- Write mostly FS
- Writes smaller then one block wait
- Adding/removing nodes requires restart of the cluster

SOFTWARE UNIVERSITY

GlusterFS

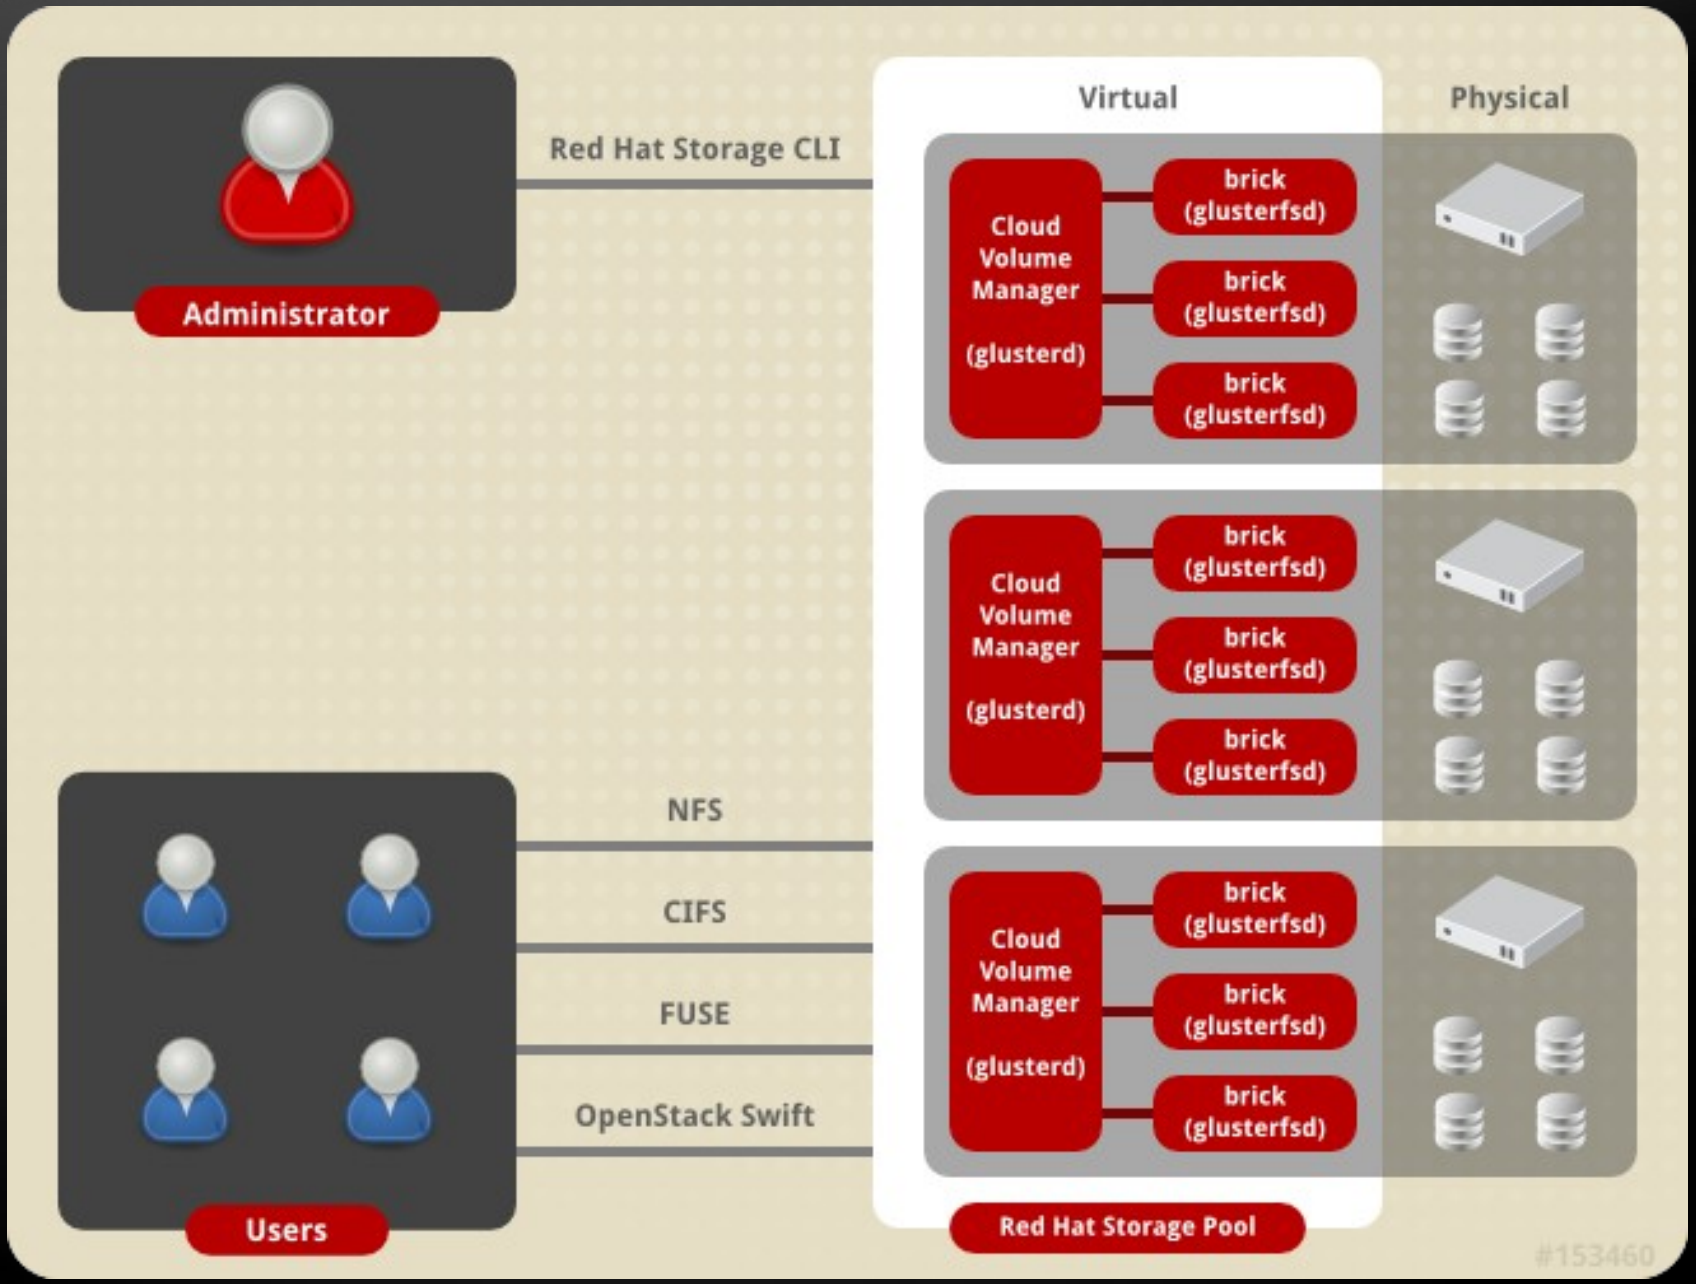

SOFTWARE UNIVERSITY

GlusterFS

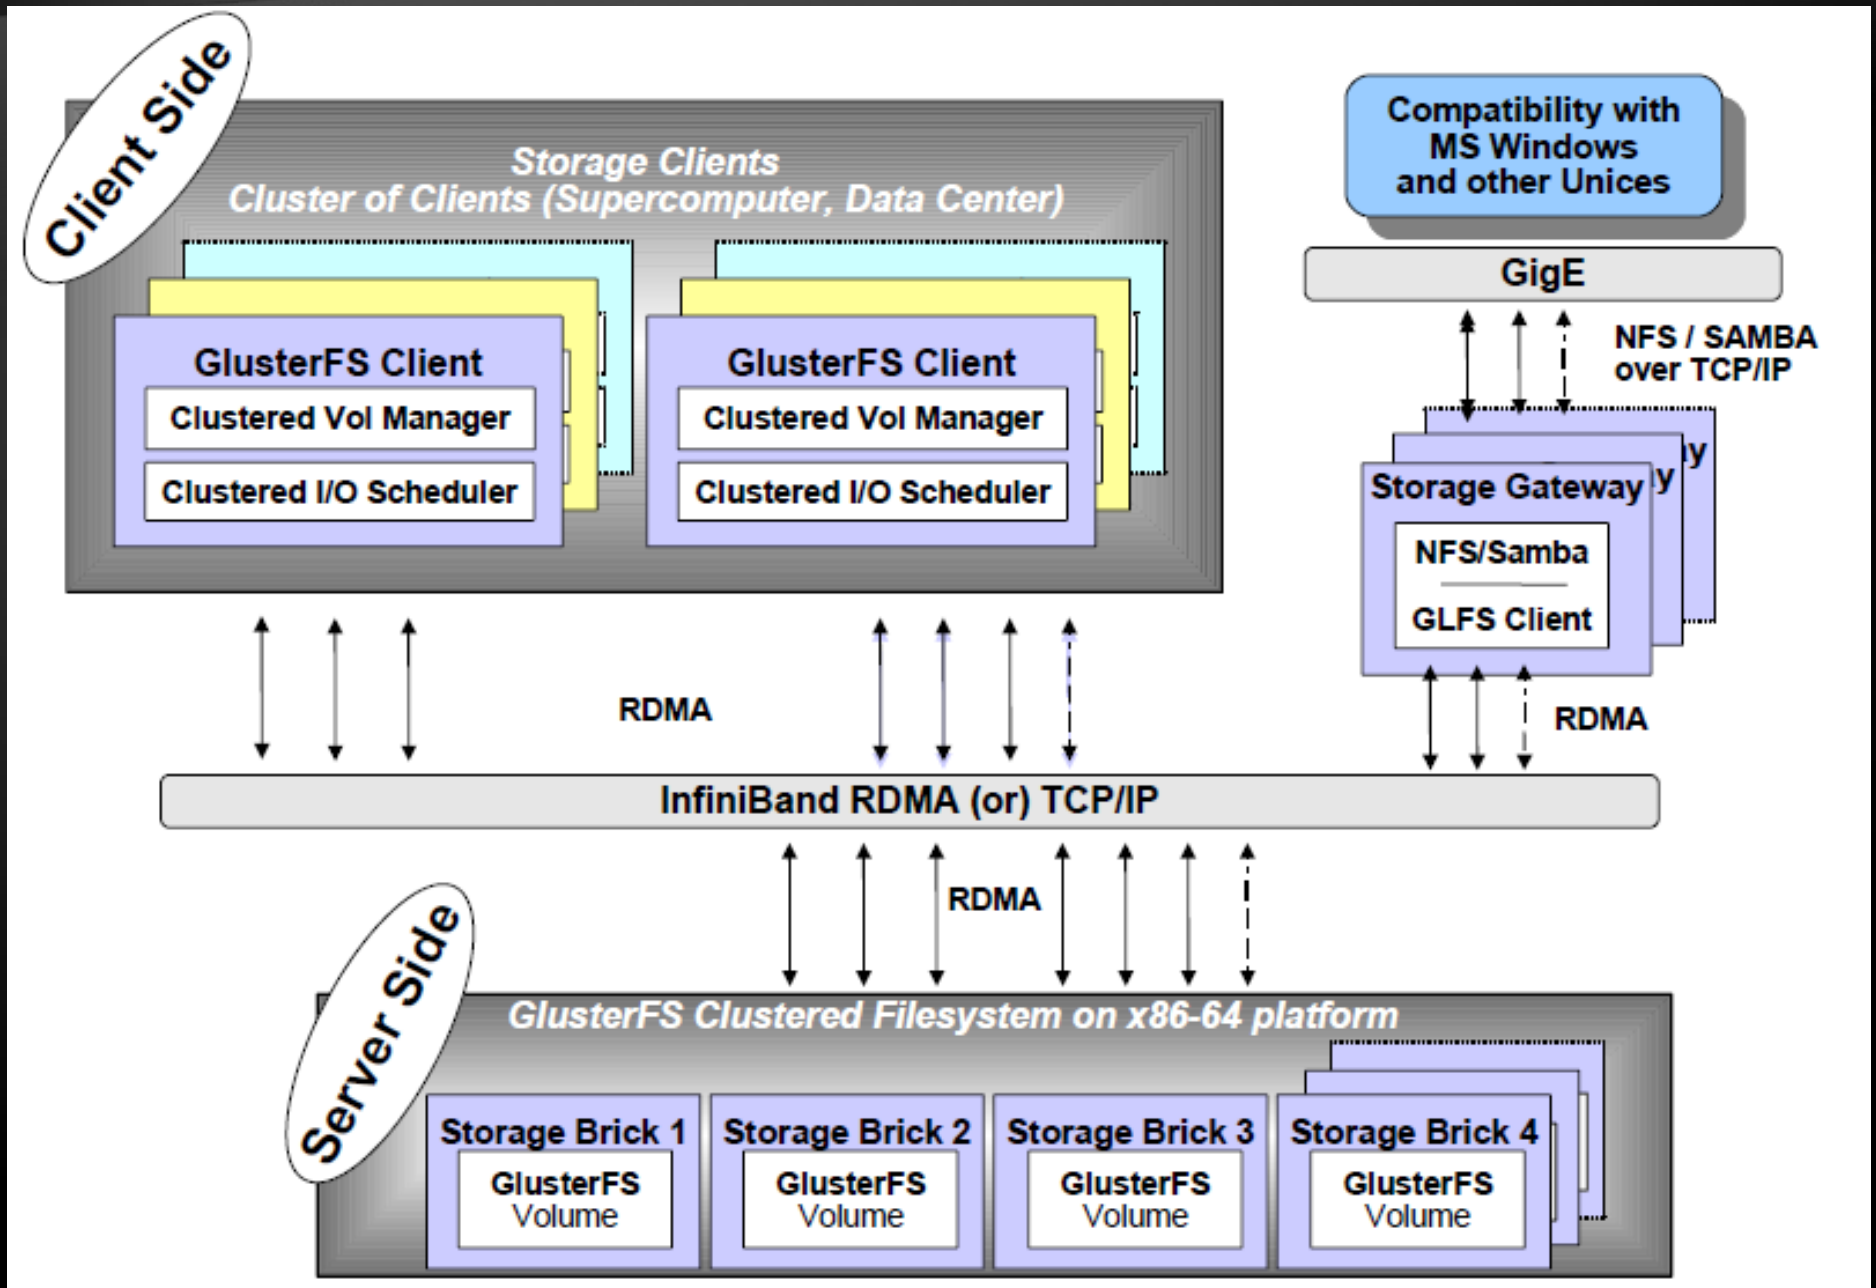

SOFTWARE UNIVERSITY

GFarm

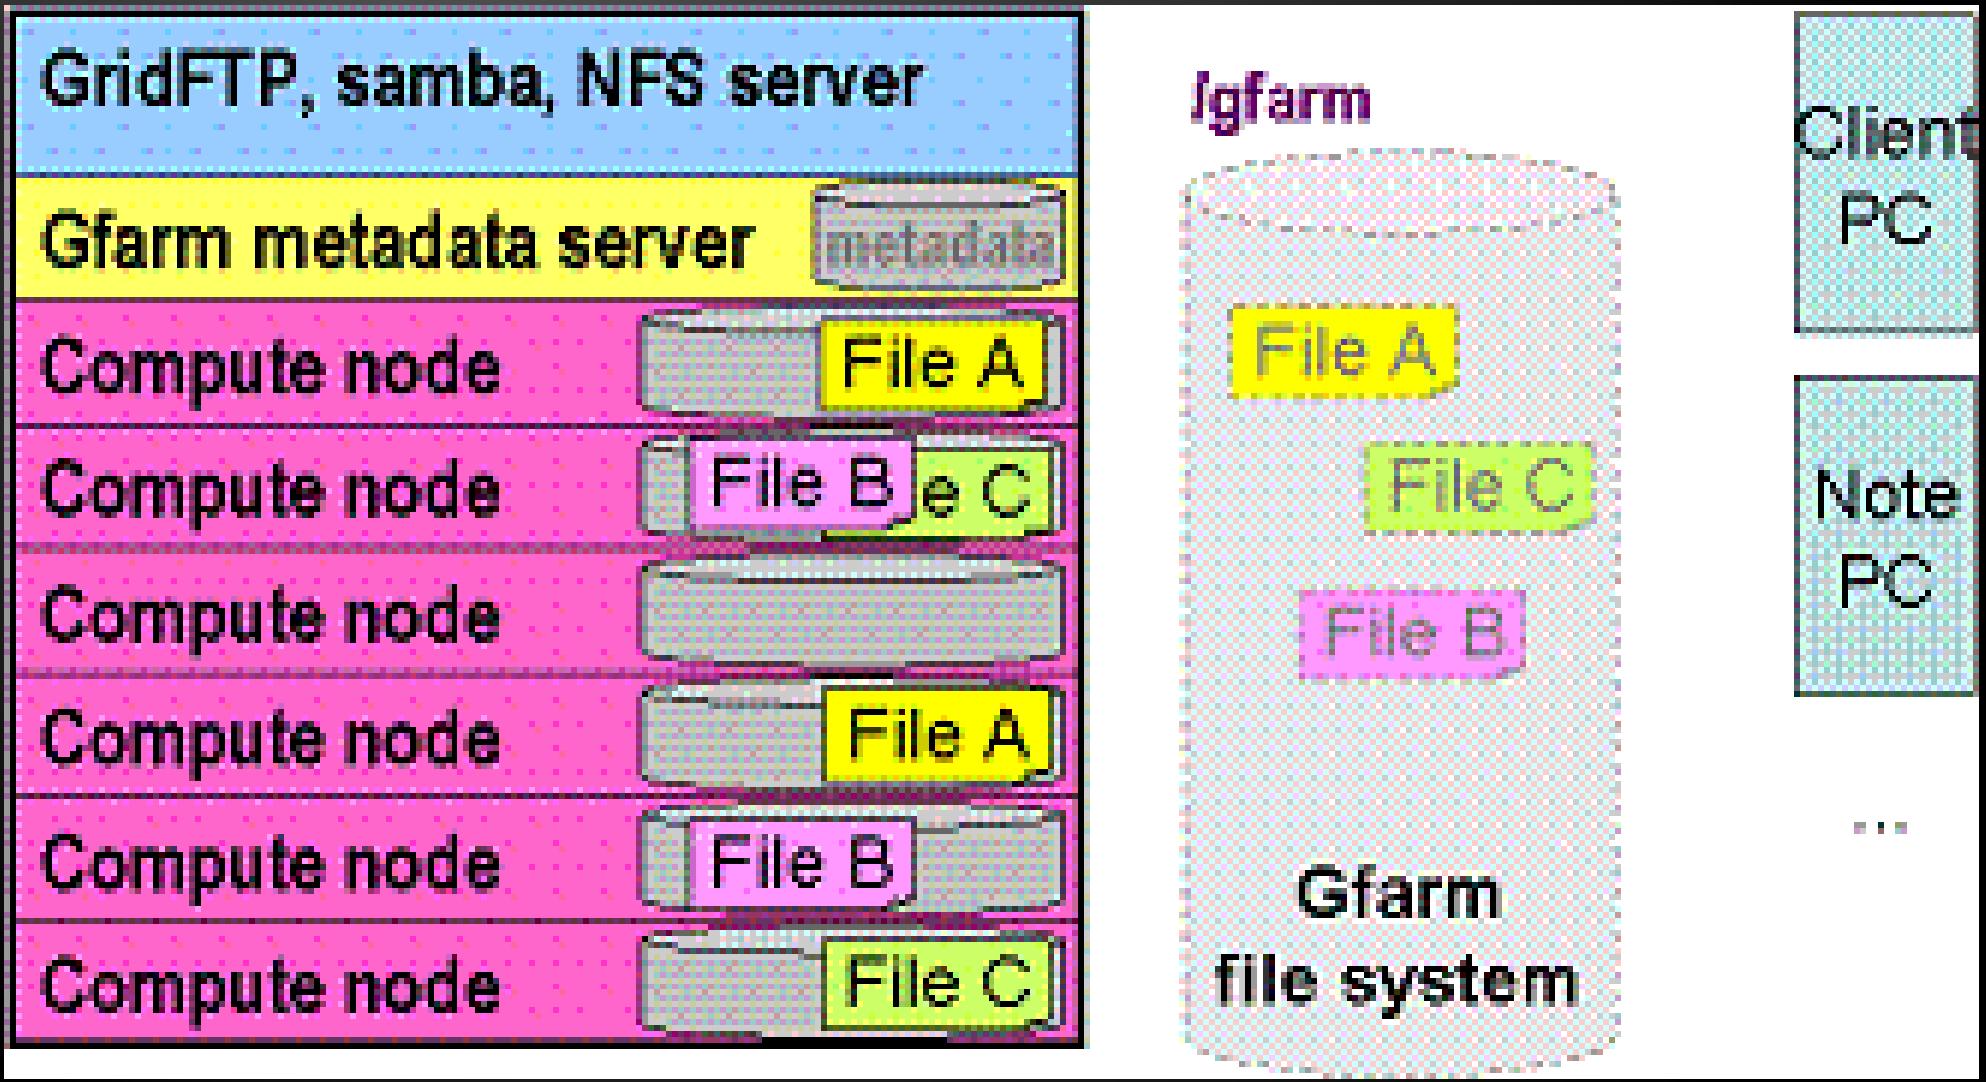

SOFTWARE UNIVERSITY

GFarm

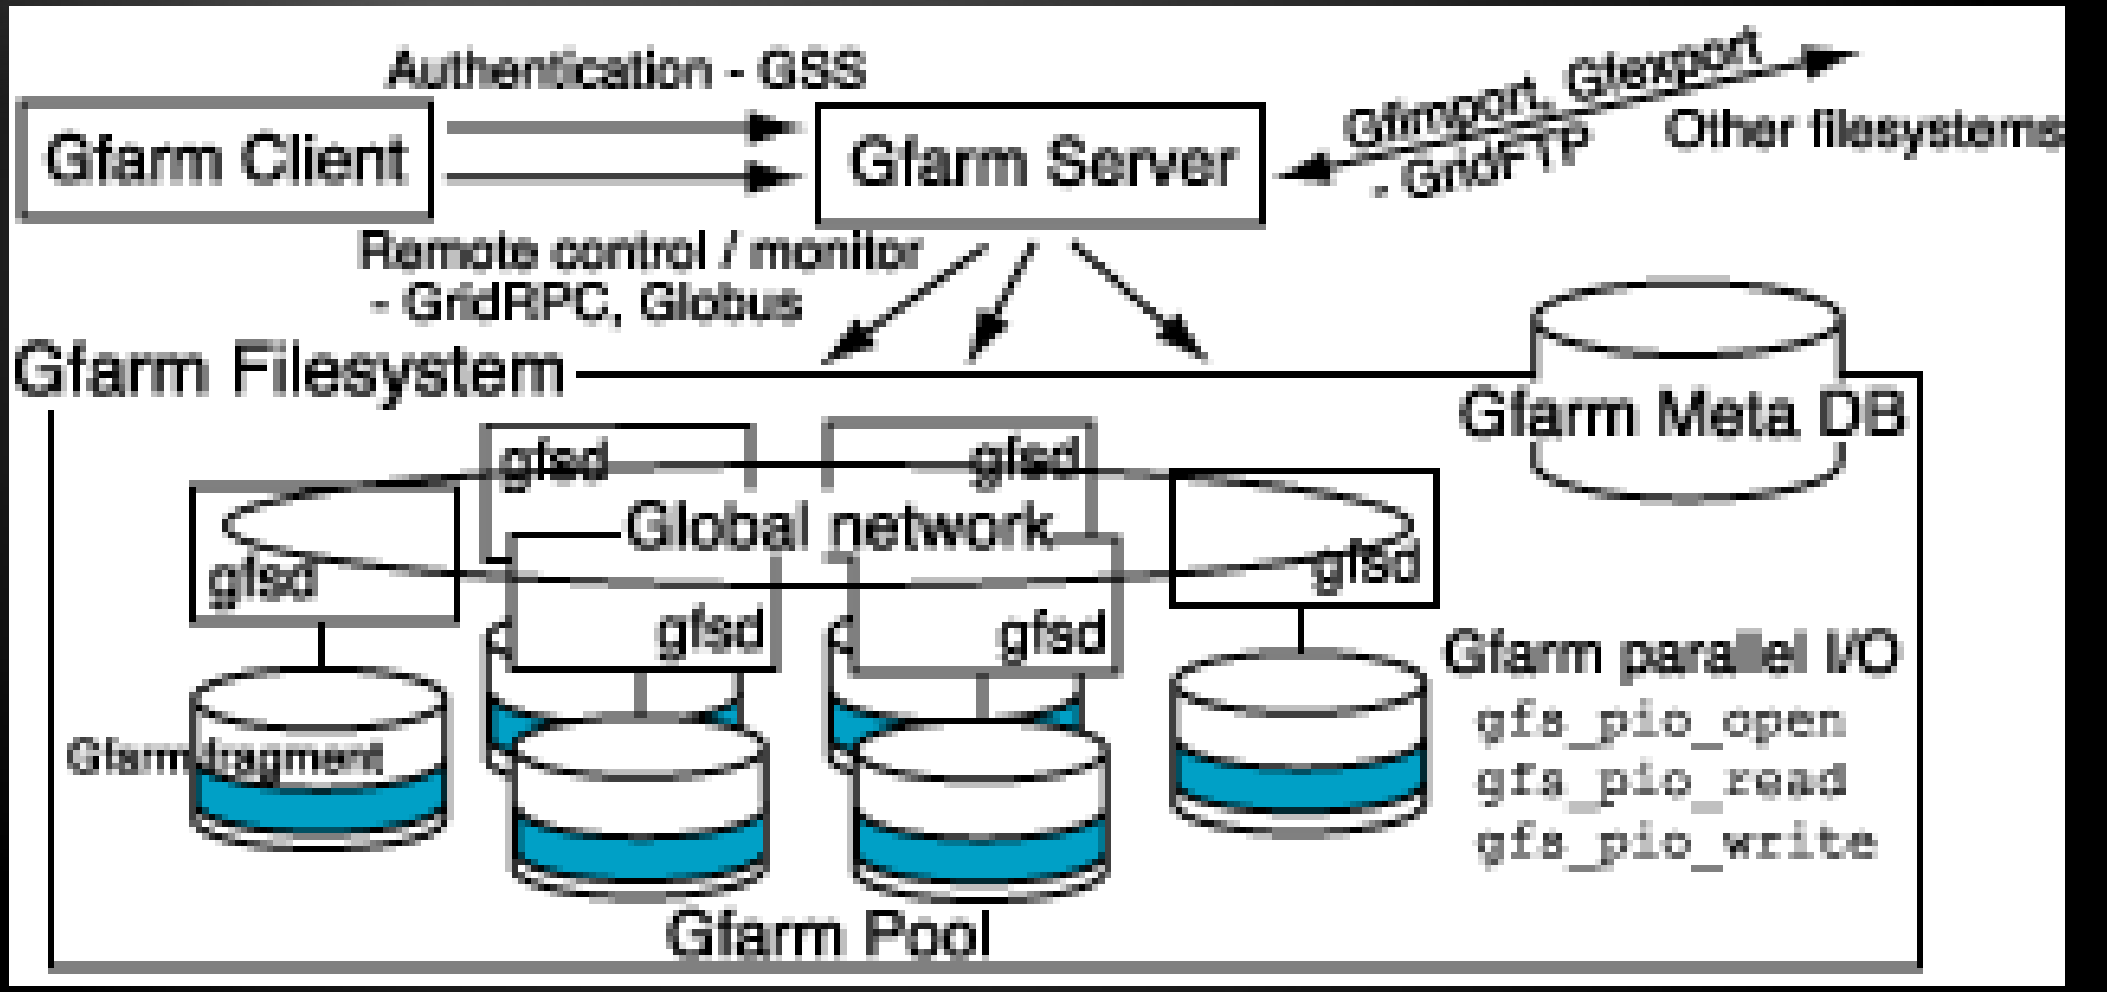

SOFTWARE UNIVERSITY

Linux Directory Structure

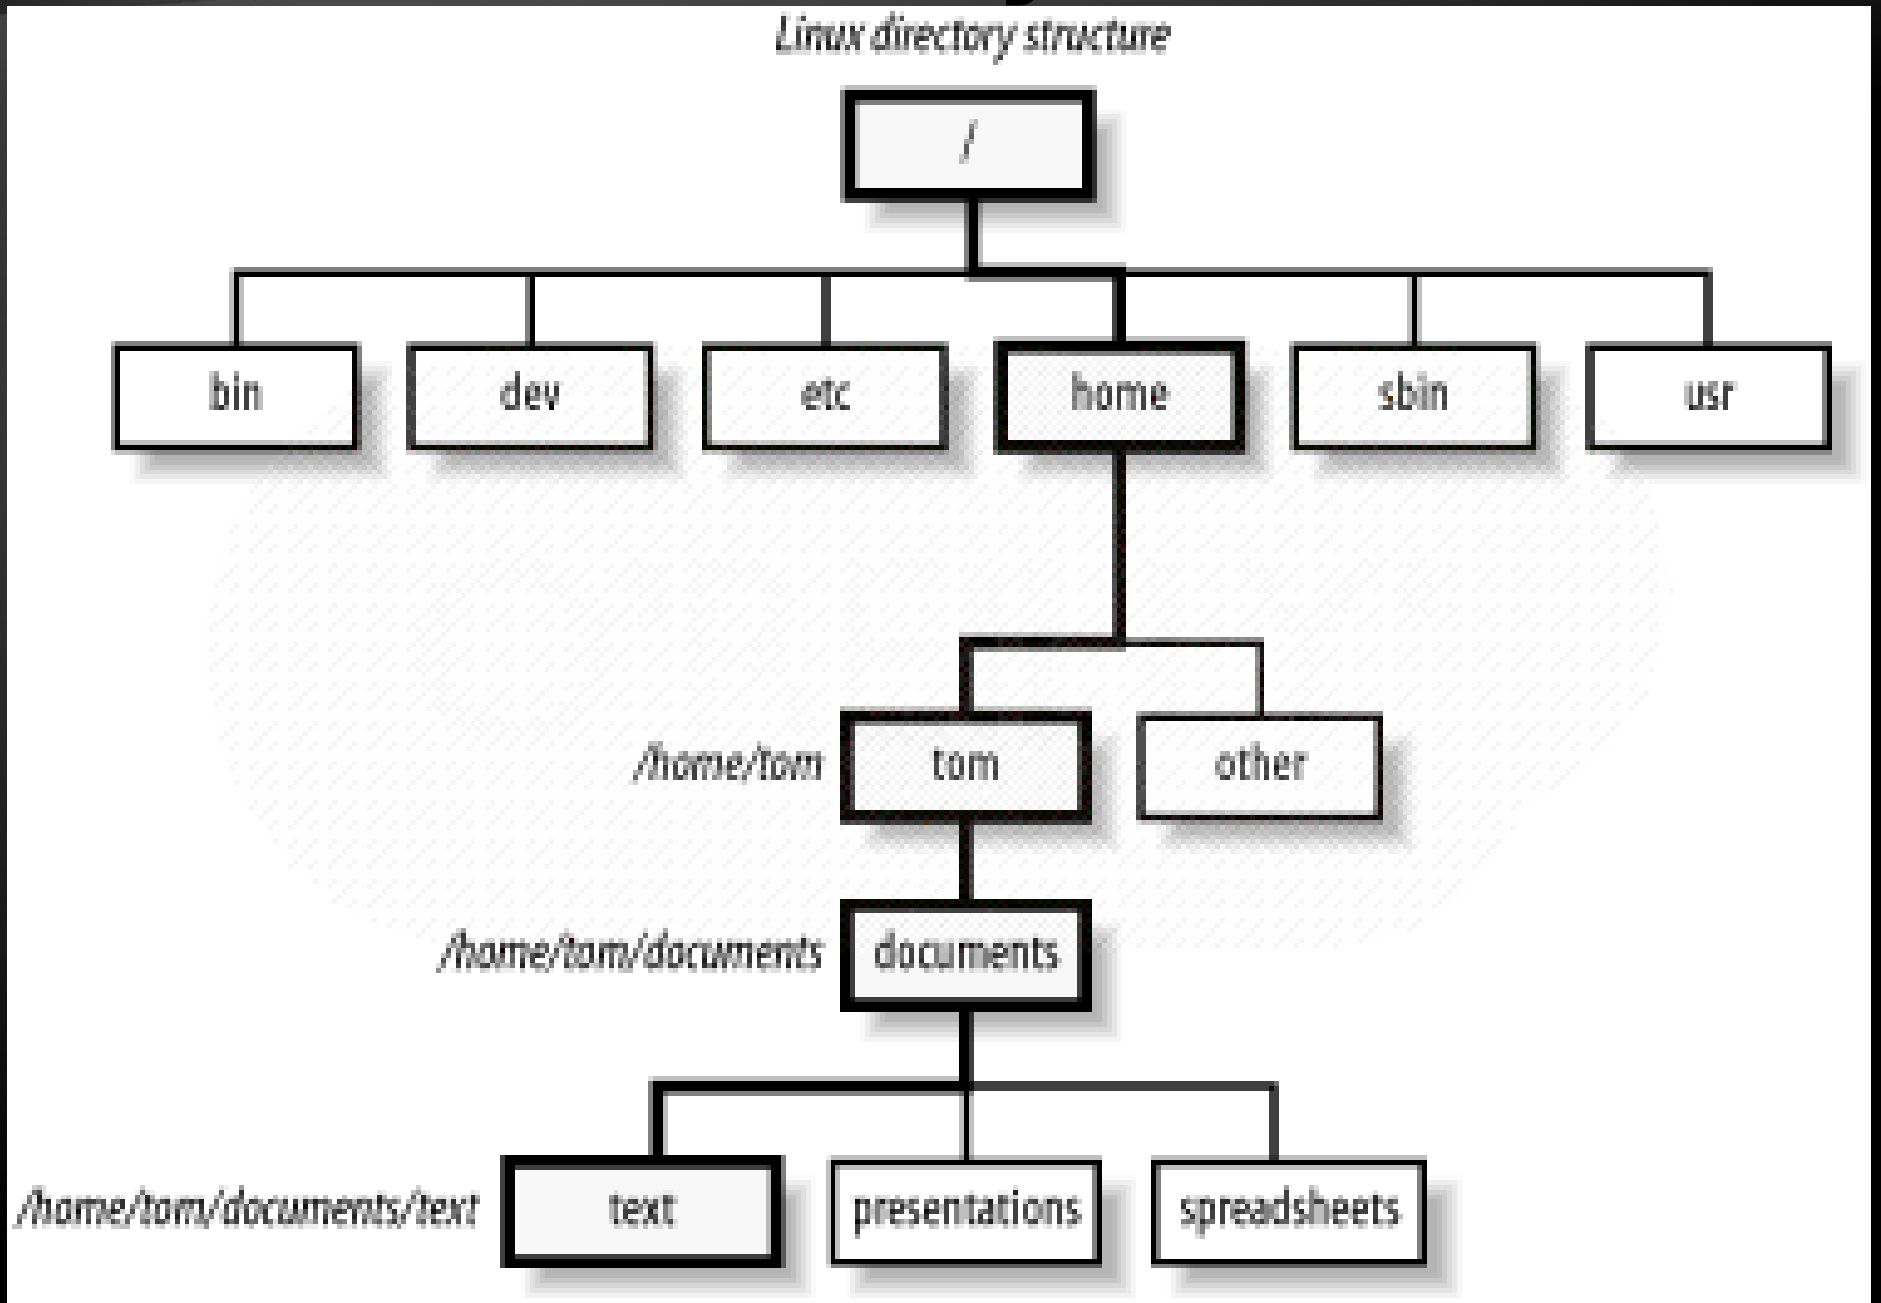

SOFTWARE UNIVERSITY

Linux Directory Structure

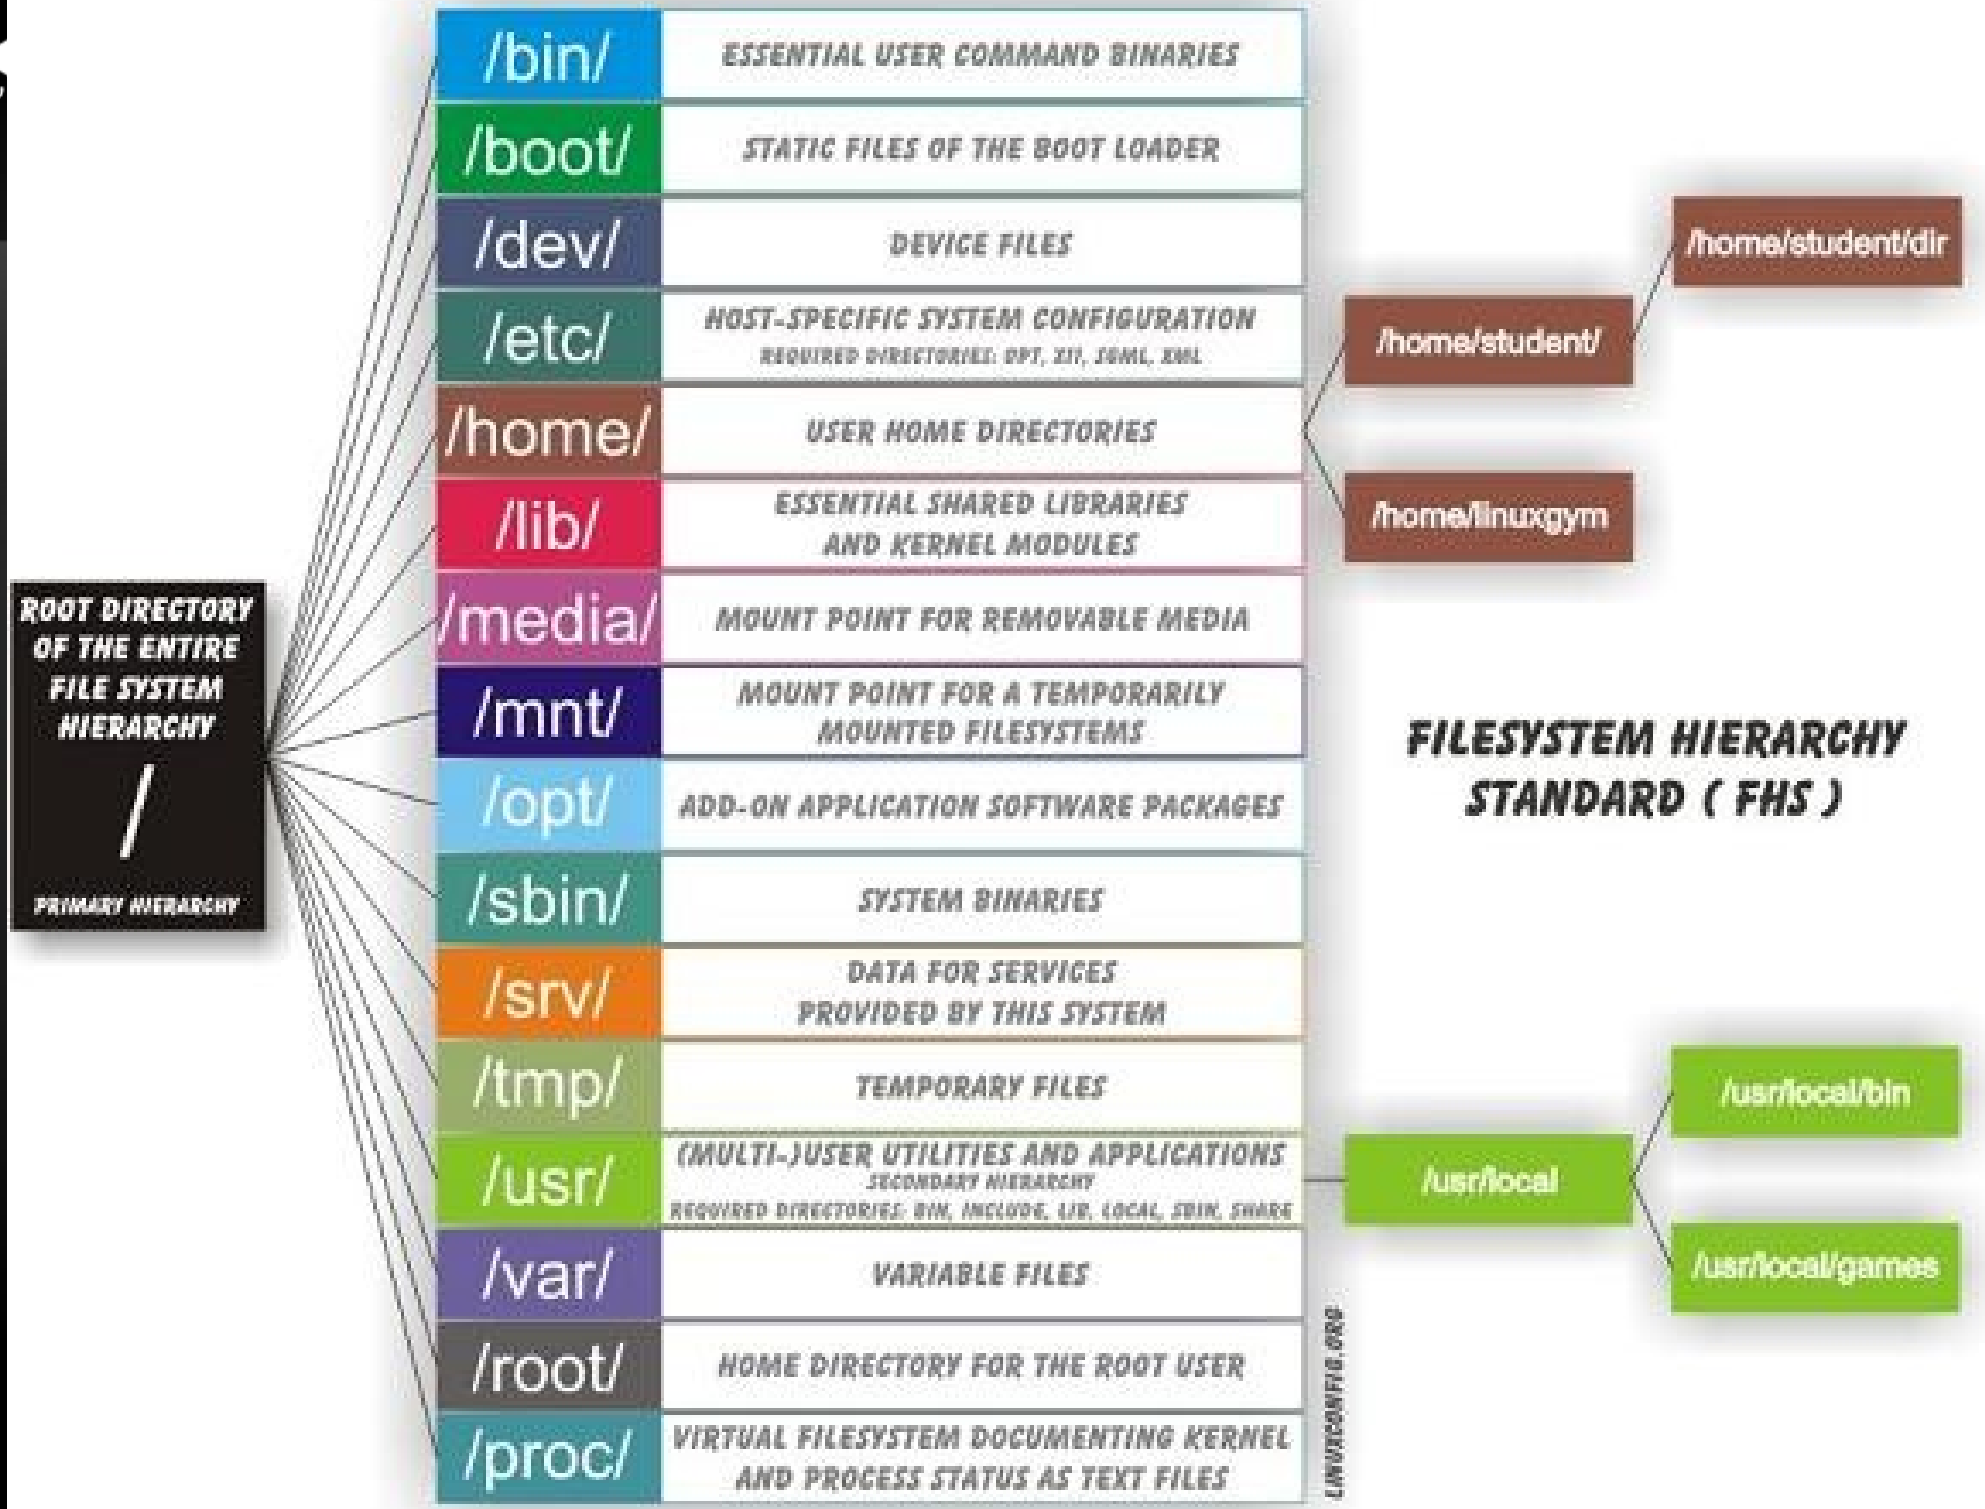
ROOT DIRECTORY OF THE ENTIRE FILE SYSTEM HIERARCHY
/
PRIMARY HIERARCHY
/bin/
ESSENTIAL USER COMMAND BINARIES
/boot/
STATIC FILES OF THE BOOT LOADER
/dev/
DEVICE FILES
/etc/
HOST-SPECIFIC SYSTEM CONFIGURATION
REQUIRED DIRECTORIES: OPT, X11, SGML, XML
/home/
USER HOME DIRECTORIES
/lib/
ESSENTIAL SHARED LIBRARIES AND KERNEL MODULES
/media/
MOUNT POINT FOR REMOVABLE MEDIA
/mnt/
MOUNT POINT FOR A TEMPORARILY MOUNTED FILESYSTEMS
/opt/
ADD-ON APPLICATION SOFTWARE PACKAGES
/sbin/
SYSTEM BINARIES
/srv/
DATA FOR SERVICES PROVIDED BY THIS SYSTEM
/tmp/
TEMPORARY FILES
/usr/
(MULTI-)USER UTILITIES AND APPLICATIONS
SECONDARY HIERARCHY
REQUIRED DIRECTORIES: BIN, INCLUDE, LIB, LOCAL, SBIN, SHARE
/var/
VARIABLE FILES
/root/
HOME DIRECTORY FOR THE ROOT USER
/proc/
VIRTUAL FILESYSTEM DOCUMENTING KERNEL AND PROCESS STATUS AS TEXT FILES
/home/student/
/home/student/dir
/home/linuxgym
/usr/local
/usr/local/bin
/usr/local/games
FILESYSTEM HIERARCHY STANDARD (FHS)
LINUXCONFIG.ORG

SOFTWARE UNIVERSITY

Links

- Wikipedia - Comparison of file systems
- Ext2 and OCFS2 on-disk layout
- Ext2 on-disk layout
- XFS on-disk structure
- ReiserFS on-disk structure
- RFSTool for Windows
- XFS Scalability
- BtrFS on-disk structure
- NILFS2 the new kid on the block
- **Usenix paper on Log-Structured File Systems**

http://softuni.bg/trainings/fasttracks/details/1060

SOFTWARE UNIVERSITY

Linux File Systems

http://softuni.bg/trainings/fasttracks/details/1060

THAT TIME AGAIN?

Yes, then we'll have a beer